

# Applied Virtual Networks

## COMP 4912

Instructor: **Dawood Sajjadi**

PhD, SMIEEE, CISSP

[ssajjaditorshizi@bcit.ca](mailto:ssajjaditorshizi@bcit.ca)

Winter-Spring 2025

**Week #1**



# Class Introduction (15 sec per student)

1. Your Name
2. You Program
3. Windows/Mac/Linux
4. What Expect to Learn



# A Brief History of Time

**Universe** created ~13.7 billion years ago.

**Planet Earth** created ~4.5 billion years ago.

**Modern Humans** showed up ~300,000 years ago.

**First Civilizations** emerged ~12,000 years ago.

**First Written records** created ~5,000 years ago.

**First Analog Computer** created ~2,200 years ago.

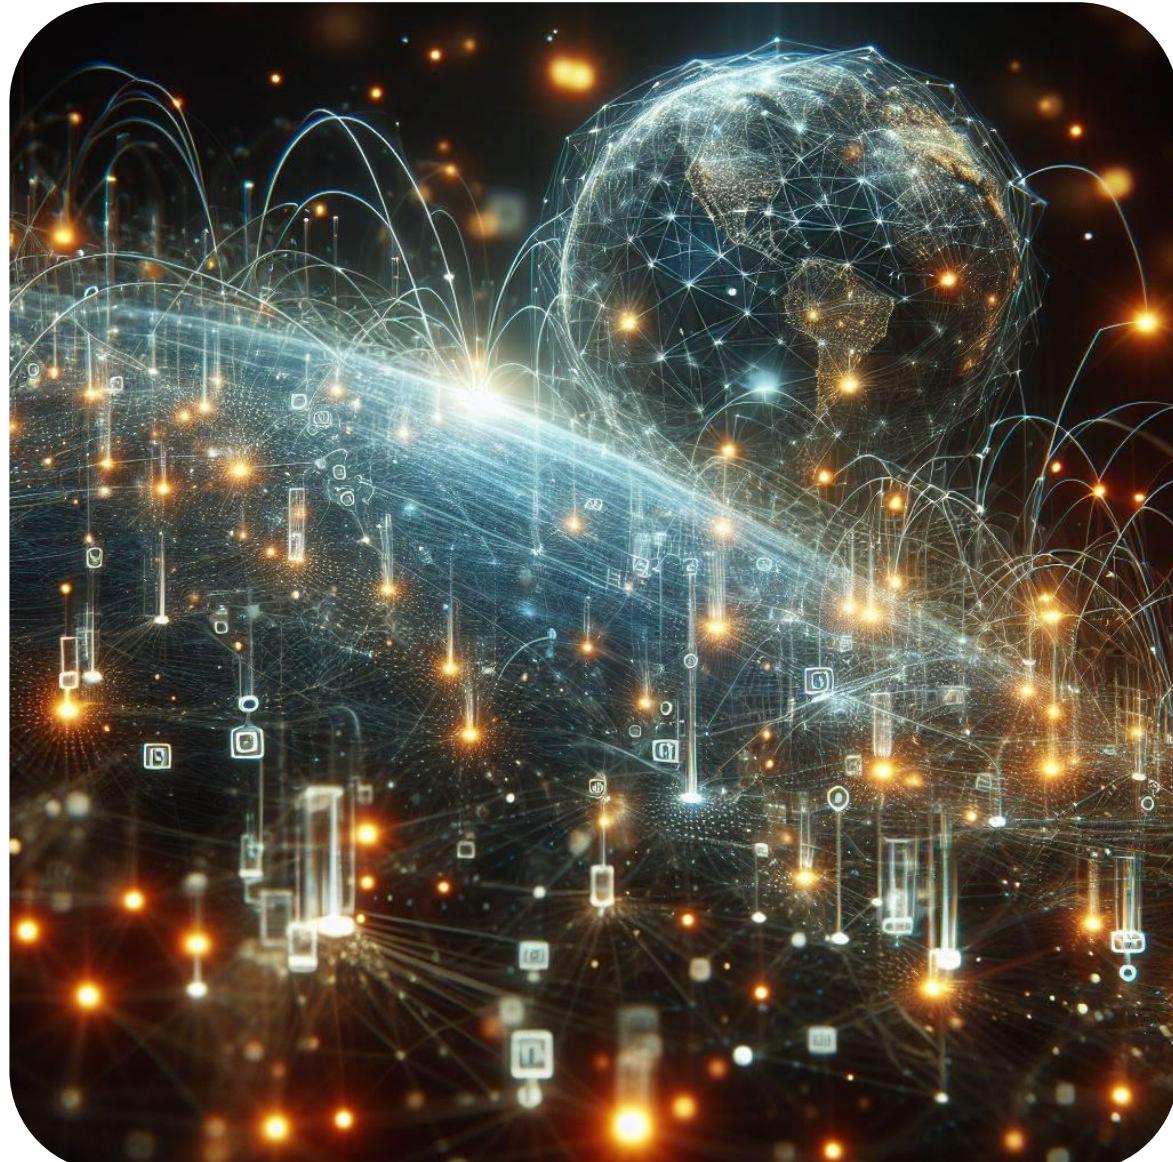
**First Digital Computer** created ~80 years ago.

**First version of the Internet** invented ~50 years ago.



# Course Logistics

- ✓ Evaluation and course components
- ✓ Lectures and labs
- ✓ Attendance
- ✓ Academic integrity
- ✓ Learning Outcome



# Evaluation

Criteria	%	Comments
Quiz	10	8 quizzes
Project	10	On a selected topic
Lab Assignments	25	5 assignments
Midterm Exam	20	Online
Final Exam	35	Online

# Lectures and Labs

- ✓ All lectures and labs are **Online**.
- ✓ Lab and lecture materials are posted **Weekly on the Learning Hub**.
- ✓ Some complementary notes may be posted on the Learning Hub during the term and students have to study them.
- ✓ For quizzes, the **Quiz tool of the Learning Hub** will be used.
- ✓ **All quizzes and exams are closed-book** unless stated otherwise.

# Attendance Requirements

- ✓ Regular **attendance in lectures and labs** is **mandatory** unless approved by the instructor.
- ✓ **Unapproved absence of 2 or more classes**  
**may result in failure or forced withdrawal from the course or program.**
- ✓ Please see Policy 5101  
Student Regulations: <http://www.bcit.ca/files/pdf/policies/5101.pdf>

# Academic Integrity

- ✓ Violation of academic integrity, including plagiarism, dishonesty in assignments, examinations, or other academic performances are prohibited and will be handled in accordance with

**Policy 5104 - Academic Integrity and Appeals, and accompanying procedures.**

- ✓ Please do not share the course materials with people outside of your class.

# How the Internet Works?

## Networking Protocols

BGP	OSPF	FTP	SMTP	
DNS	IPSec	TLS	SSH	IMAP
	SIP		HTTPs	
			SNMP	



## Networking/Computing Devices

Servers  
Containers  
**Firewalls**  
WAF

VMs  
Storage  
UTM  
VPN

Switch  
Router  
Load Balancer

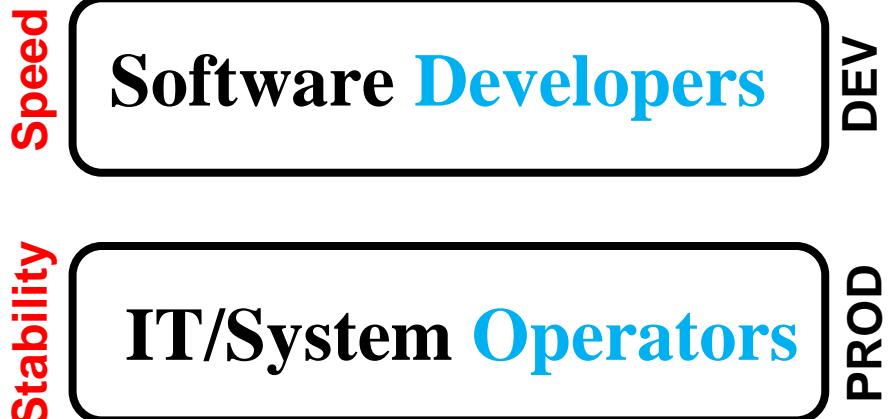
## Internet Applications/Sites



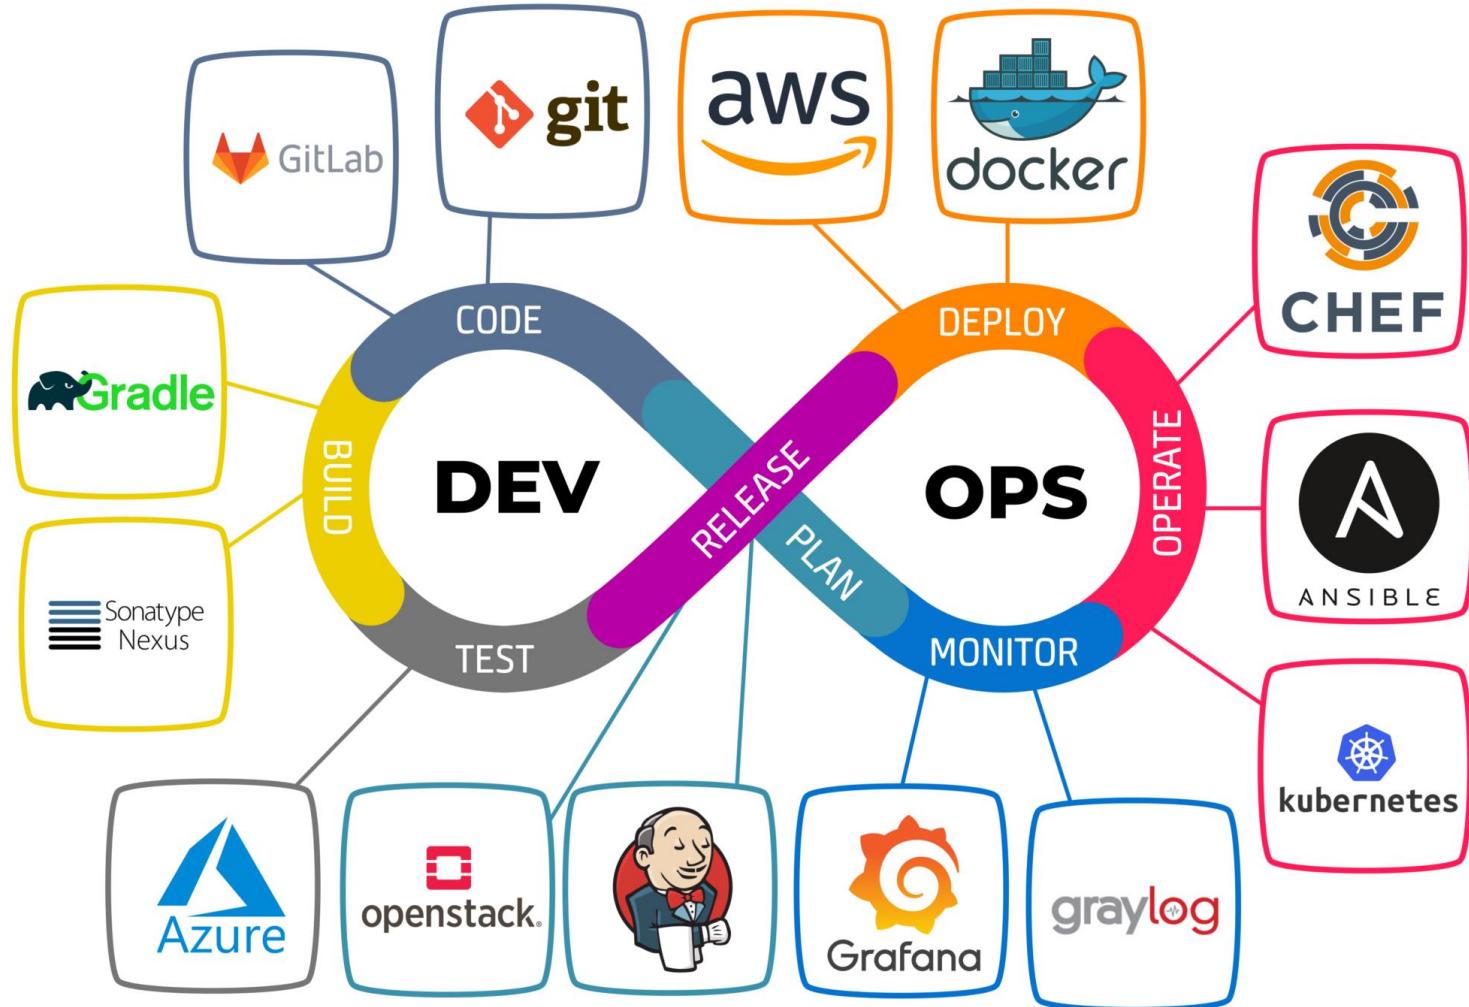
## Data Centers as Engines of the Internet



# What is DevOps/SRE?



A **DevOps Engineer** bridges the gap between software development and IT operations to make software delivery faster, more reliable, and efficient.



# What You Learn From This Course Will Help You to ...

Devops in Canada  
7,272 results

Job Alert Off

 OneSpan  
OneSpan  
Montreal, QC  
1 company alum works here  
Promoted • 18 applicants

 Azure Devops, Infrastructure and Security Specialist  
GS1 Canada  
Montreal, QC  
1 alum works here  
Promoted • 5 applicants

 DevOps Engineer  
Avanade  
Vancouver, BC  
4 alumni work here  
Promoted • 4 applicants

 Cloud Engineer, Elastic Engineering  
Onica, a Rackspace Technology Company  
Vancouver, BC  
2 alumni work here  
4 weeks ago • 3 applicants

 DevOps Engineering Director - K8s and Python  
Canonical  
Vancouver, BC Remote  
1 connection works here  
1 month ago • 0 applicants

 Intermediate DevOps/SRE Engineer (Linux)  
Global Relay  
Vancouver, BC

Devops in United States  
72,023 results

Job Alert Off

 Senior Staff DevOps Engineer (Information Security)  
Palo Alto Networks  
Santa Clara, CA  
2 alumni work here  
Promoted • 2 applicants • Easy Apply

 DevOps, Infrastructure Manager  
Optello  
Boise, ID  
Medical, Vision, Dental, 401(k)  
Actively recruiting  
Promoted • 3 applicants • Easy Apply

 Principal Consultant - DevOps - Nationwide Opportunities  
Amazon Web Services (AWS)  
San Francisco, CA  
Medical benefit  
6 connections work here  
Promoted • 0 applicants

 Security Infrastructure DevOps Engineer  
Apple  
Cupertino, CA  
2 connections work here  
3 hours ago

 Platform Services QA Lead and DevOps Engineer  
Viasat Inc.  
Remote, OR  
1 alum works here  
2 weeks ago • 0 applicants

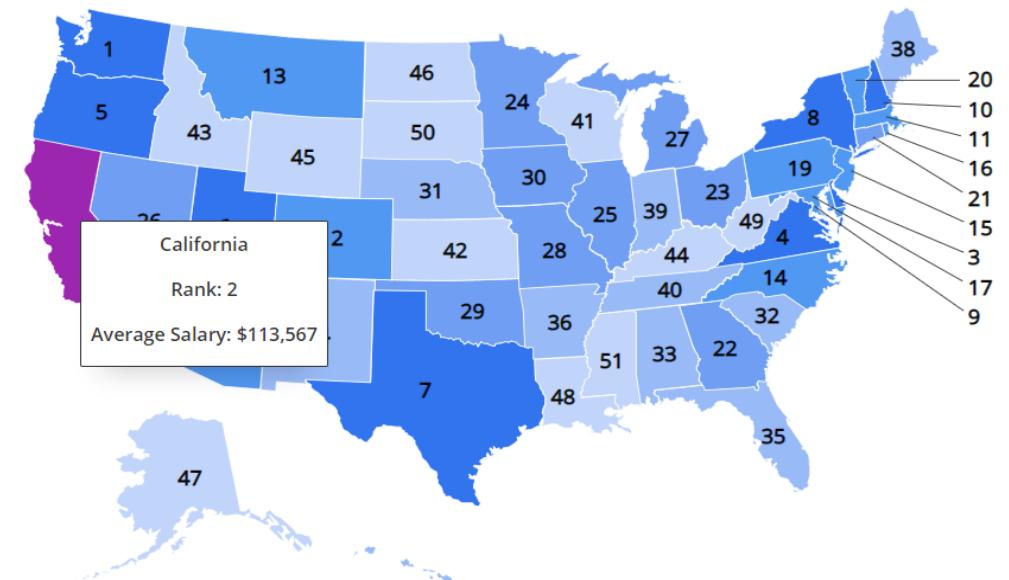
 Senior CI/CD DevOps Engineer  
Microsoft

## DevOps/SRE



## BEST STATES FOR A DEVOPS ENGINEER

Rank 1 - 10    11 - 20    21 - 30    31 - 40    41 - 51



# Learning Outcomes of the course

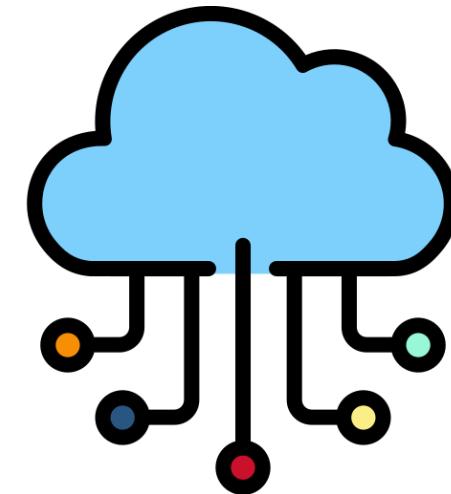
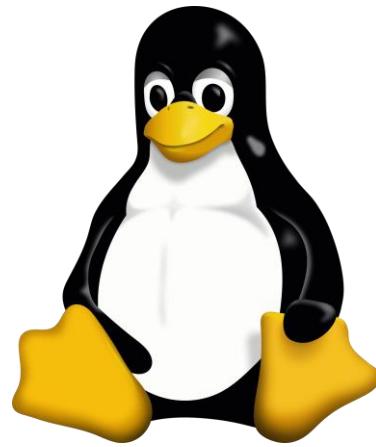
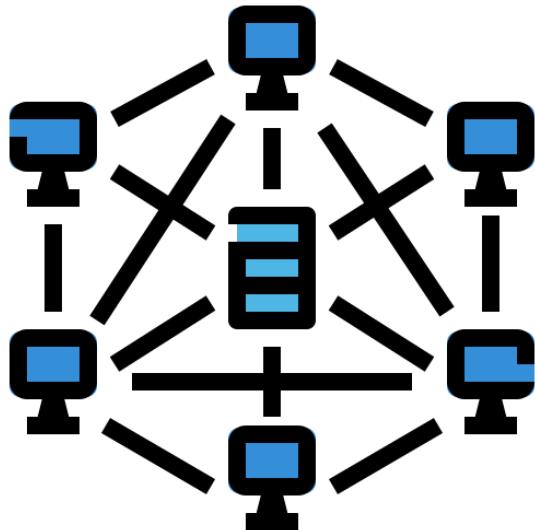
- ✓ Describe the different types of **Hypervisors** and Cloud technologies.
- ✓ Determine environment load potential and the methods required to **scale**.
- ✓ Describe the different technologies and protocols used for **network** and **storage virtualization**.
- ✓ Identify common **performance** bottlenecks and different ways to detect and prevent them.
- ✓ Explain fundamental concepts surrounding **High availability** (HA) and **Disaster Recovery** (DR) in a virtual environment.
- ✓ Describe a **backup strategy** for a virtual environment.
- ✓ Explain various monitoring protocols (**SNMP** and **syslog**) and tools.
- ✓ Describe fundamental concepts in **Public and Private Cloud** environments.
- ✓ Conduct a performance assessment combined with a review of **security baselines** in cloud environments.
- ✓ Describe **container management and orchestration** for cloud-based network solutions.
- ✓ Automate **resource provisioning** process in virtual environments.
- ✓ Explain the **Infrastructure-as-a-Code (IaaS)** concept for real-world applications.

# Learning Outcomes of Week #1

1. Describe how Internet works?
2. Explain how big companies serve **Billions** of people on the Internet.
3. Describe the general concept of Virtualization on the Internet.
4. Understand the Concept of Cloud Computing/Technology.
5. Explain different types of Cloud environments/architecture.

# Quick Recap on Requirements of this Course

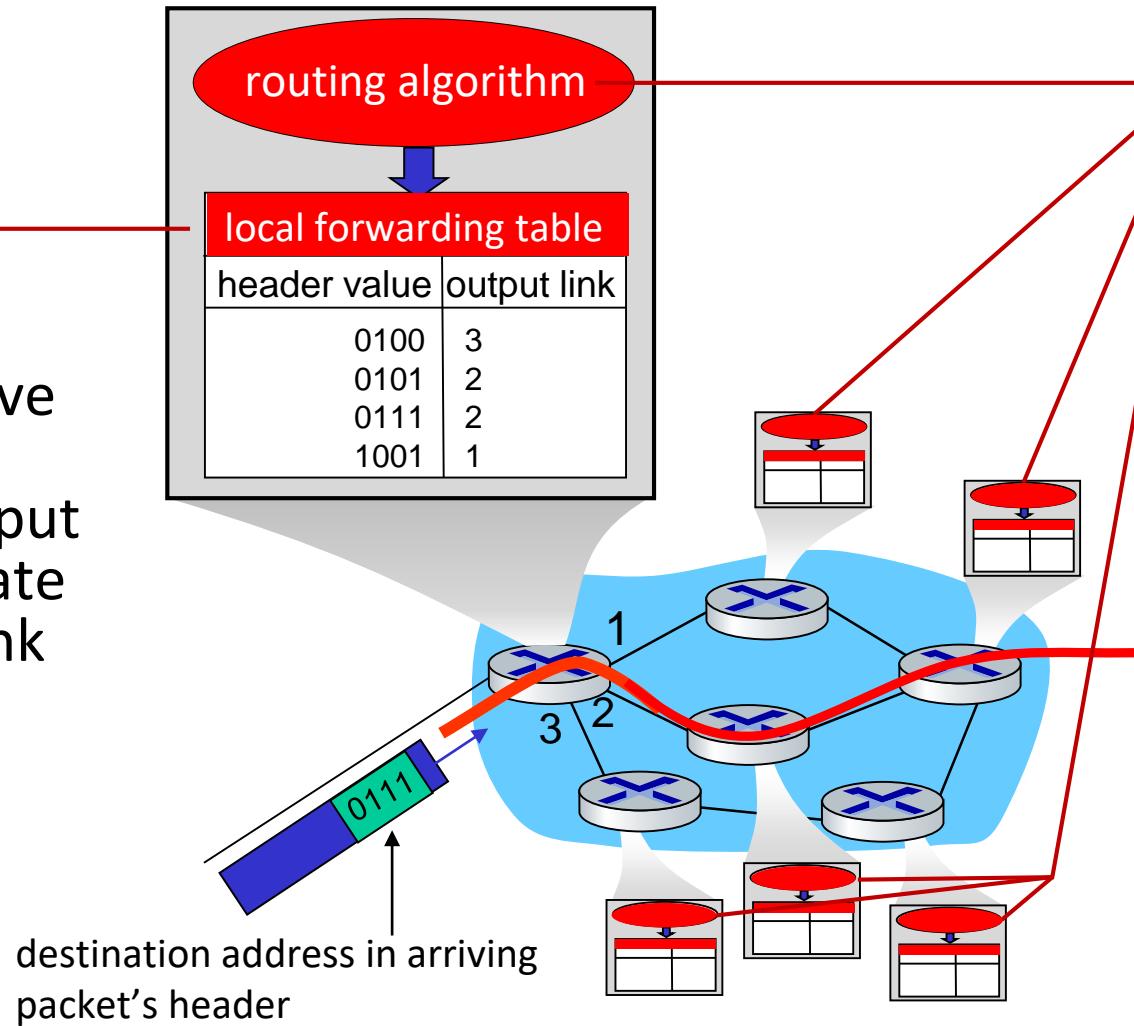
1. Basics of Computer Networking
2. Basics of Linux Operating System
3. Cloud Computing Platforms



# Two key network-core functions

## Forwarding:

- aka “switching”
- *local* action: move arriving packets from router’s input link to appropriate router output link



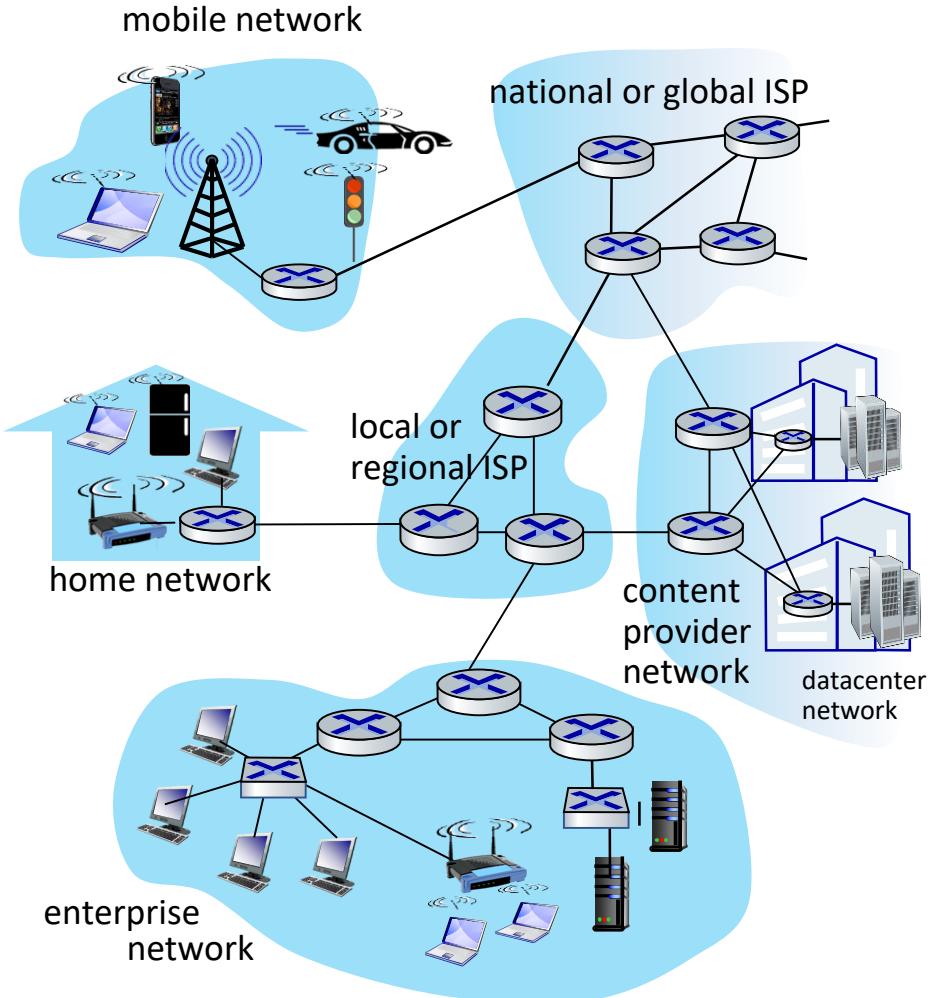
## Routing:

- *global* action: determine source-destination paths taken by packets
- routing algorithms

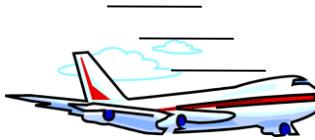




1. Visit **ping.pe** web site
2. Enter **Google.com** in the box
3. Check the 2<sup>nd</sup> column from the right (**show**) for different Geo locations



# Example: organization of air travel



*end-to-end transfer of person plus baggage*

ticket (purchase)

baggage (check)

gates (load)

runway takeoff

airplane routing

ticket (complain)

baggage (claim)

gates (unload)

runway landing

airplane routing

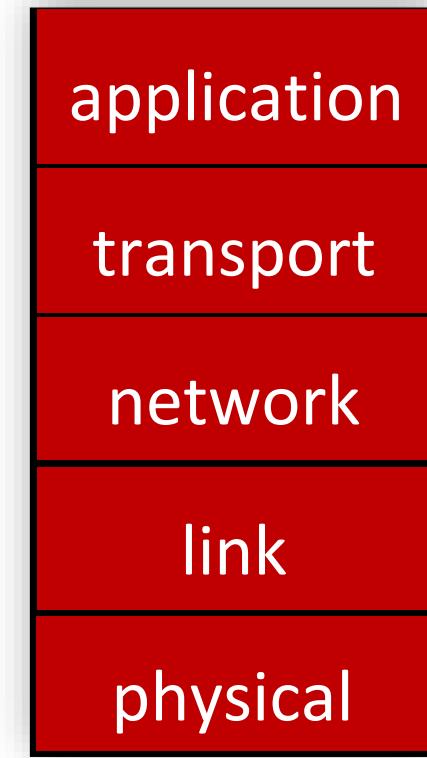
airplane routing

How would you *define/discuss* the system of airline travel?

- a series of steps, involving many services

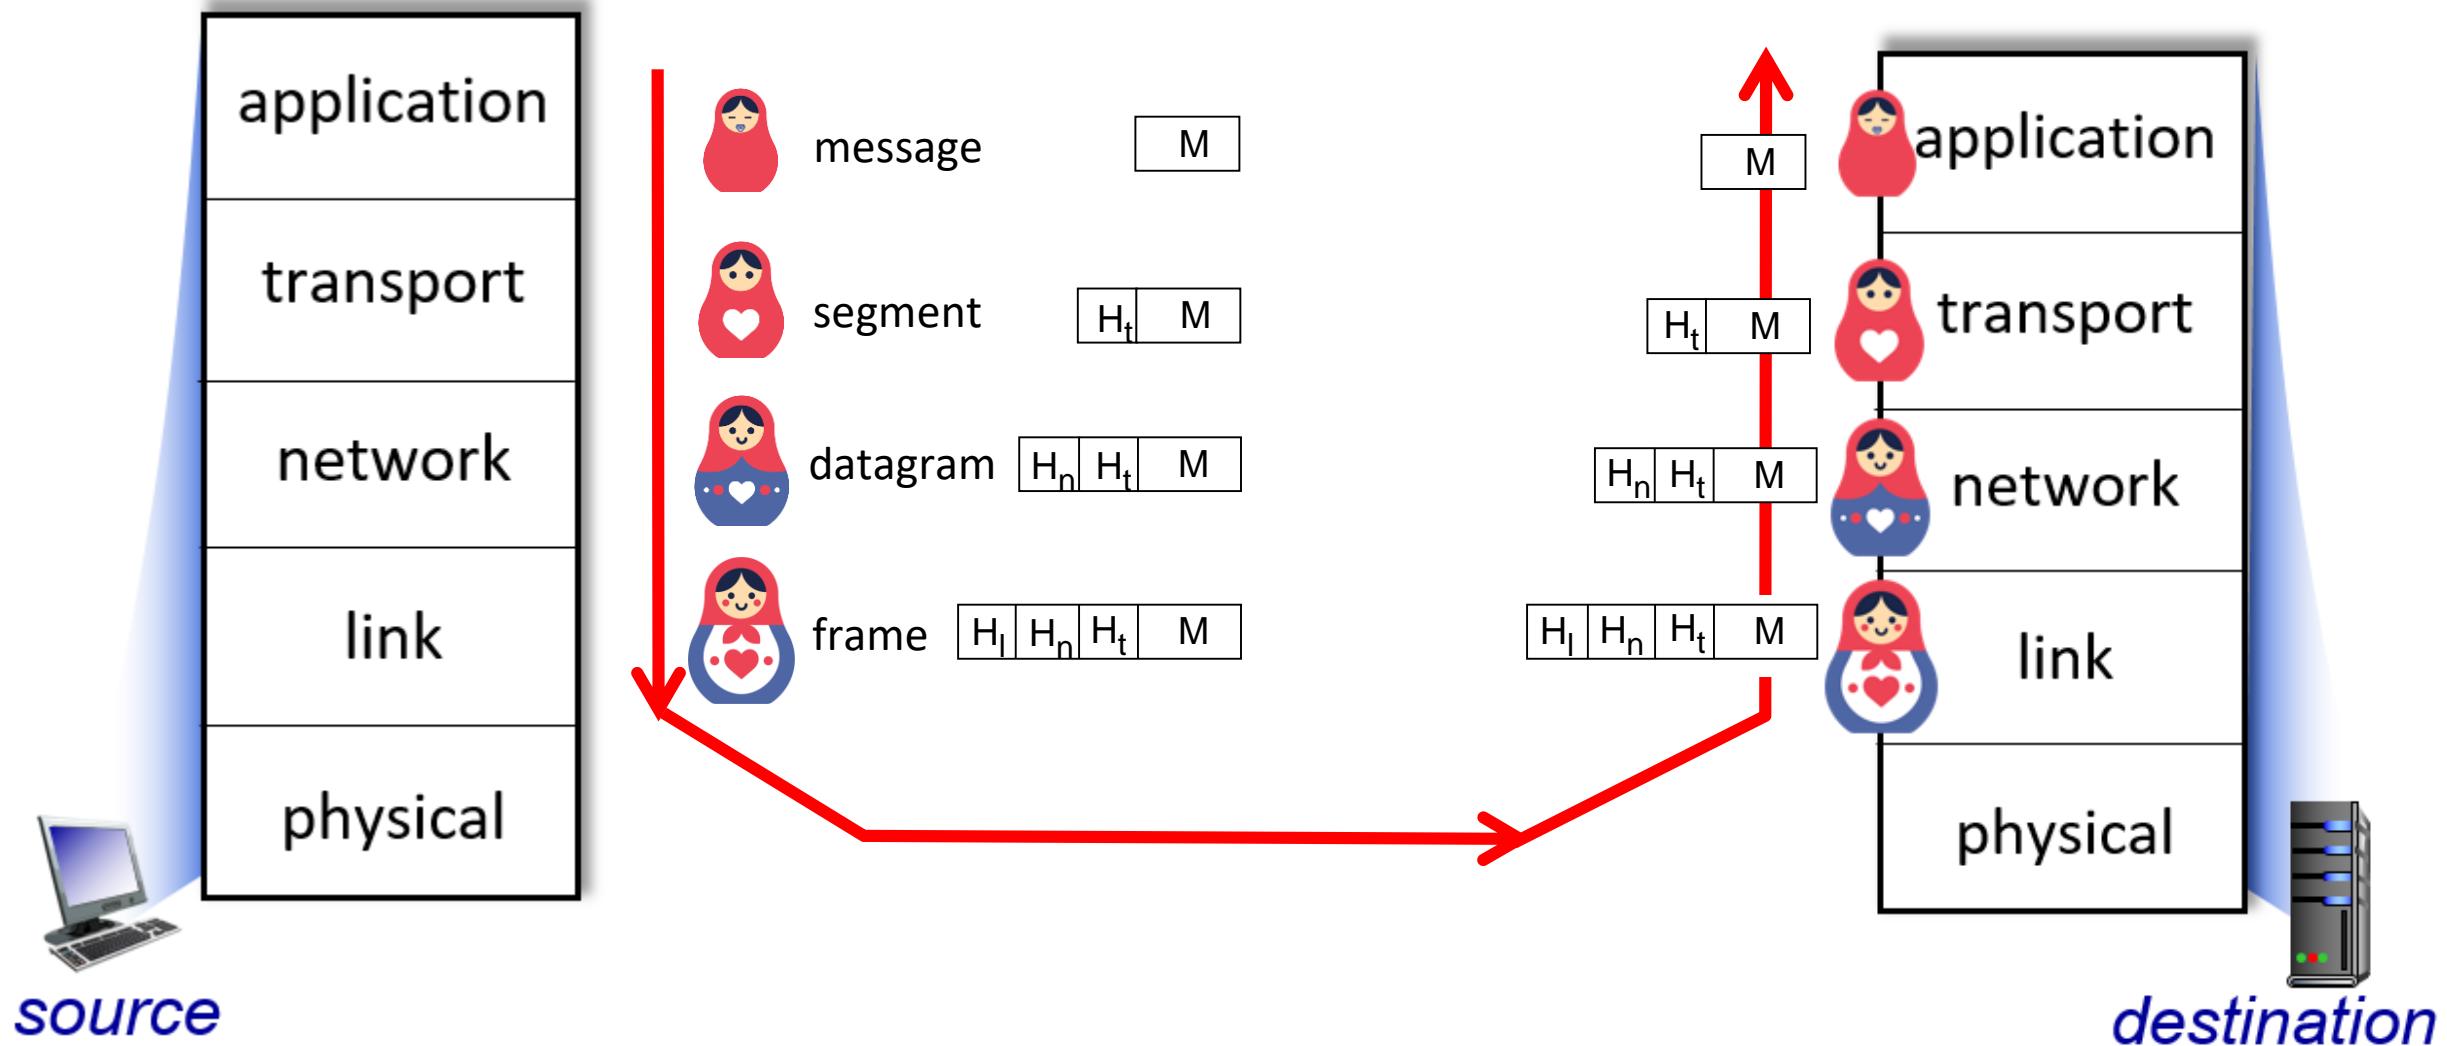
# Layered Internet protocol stack

- *application*: supporting network applications
  - HTTP, IMAP, SMTP, DNS
- *transport*: process-process data transfer
  - TCP, UDP
- *network*: routing of datagrams from source to destination
  - IP, routing protocols
- *link*: data transfer between neighboring network elements
  - Ethernet, 802.11 (WiFi), PPP
- *physical*: bits “on the wire”



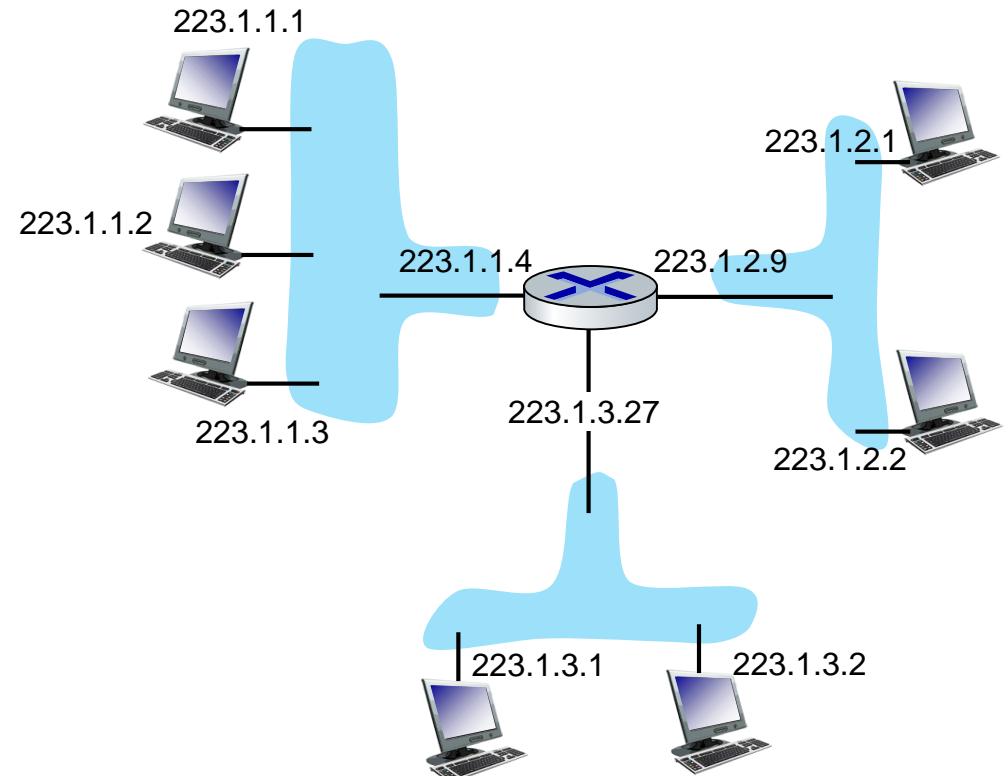
Five-Layer Internet Protocol Stack

# Services, Layering and Encapsulation



# IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)



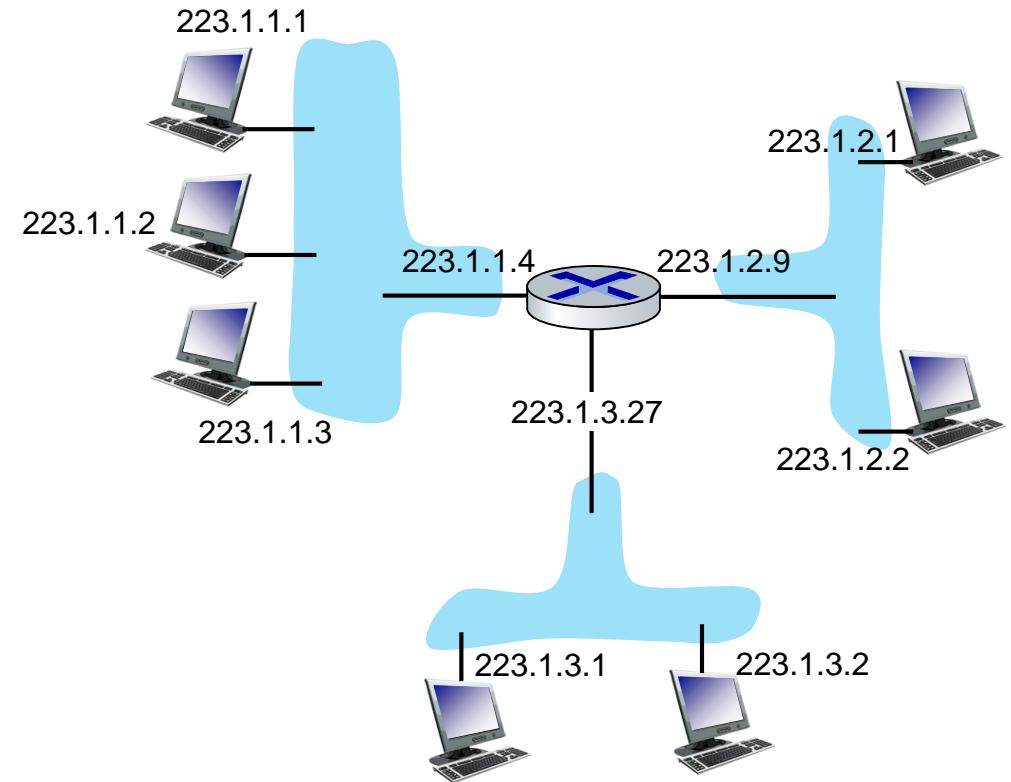
dotted-decimal IP address notation:

$223.1.1.1 = \underline{11011111} \underline{00000001} \underline{00000001} \underline{00000001}$

223      1      1      1

# IP addressing: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)



dotted-decimal IP address notation:

223.1.1.1 =   
                  |                |                |  
          223      1      1      1

# IP addressing: CIDR

CIDR: Classless InterDomain Routing (pronounced “cider”)

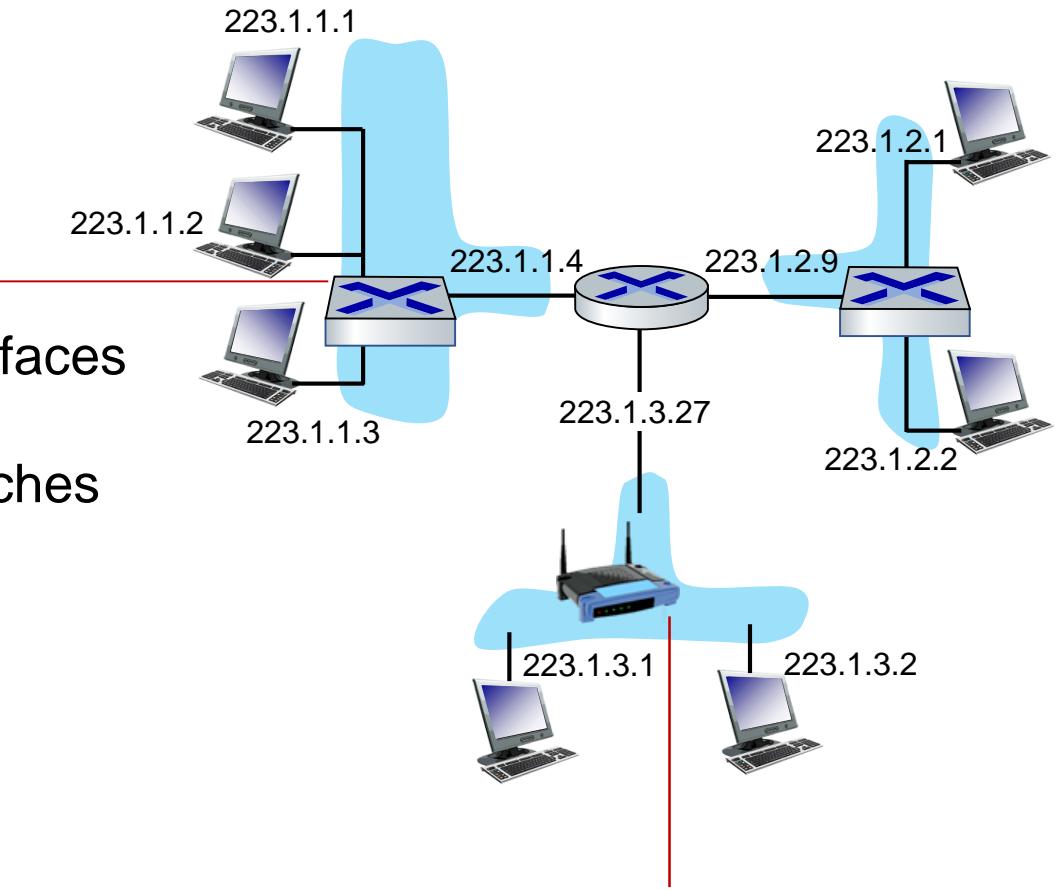
- subnet portion of address of arbitrary length
- Variable Length Subnet Masks (VLSM)
- address format:  $a.b.c.d/x$ , where  $x$  is # bits in subnet portion of address



# IP addressing: introduction

Q: how are interfaces actually connected?

A: wired  
Ethernet interfaces  
connected by  
Ethernet switches



A: wireless WiFi interfaces  
connected by WiFi base station

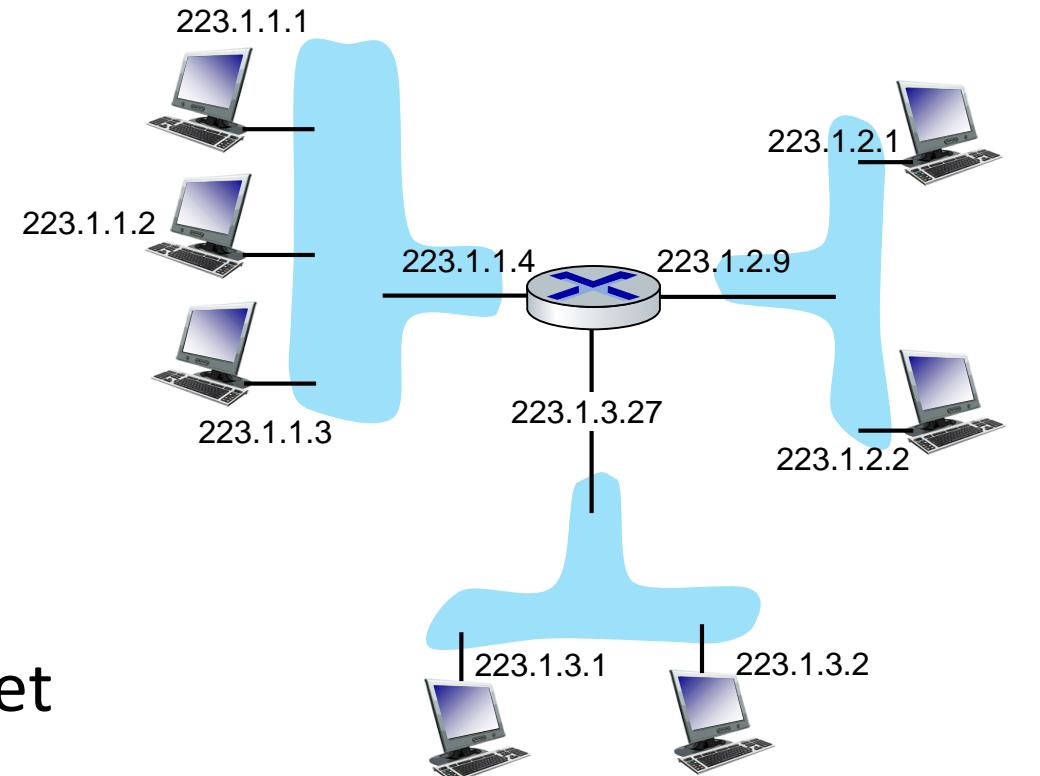
# Subnets

- *What's a subnet ?*

- device interfaces that can physically reach each other **without passing through an intervening router**

- IP addresses have structure:

- **subnet part:** devices in same subnet have common high order bits
  - **host part:** remaining low order bits

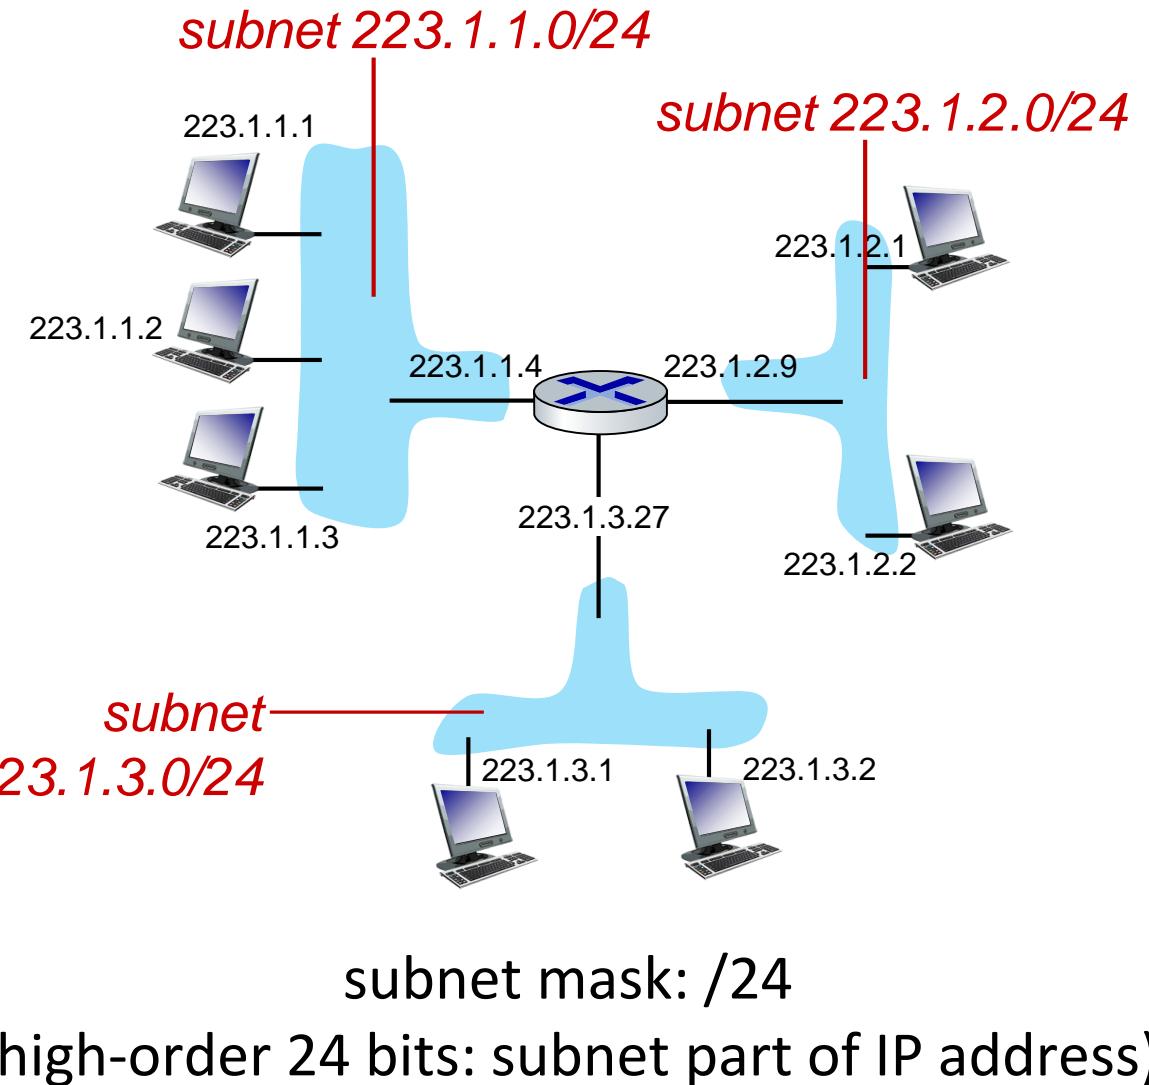


network consisting of 3 subnets

# Subnets

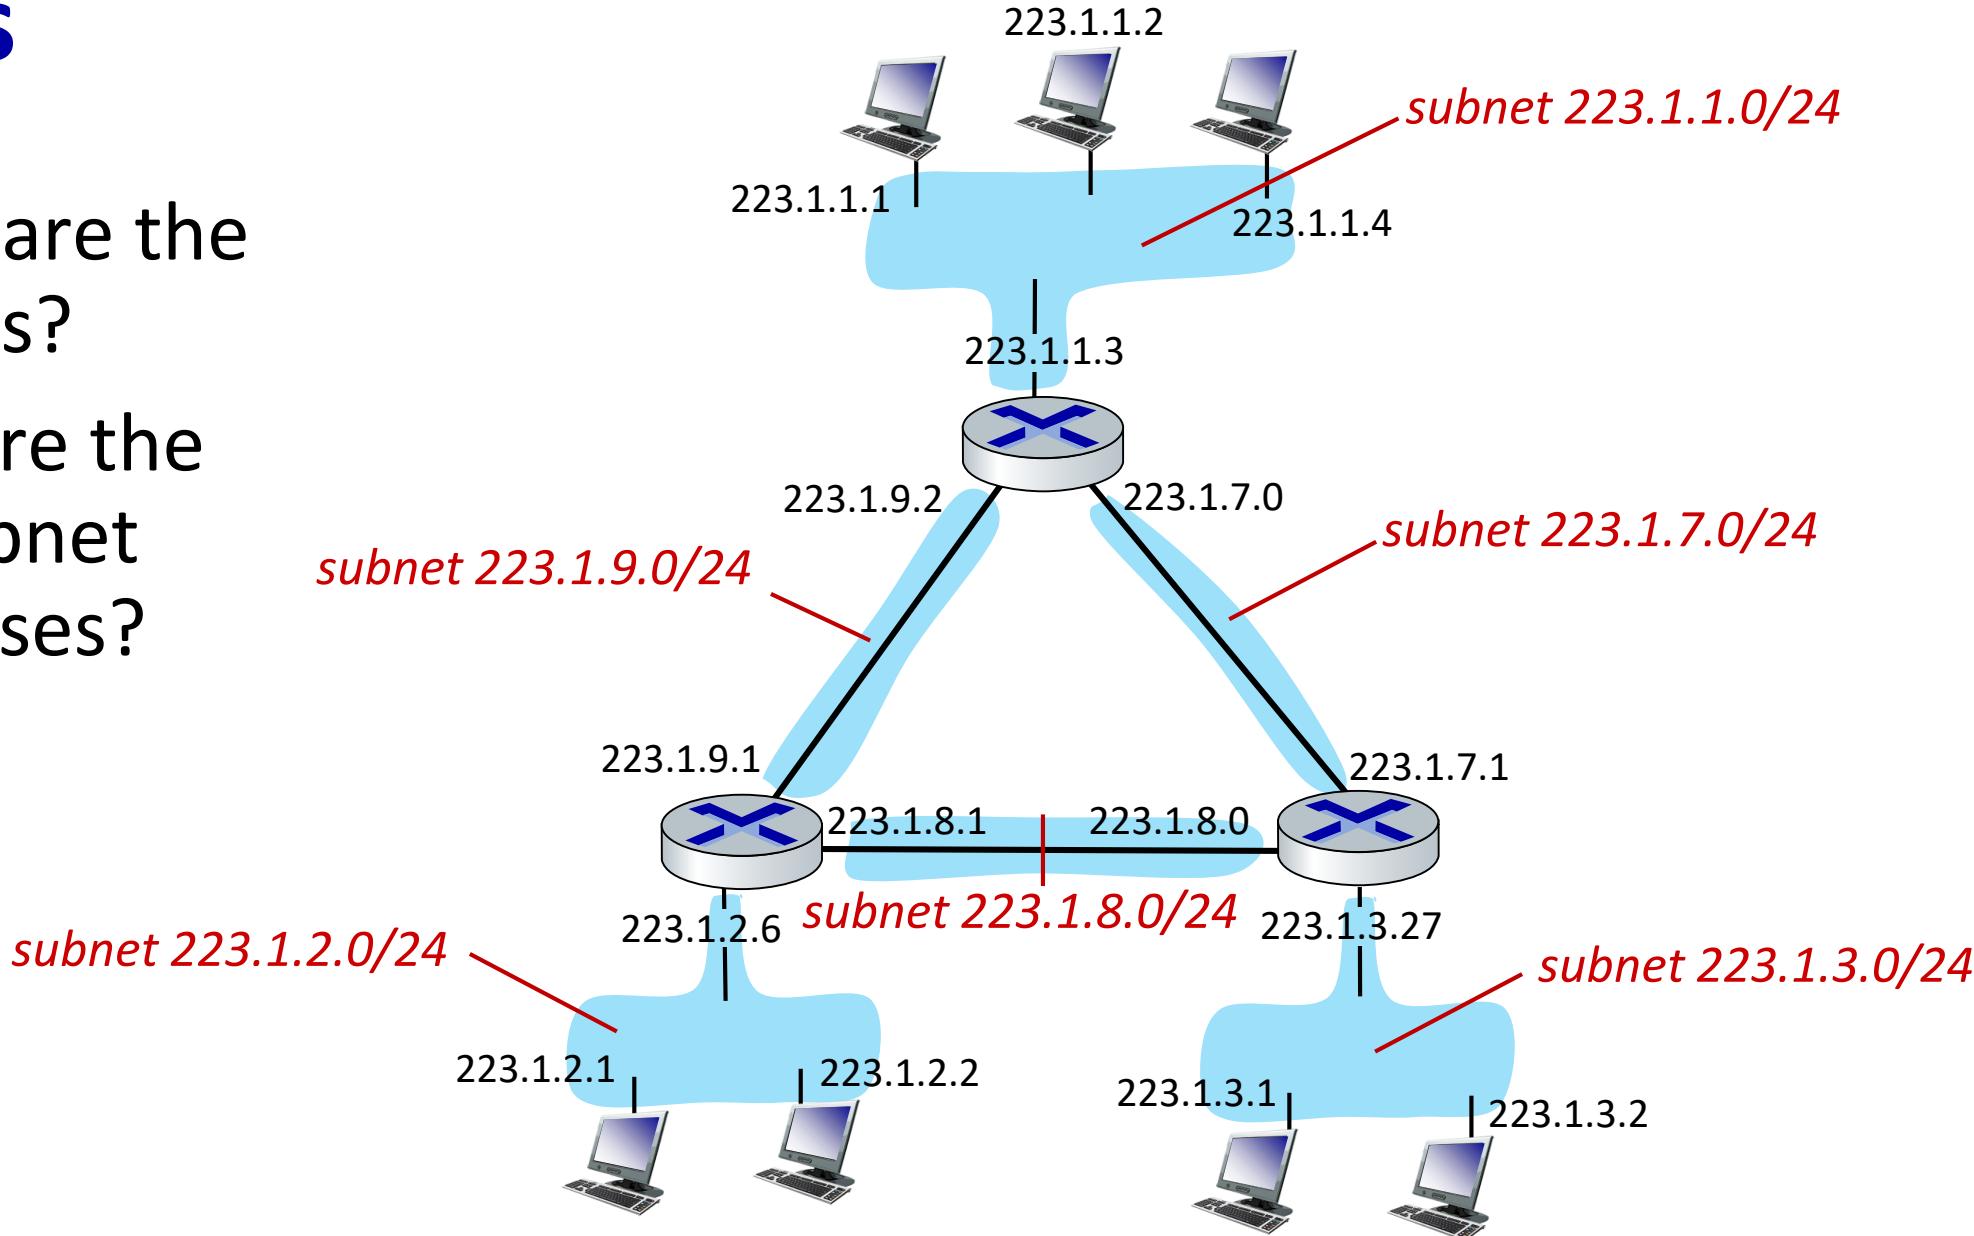
*Recipe for defining subnets:*

- detach each interface from its host or router, creating “islands” of isolated networks
- each isolated network is called a *subnet*



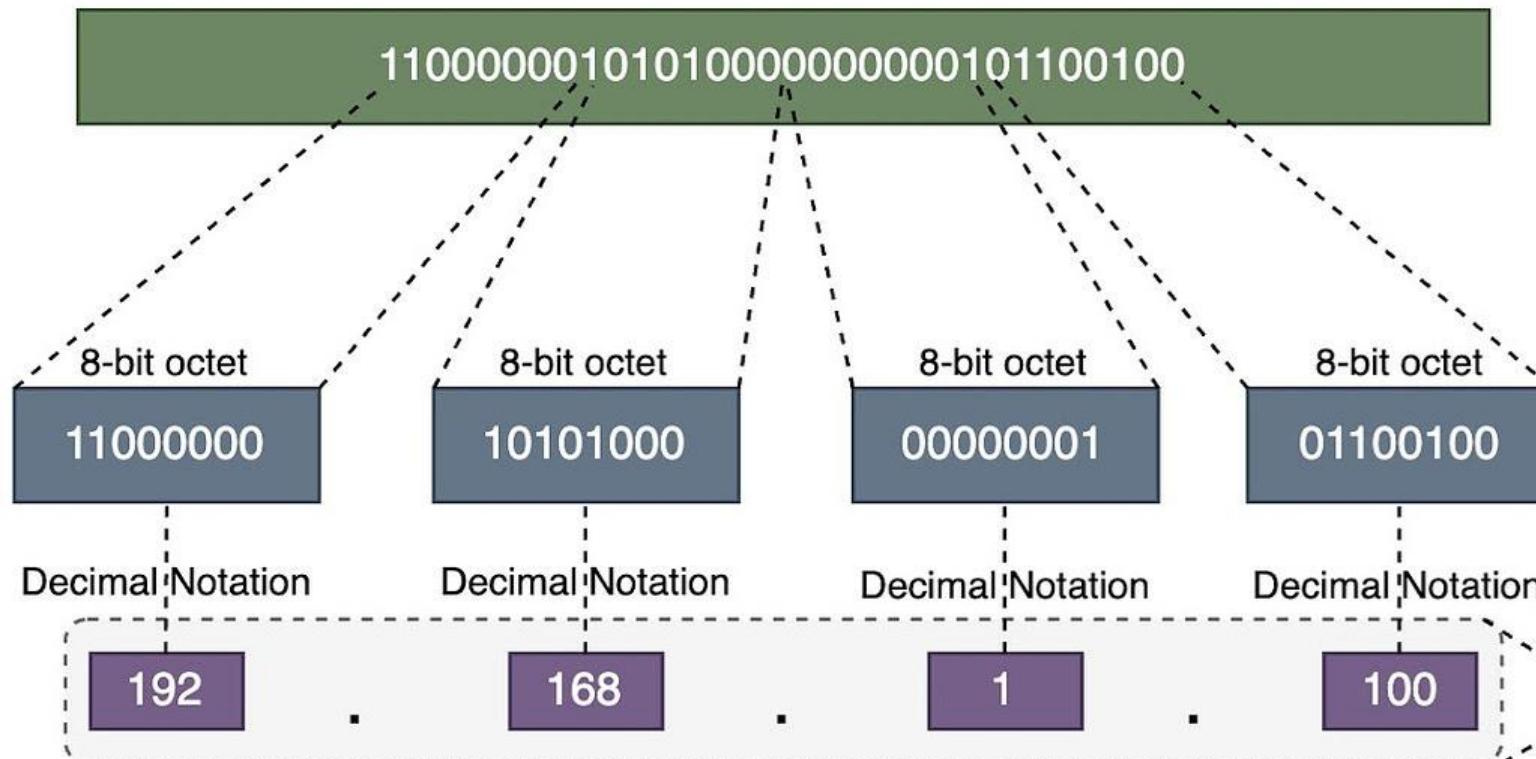
# Subnets

- where are the subnets?
- what are the /24 subnet addresses?



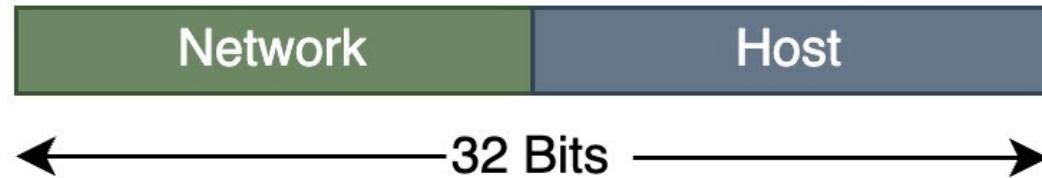
128	64	32	16	8	4	2	1	Decimal Value
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	128
1	1	0	0	0	0	0	0	192
1	1	1	0	0	0	0	0	224
1	1	1	1	0	0	0	0	240
1	1	1	1	1	0	0	0	248
1	1	1	1	1	1	0	0	252
1	1	1	1	1	1	1	0	254
1	1	1	1	1	1	1	1	255

### 32-bit IPv4 address

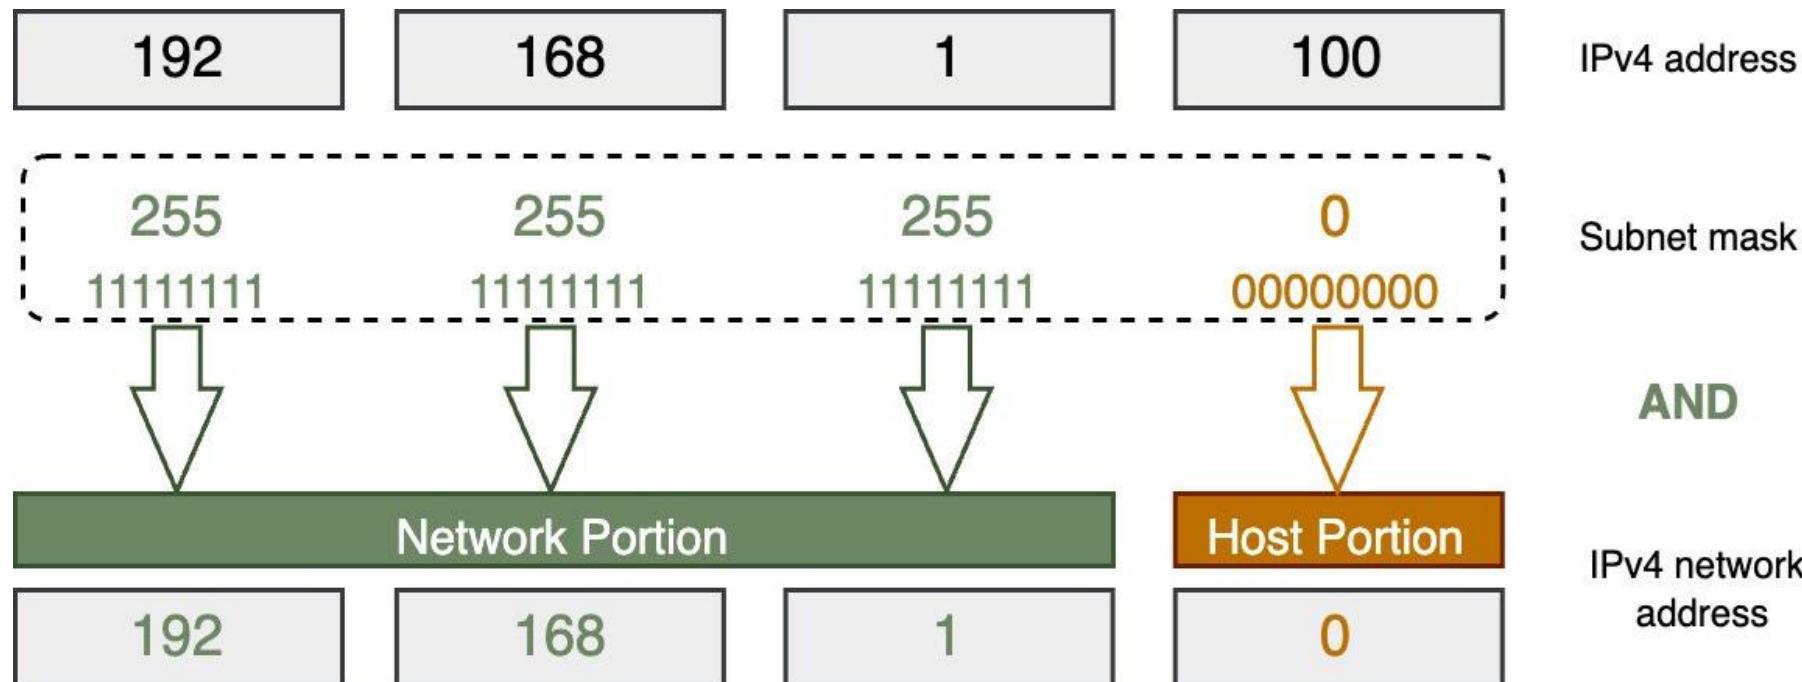


# IP addressing: introduction

## IPv4 Address

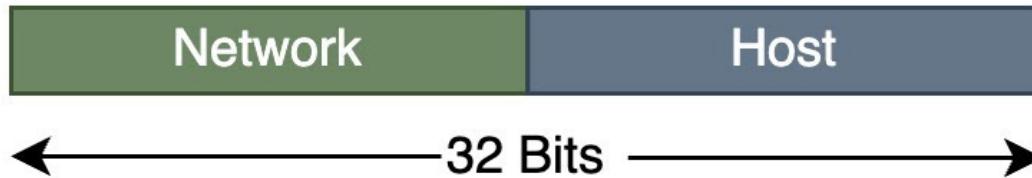


128	64	32	16	8	4	2	1	Decimal Value
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	128
1	1	0	0	0	0	0	0	192
1	1	1	0	0	0	0	0	224
1	1	1	1	0	0	0	0	240
1	1	1	1	1	0	0	0	248
1	1	1	1	1	1	0	0	252
1	1	1	1	1	1	1	0	254
1	1	1	1	1	1	1	1	255

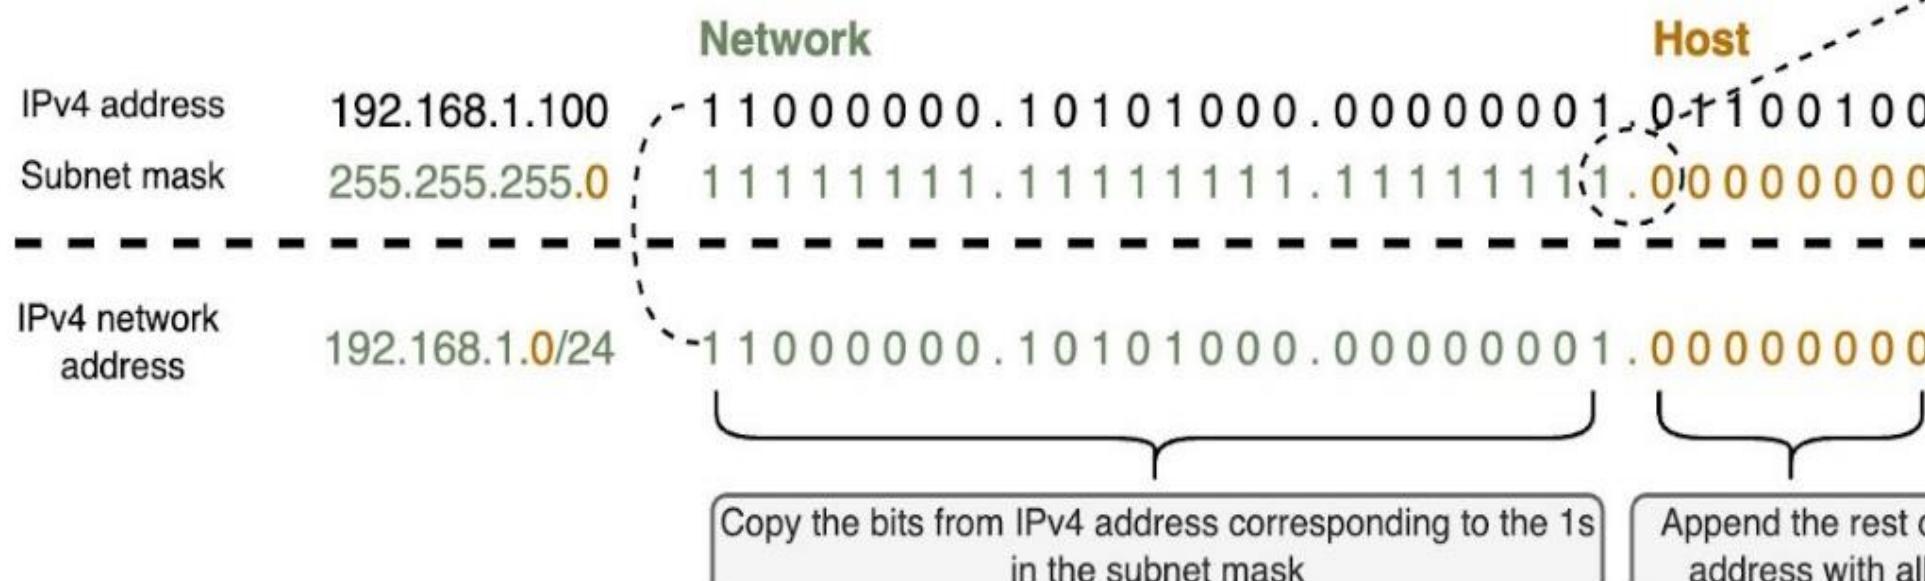


# IP addressing: introduction

## IPv4 Address



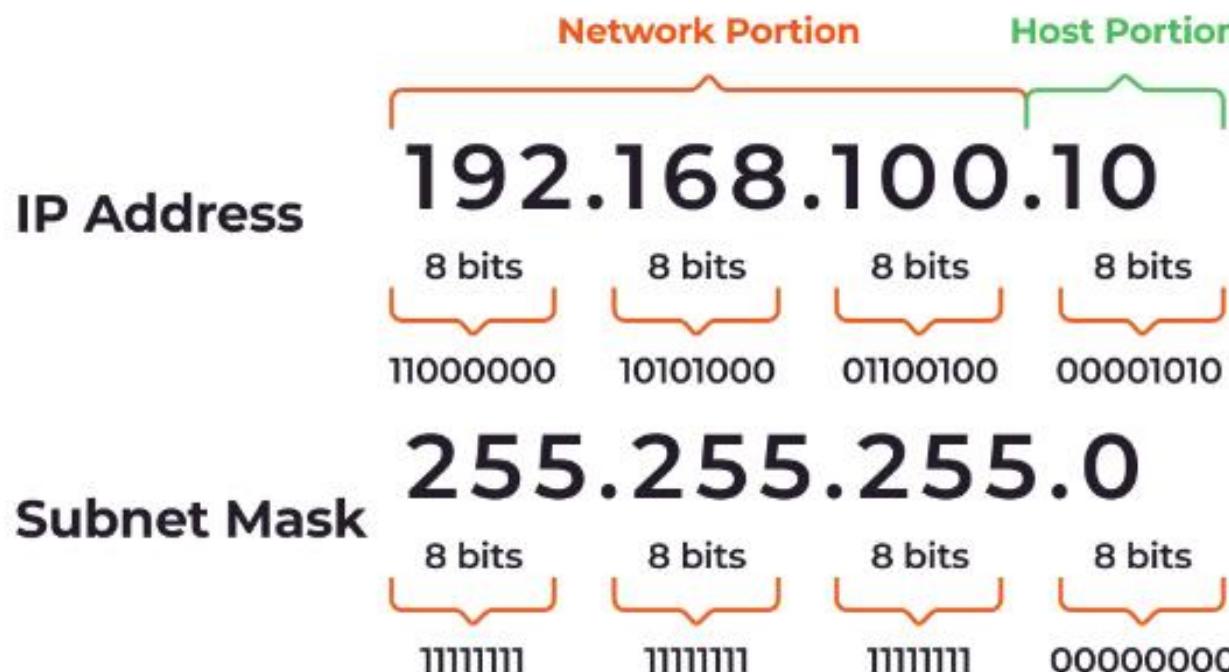
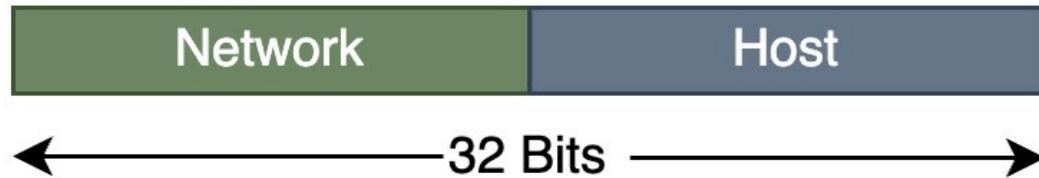
128	64	32	16	8	4	2	1	Decimal Value
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	128
1	1	0	0	0	0	0	0	192
1	1	1	0	0	0	0	0	224
1	1	1	1	0	0	0	0	240
1	1	1	1	1	0	0	0	248
1	1	1	1	1	1	0	0	252
1	1	1	1	1	1	1	0	254
1	1	1	1	1	1	1	1	255



## Network Address Calculation

# IP addressing: introduction

## IPv4 Address



128	64	32	16	8	4	2	1	Decimal Value
0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	128
1	1	0	0	0	0	0	0	192
1	1	1	0	0	0	0	0	224
1	1	1	1	0	0	0	0	240
1	1	1	1	1	0	0	0	248
1	1	1	1	1	1	0	0	252
1	1	1	1	1	1	1	0	254
1	1	1	1	1	1	1	1	255

Network IP → 192.168.100.0

Broadcast IP → 192.168.100.255

First Useable IP → 192.168.100.1

Last Useable IP → 192.168.100.254

# IP addressing: introduction

cidr.xyz

## CIDR.xyz AN INTERACTIVE IP ADDRESS AND CIDR RANGE VISUALIZER

CIDR is a notation for describing blocks of IP addresses and is used heavily in various networking configurations. IP addresses contain 4 octets, each consisting of 8 bits giving values between 0 and 255. The decimal value that comes after the slash is the number of bits consisting of the routing prefix. This in turn can be translated into a netmask, and also designates how many available addresses are in the block.

200 . 23 . 16 . 1 / 28



255.255.255.240 NETMASK	200.23.16.0 CIDR BASE IP	200.23.16.15 BROADCAST IP	16 COUNT
200.23.16.1 FIRST USABLE IP		200.23.16.14 LAST USABLE IP	

\* For routing mask values <= 30, first and last IPs are base and broadcast addresses and are unusable.

Created by [Yuval Adam](#). Source available on [Github](#).

<https://cidr.xyz>

# IP addresses: how to get one?

That's actually **two** questions:

1. Q: How does a *host* get IP address within its network (host part of address)?
2. Q: How does a *network* get IP address for itself (network part of address)

How does *host* get IP address?

- hard-coded by sysadmin in config file (e.g., /etc/rc.config in UNIX)
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from server
  - “plug-and-play”

# DHCP: Dynamic Host Configuration Protocol

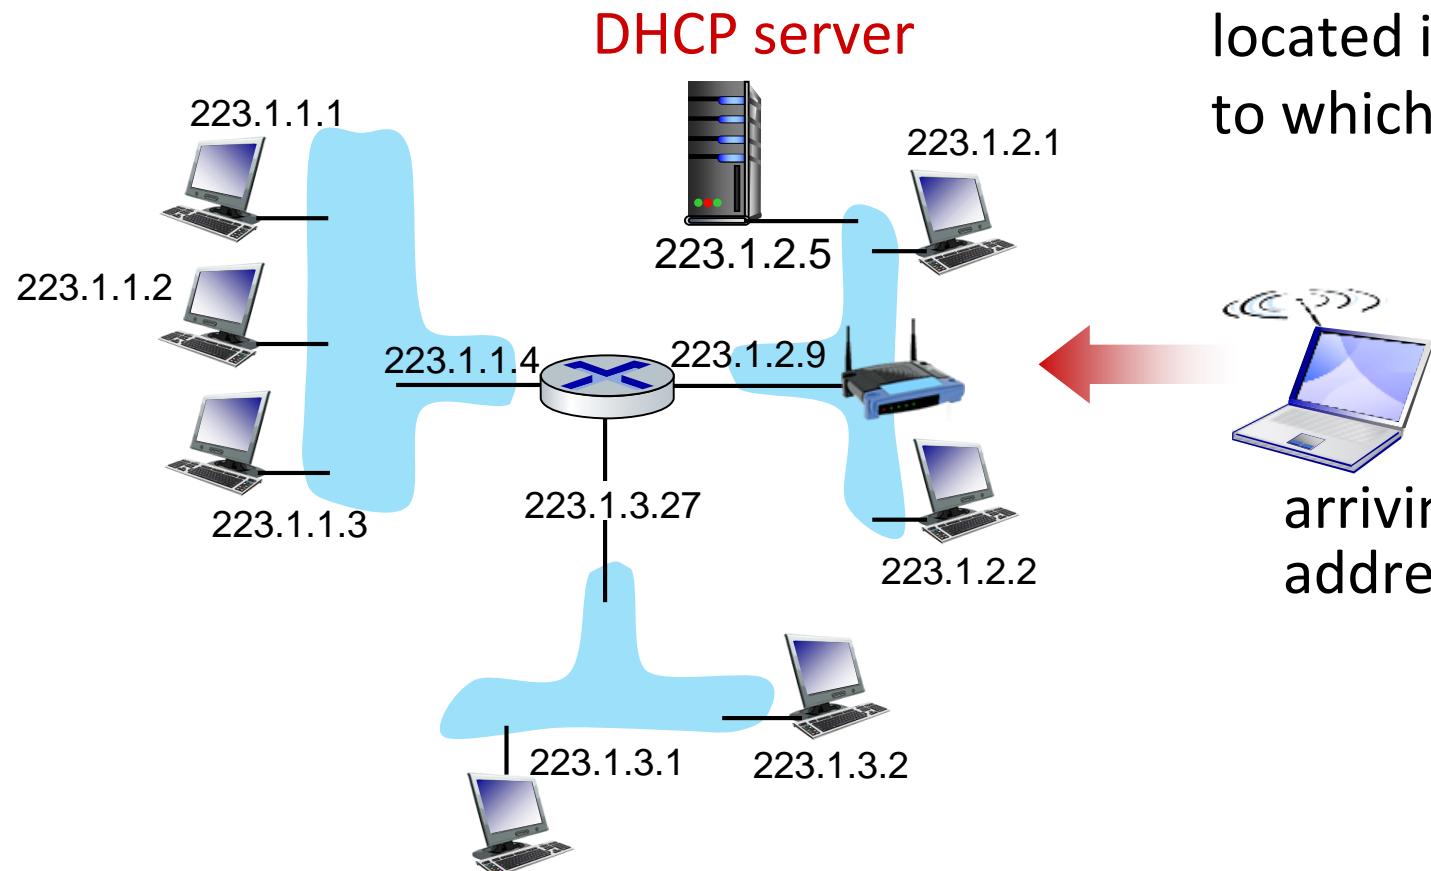
**goal:** host *dynamically* obtains IP address from network server when it “joins” network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/on)
- support for mobile users who join/leave network

## DHCP overview:

- host broadcasts **DHCP discover** msg [optional]
- DHCP server responds with **DHCP offer** msg [optional]
- host requests IP address: **DHCP request** msg
- DHCP server sends address: **DHCP ack** msg

# DHCP client-server scenario



Typically, DHCP server will be co-located in router, serving all subnets to which router is attached

arriving **DHCP client** needs address in this network

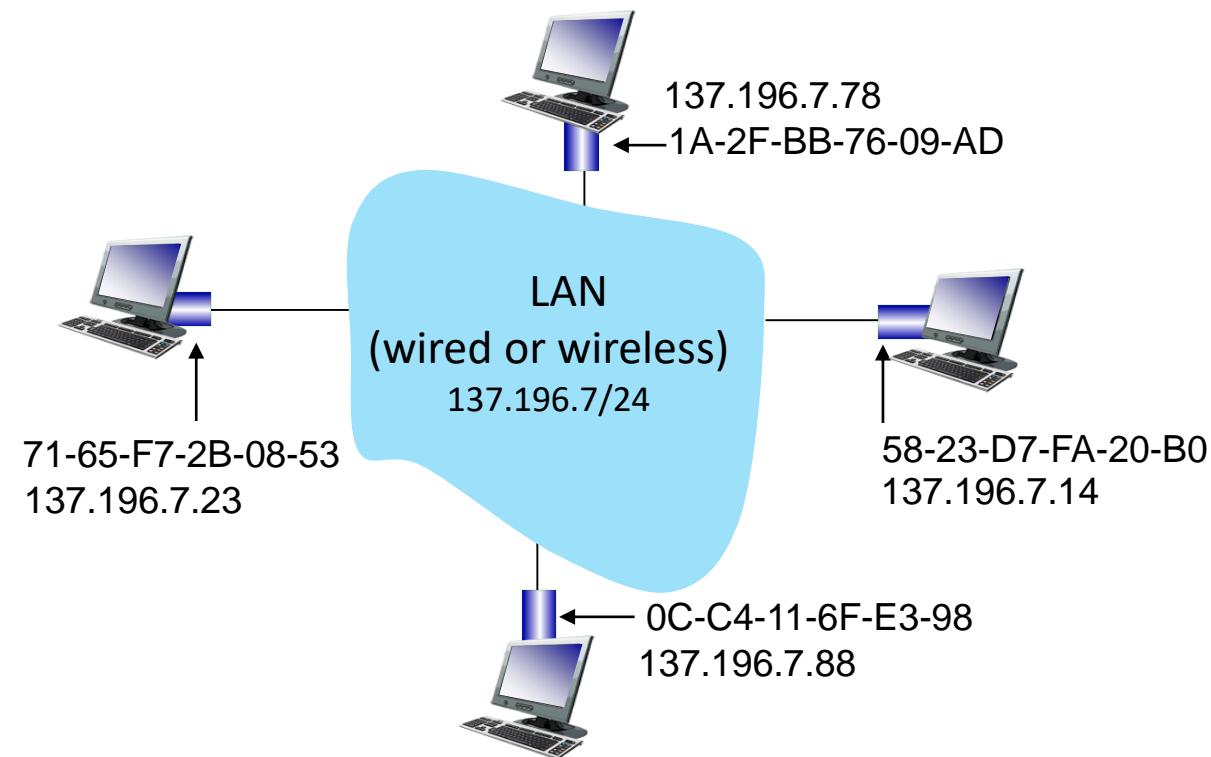
# MAC addresses

- 32-bit IP address:
    - *network-layer* address for interface
    - used for layer 3 (network layer) forwarding
    - e.g.: 128.119.40.136
  - MAC (or LAN or physical or Ethernet) address:
    - function: used “locally” to get frame from one interface to another physically-connected interface (same subnet, in IP-addressing sense)
    - 48-bit MAC address (for most LANs) burned in NIC ROM, also sometimes software settable
    - e.g.: 1A-2F-BB-76-09-AD
- hexadecimal (base 16) notation  
(each “numeral” represents 4 bits)*

# MAC addresses

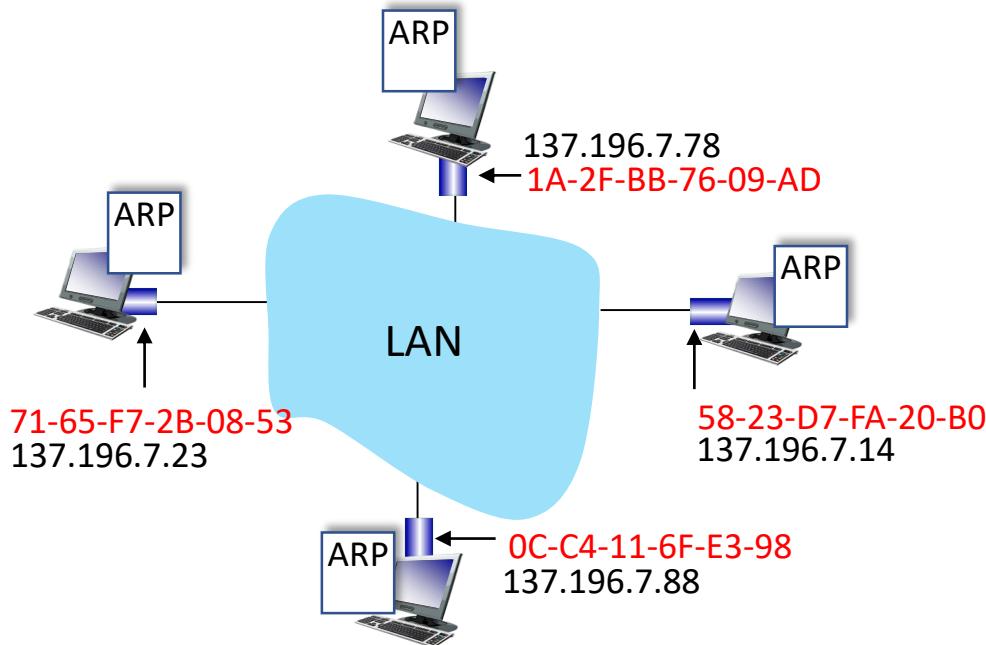
each interface on LAN

- has unique 48-bit **MAC** address
- has a locally unique 32-bit IP address (as we've seen)



# ARP: address resolution protocol

*Question:* how to determine interface's MAC address, knowing its IP address?



**ARP table:** each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:  
<IP address; MAC address; TTL>
- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)

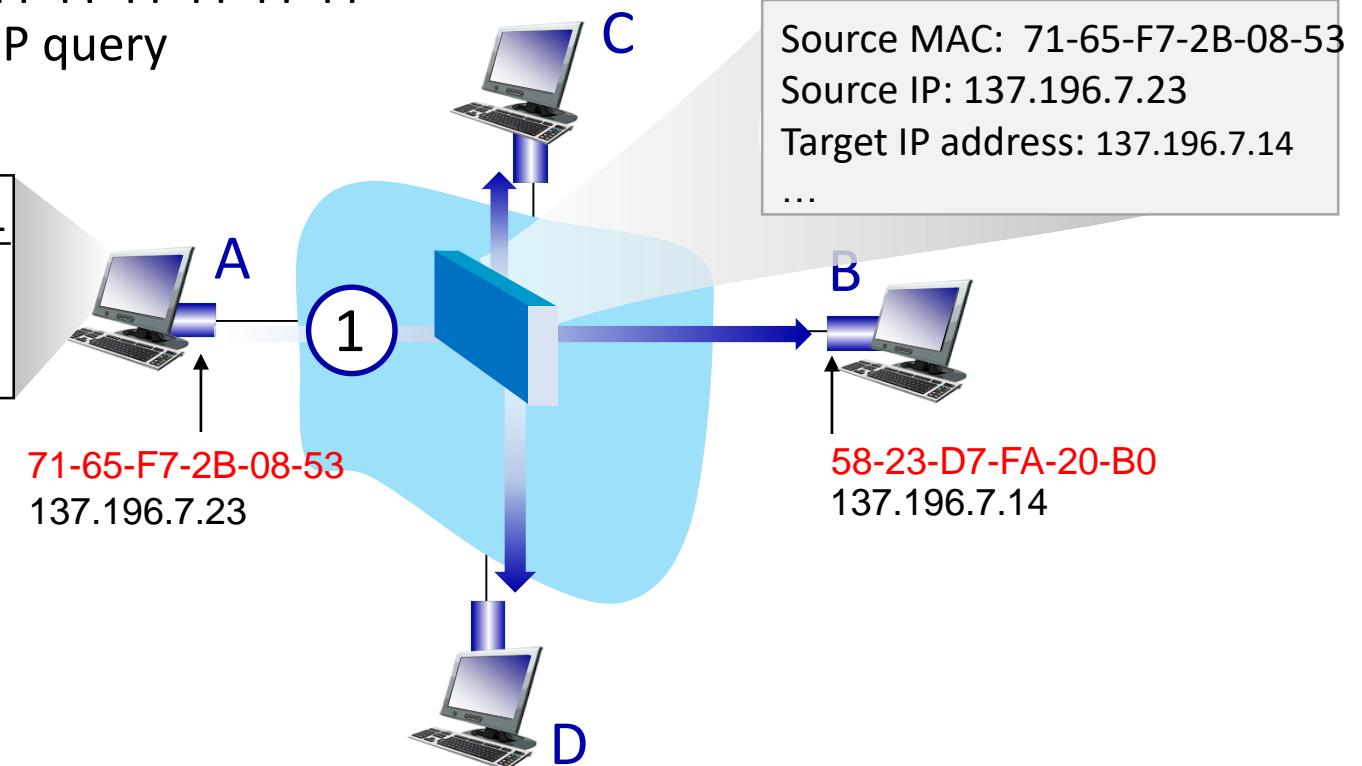
# ARP protocol in action

example: A wants to send datagram to B

- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address

- 1 A broadcasts ARP query, containing B's IP addr
- destination MAC address = FF-FF-FF-FF-FF-FF
  - all nodes on LAN receive ARP query

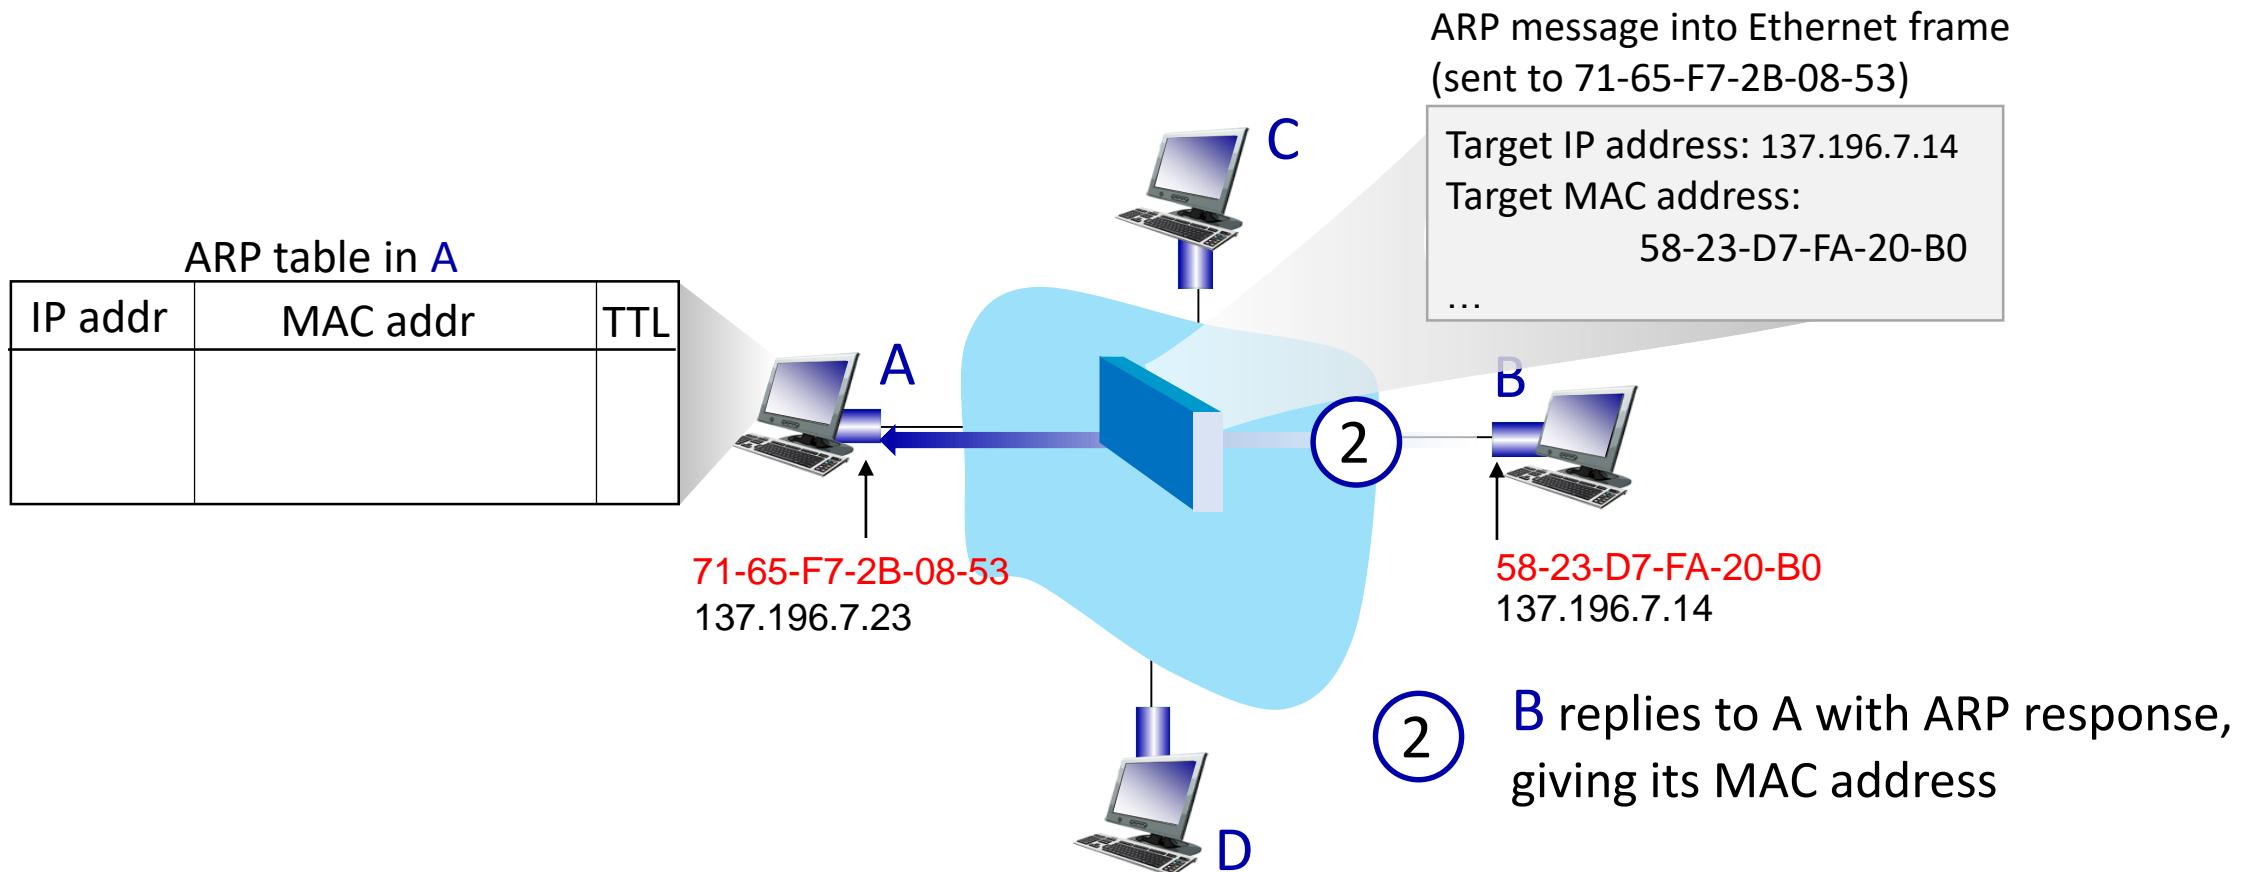
ARP table in A		
IP addr	MAC addr	TTL



# ARP protocol in action

example: A wants to send datagram to B

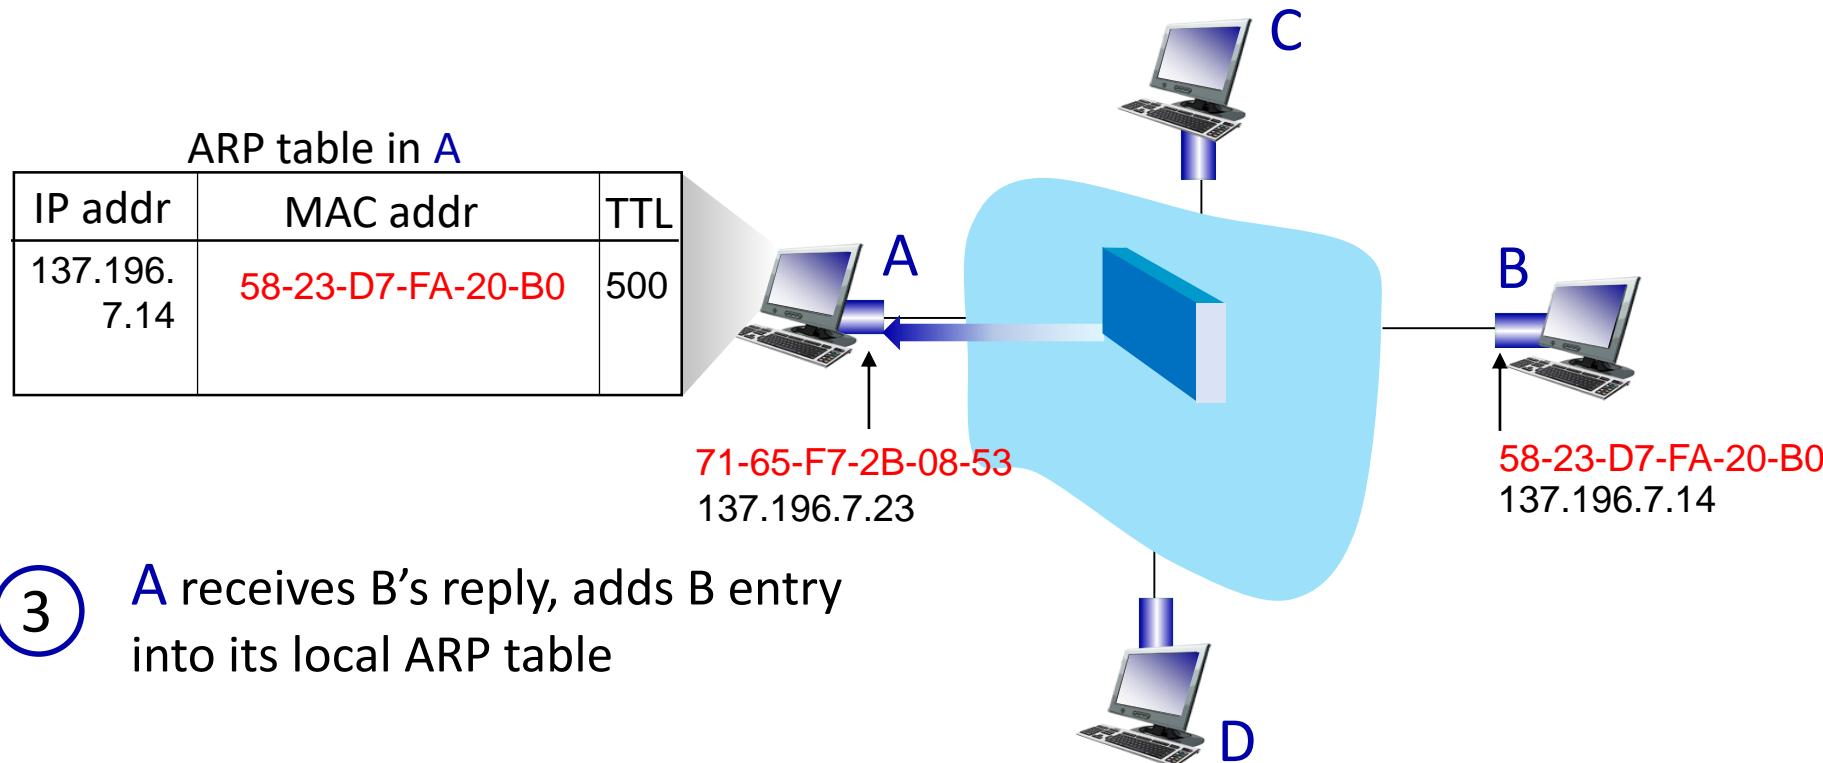
- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address



# ARP protocol in action

example: A wants to send datagram to B

- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address

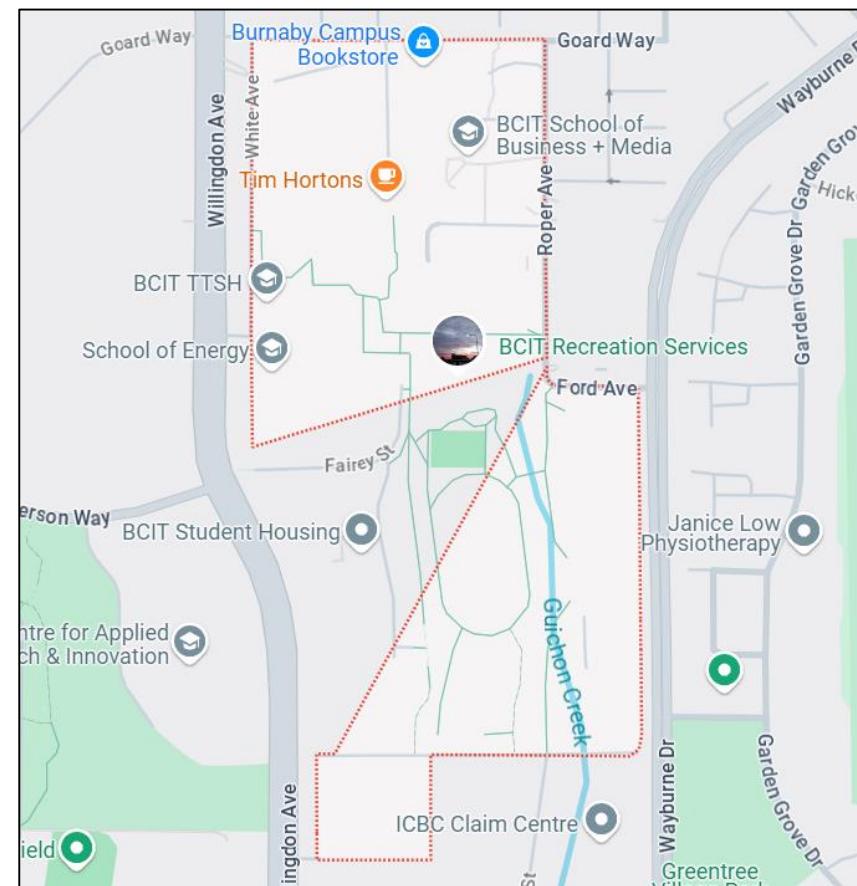


# IP address vs MAC address

House/Unit/Apt Number → MAC address



BCIT Postal Code: V5G 3H2 → IP address



Component	Postal System	Computer Network
MAC Address	Home/Mailbox Number	Unique Hardware Address (e.g., 00:1A:2B:3C:4D:5E)
IP Address	Postal Code (e.g., V3N 4Y7)	Logical Network Address (e.g., 192.168.1.10)

**First, check the IP address and MAC address of your Network Interface Card (NIC)?**

*Windows: ipconfig /all*

*Linux/macOS: ifconfig*

**Then, visit the following website to find out who is the vendor of your NIC?**

<https://macaddress.io/>

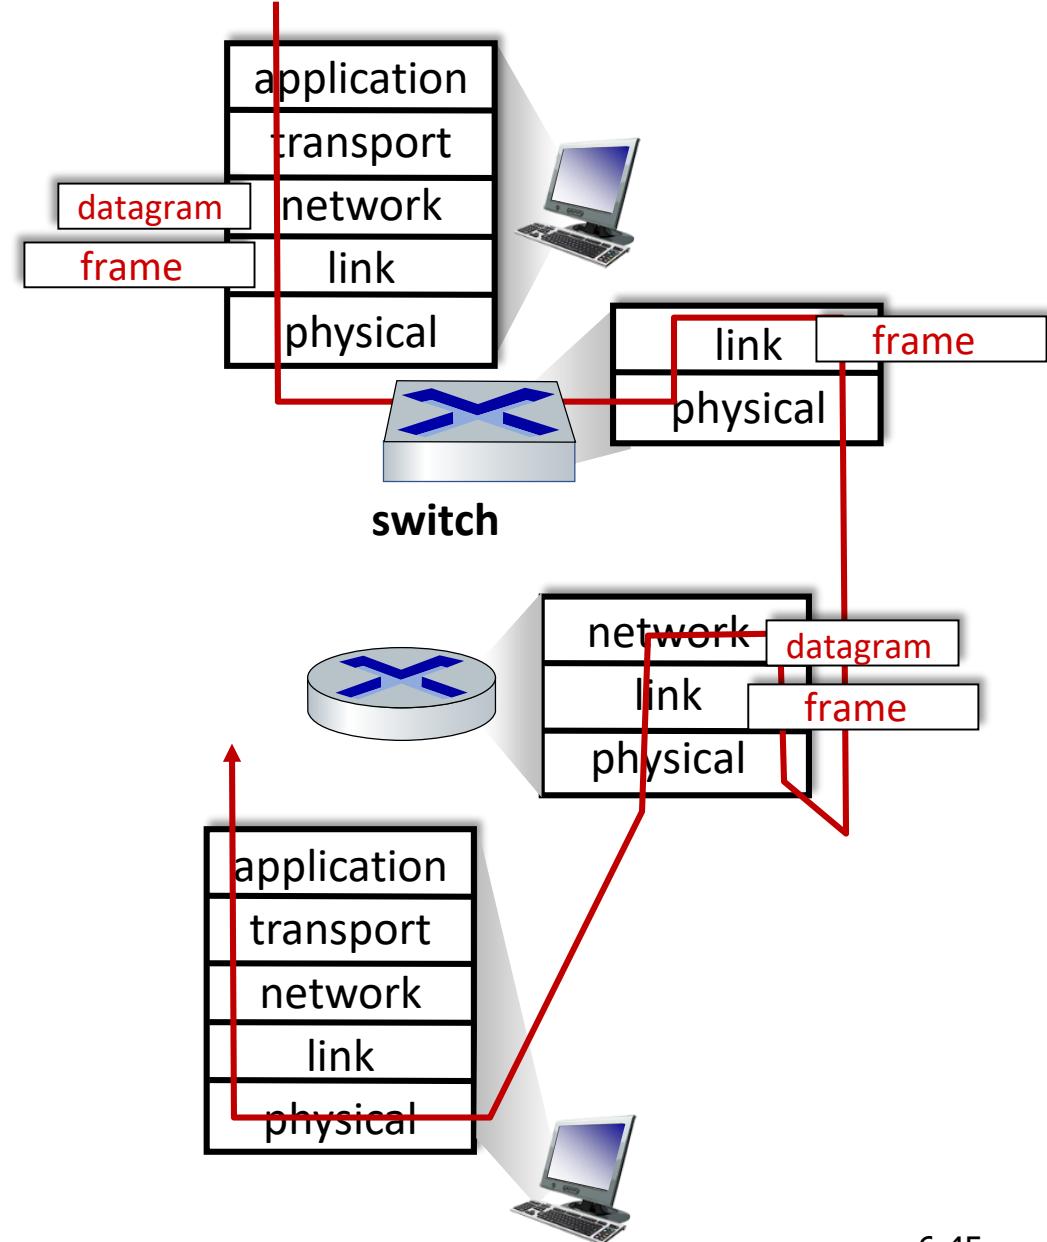
# Switches vs. routers

both are store-and-forward:

- *routers*: network-layer devices (examine network-layer headers)
- *switches*: link-layer devices (examine link-layer headers)

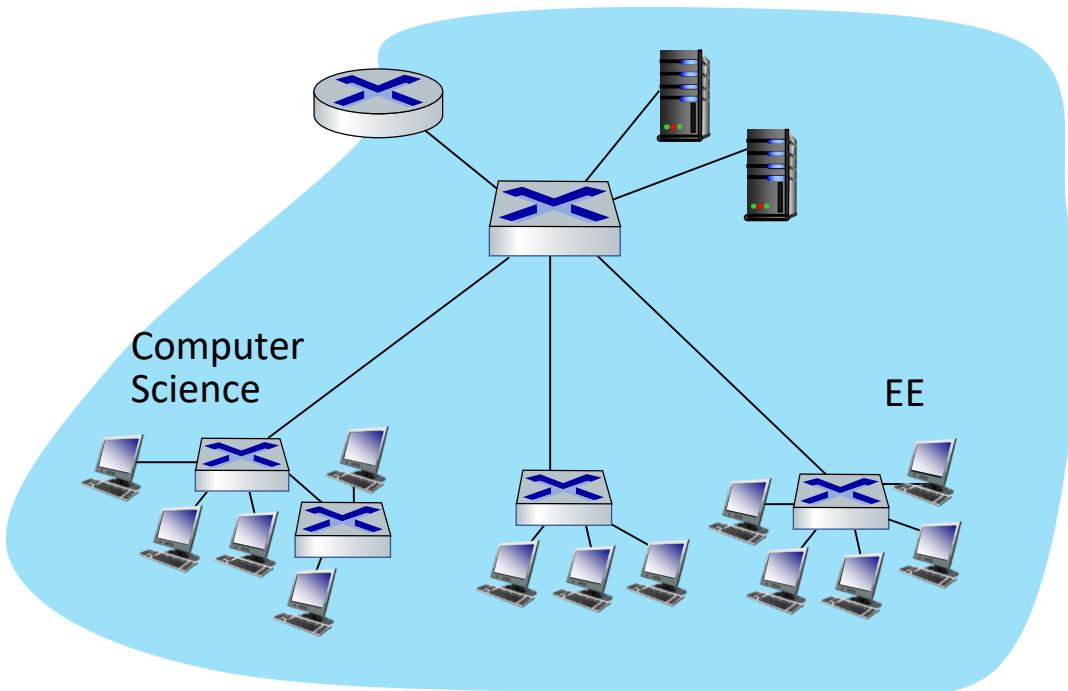
both have forwarding tables:

- *routers*: compute tables using routing algorithms, IP addresses
- *switches*: learn forwarding table using flooding, learning, MAC addresses



# Virtual LANs (VLANs): motivation

*Q:* what happens as LAN sizes scale, users change point of attachment?

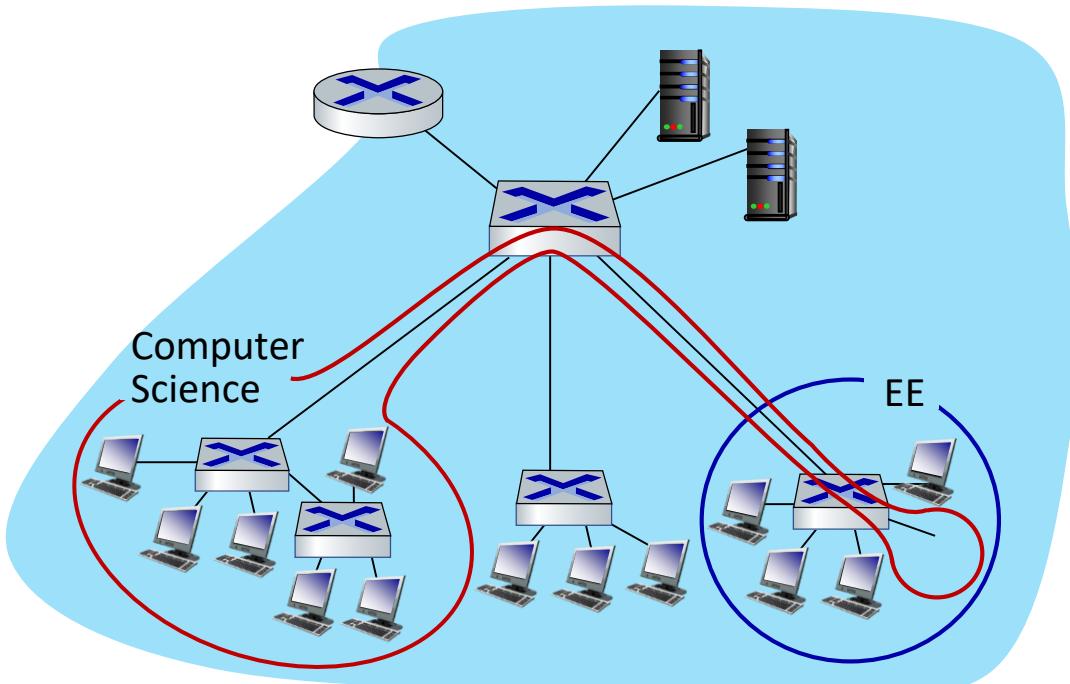


single broadcast domain:

- *scaling:* all layer-2 broadcast traffic (ARP, DHCP, unknown MAC) must cross entire LAN
- efficiency, security, privacy issues

# Virtual LANs (VLANs): motivation

Q: what happens as LAN sizes scale, users change point of attachment?



single broadcast domain:

- *scaling*: all layer-2 broadcast traffic (ARP, DHCP, unknown MAC) must cross entire LAN
- efficiency, security, privacy, efficiency issues

administrative issues:

- CS user moves office to EE - *physically* attached to EE switch, but wants to remain *logically* attached to CS switch

# IPv6: motivation

- **initial motivation:** 32-bit IPv4 address space would be completely allocated
- additional motivation:
  - speed processing/forwarding: 40-byte fixed length header
  - enable different network-layer treatment of “flows”

IPv4

**32-bit address space**

=

**4,294,967,296 ( $2^{32}$ )  
unique addresses**

IPv6

**128-bit address space**

=

**340,282,366,920,938,463,463,374,  
607,431,768,211,456 ( $2^{128}$ )  
unique addresses**

example:  
**103.81.45.28**

example:  
**2001:0DB8:0000:2F3B:02AA:00FF:FE28:9C5A**

# IPv4 vs. IPv6

## Address Space:

- **IPv4:** Limited, leading to address exhaustion.
- **IPv6:** Vast, providing a nearly unlimited number of addresses.

## Configuration:

- **IPv4:** Manual or DHCP.
- **IPv6:** Stateless Address Auto-configuration (SLAAC) and DHCPv6.

## Security:

- **IPv4:** Optional IPSec support.
- **IPv6:** Mandatory IPSec support, integrated into the protocol.

## Network Address Translation (NAT):

- **IPv4:** Commonly used due to limited address space.
- **IPv6:** Not required, supporting end-to-end connectivity.

## Header Complexity:

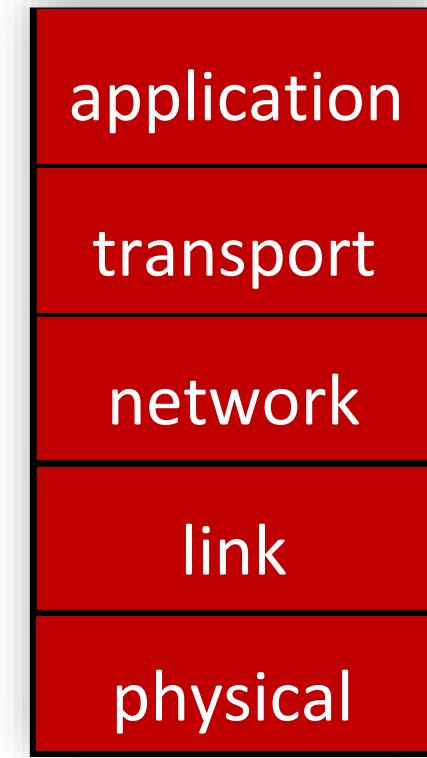
- **IPv4:** More complex with various optional fields.
- **IPv6:** Simplified header structure for more efficient processing.

## Performance:

- **IPv4:** Can be slower due to header complexity and NAT overhead.
- **IPv6:** Potentially faster due to simplified headers and elimination of NAT.

# Layered Internet protocol stack

- *application*: supporting network applications
  - HTTP, IMAP, SMTP, DNS
- *transport*: process-process data transfer
  - TCP, UDP
- *network*: routing of datagrams from source to destination
  - IP, routing protocols
- *link*: data transfer between neighboring network elements
  - Ethernet, 802.11 (WiFi), PPP
- *physical*: bits “on the wire”

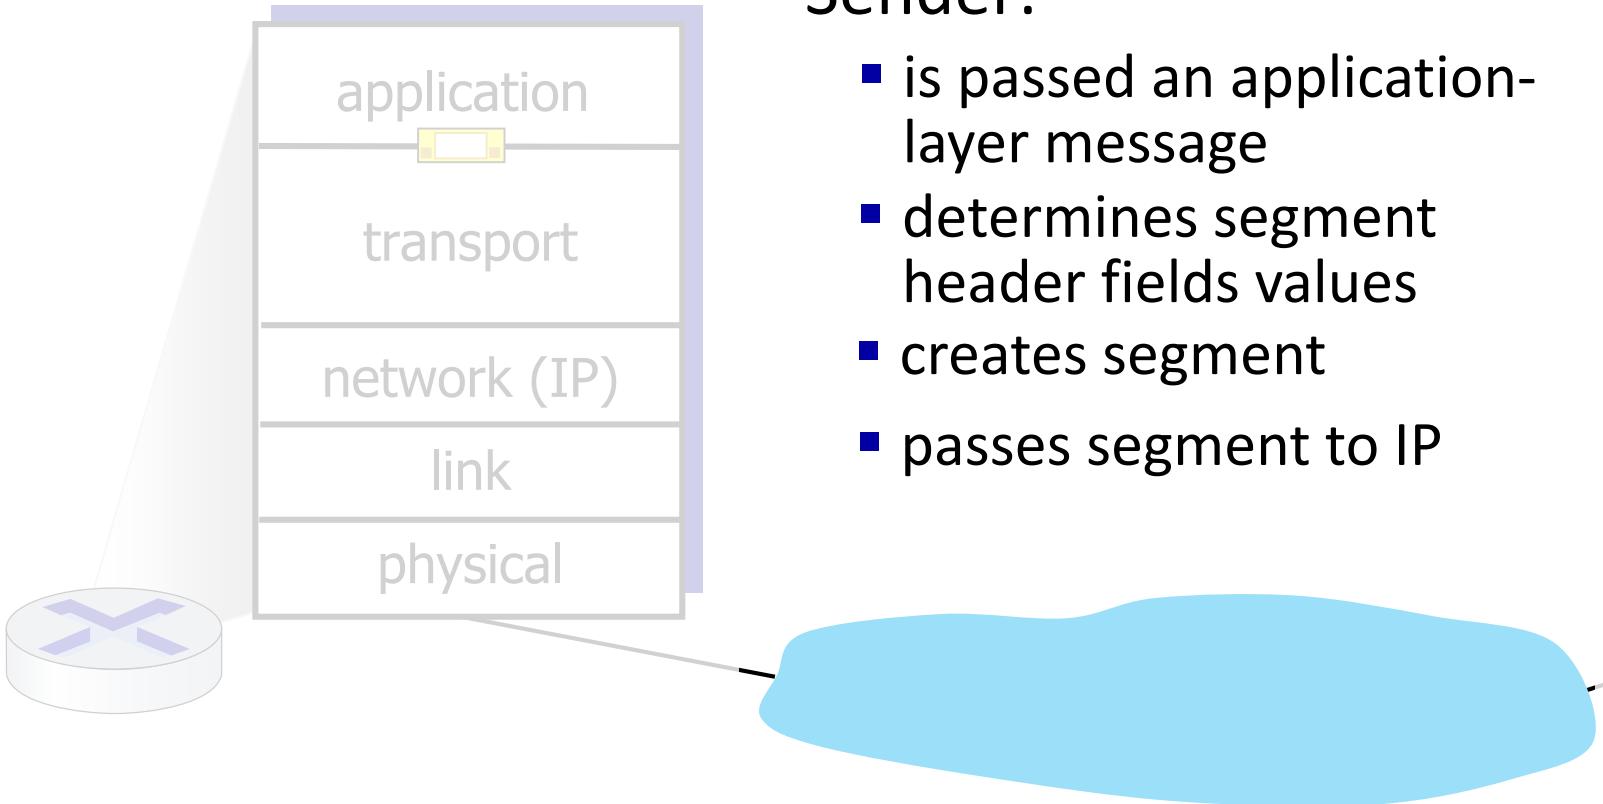


Five-Layer Internet Protocol Stack

# Transport Layer Actions

Sender:

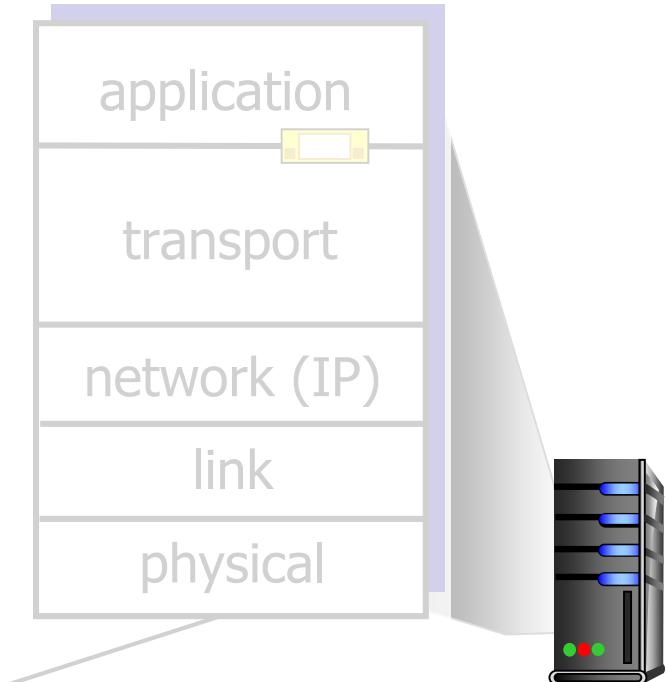
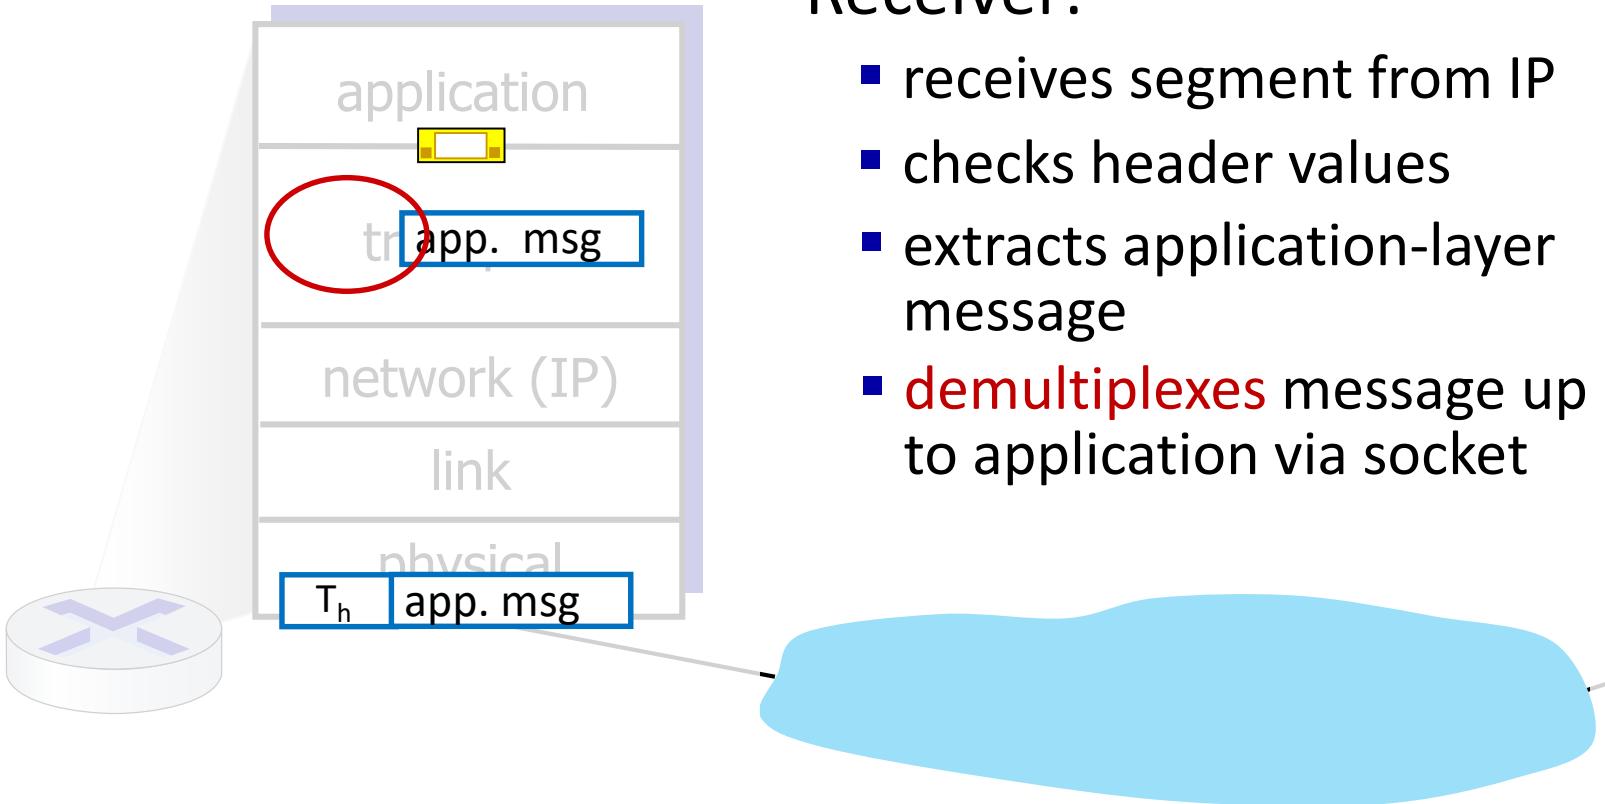
- is passed an application-layer message
- determines segment header fields values
- creates segment
- passes segment to IP



# Transport Layer Actions

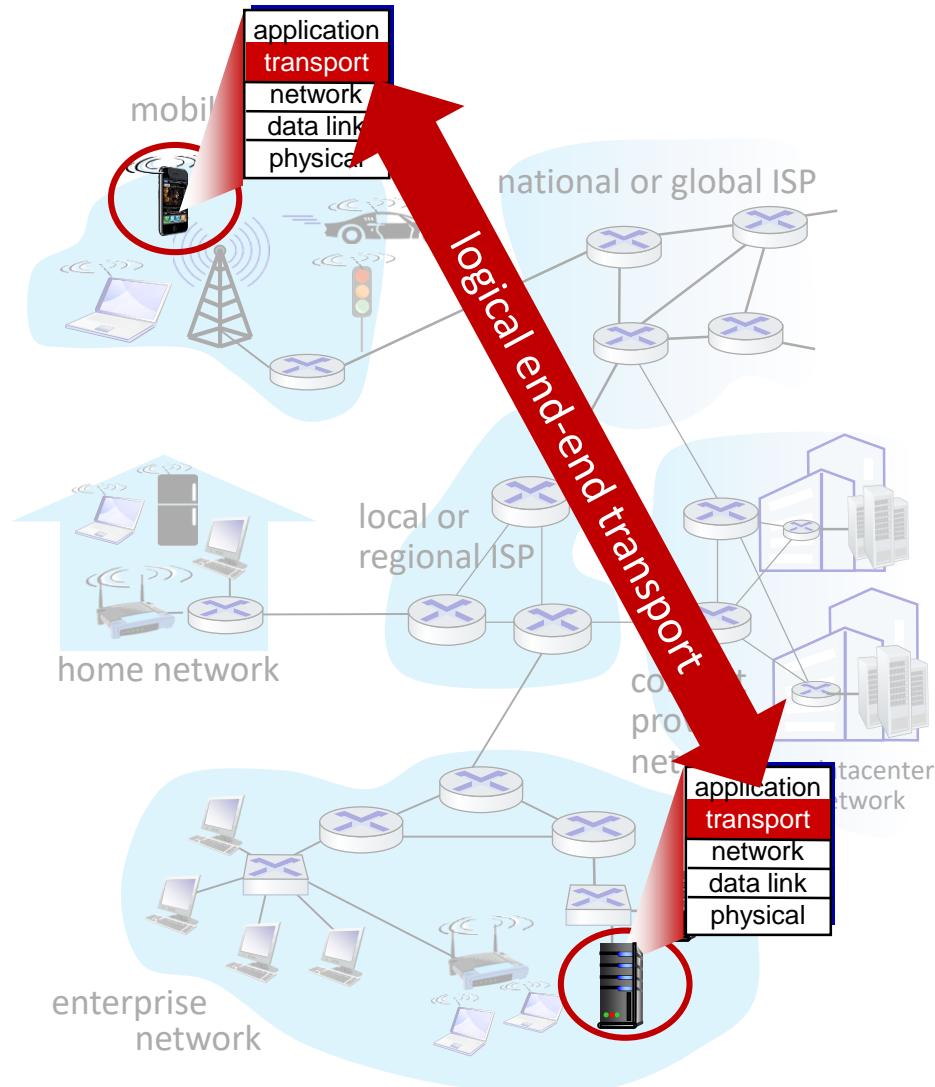
## Receiver:

- receives segment from IP
- checks header values
- extracts application-layer message
- **demultiplexes** message up to application via socket



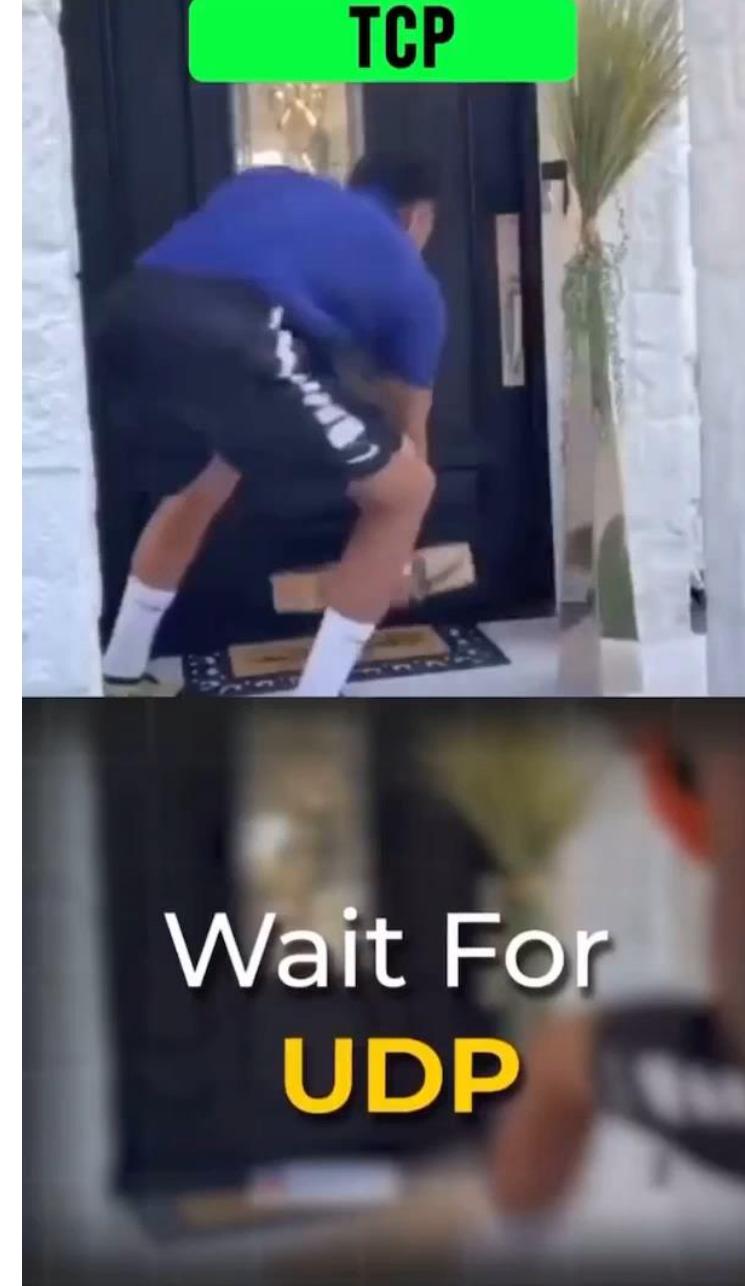
# Two principal Internet transport protocols

- **TCP:** Transmission Control Protocol
  - reliable, in-order delivery
  - congestion control
  - flow control
  - connection setup
- **UDP:** User Datagram Protocol
  - unreliable, unordered delivery
  - no-frills extension of “best-effort” IP
- services *not* available:
  - delay guarantees
  - bandwidth guarantees



# TCP vs UDP ;)

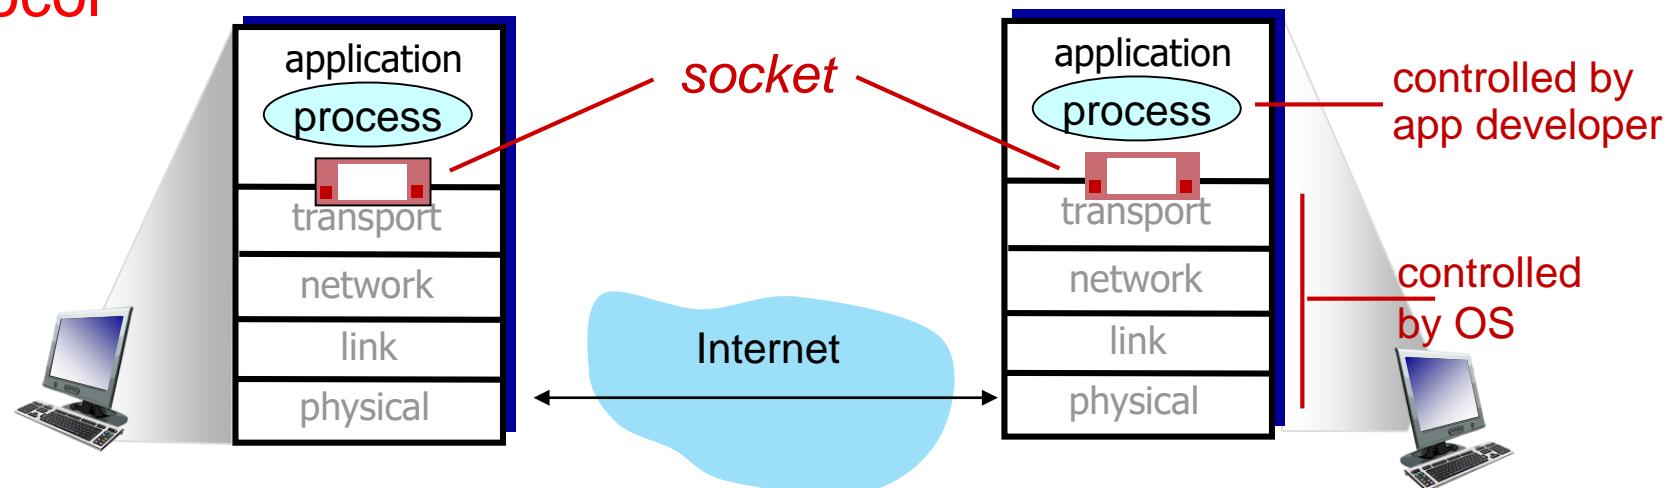
- **TCP:** Transmission Control Protocol
  - reliable, in-order delivery
  - congestion control
  - flow control
  - connection setup
- **UDP:** User Datagram Protocol
  - unreliable, unordered delivery
  - no-frills extension of “best-effort” IP
- services *not* available:
  - delay guarantees
  - bandwidth guarantees



# Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process
  - two sockets involved: one on each side

*Socket is a door between application process and end-end-transport protocol*



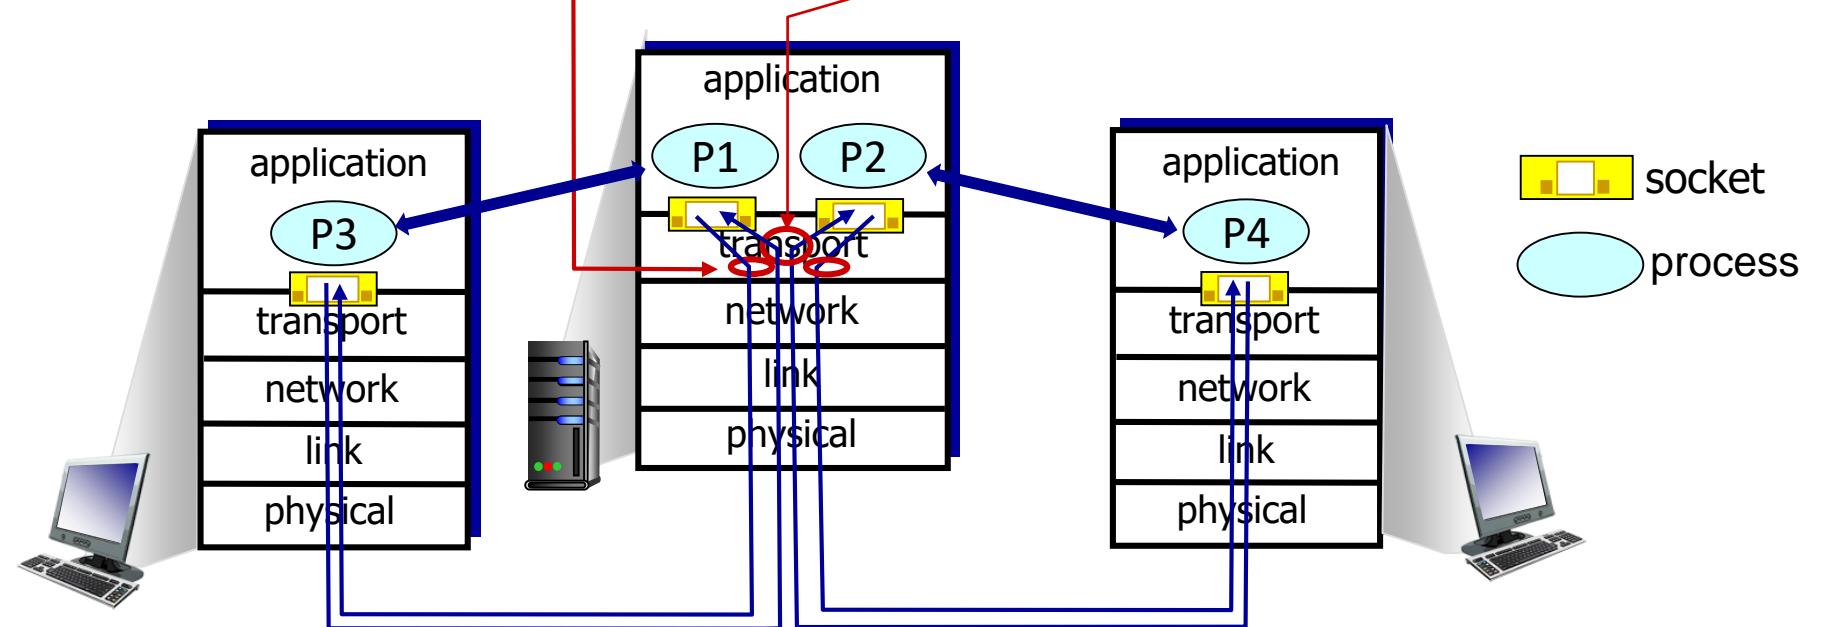
# Multiplexing/demultiplexing

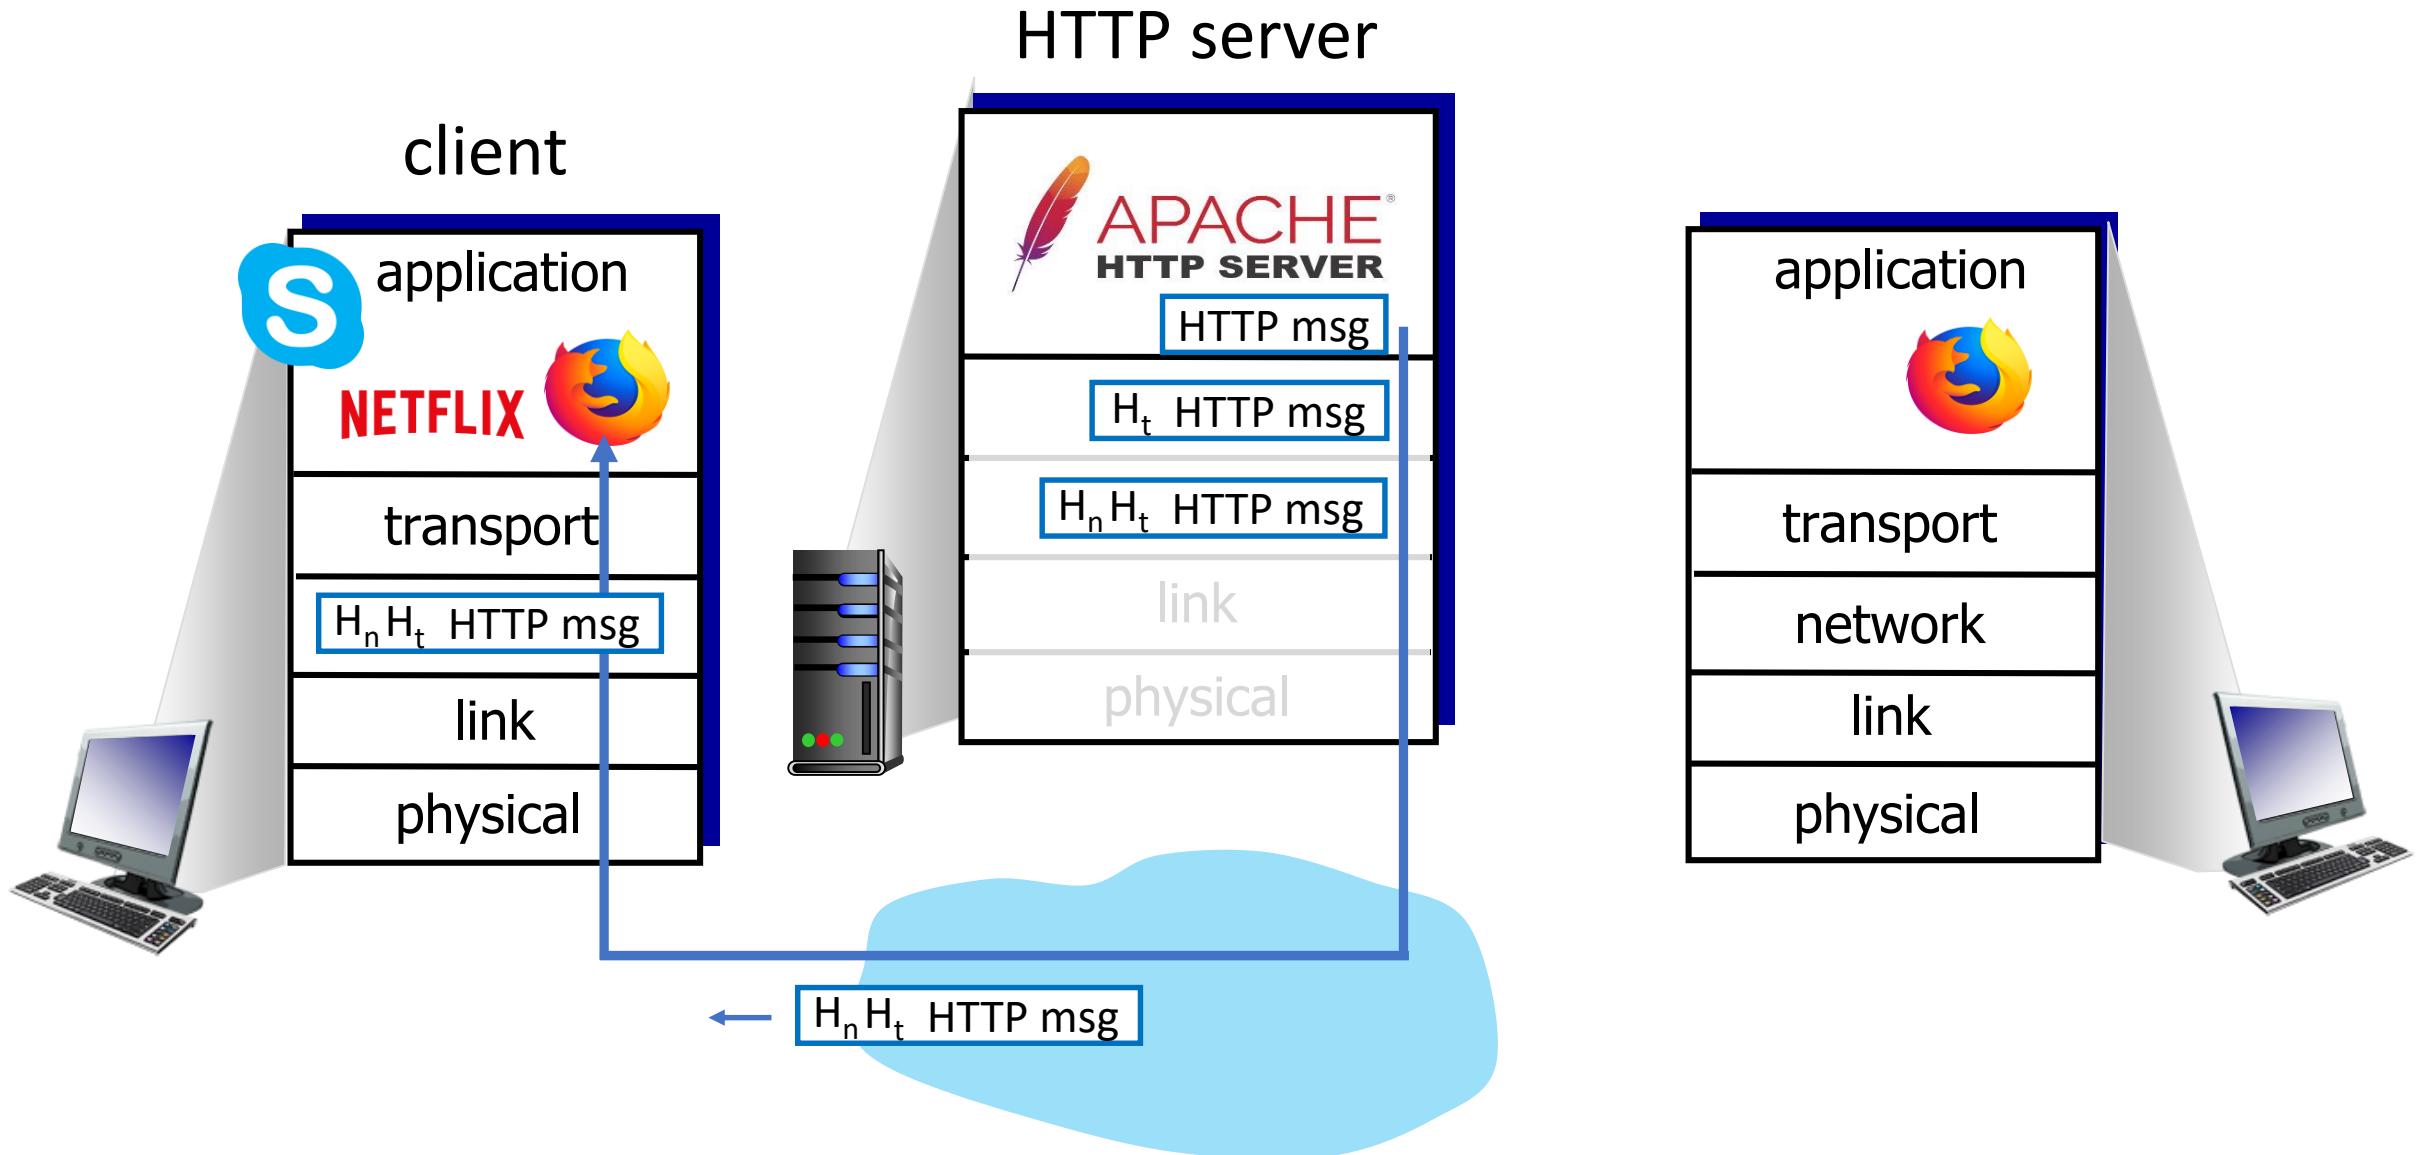
*multiplexing as sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

*demultiplexing as receiver:*

use header info to deliver received segments to correct socket





# TCP: Three-way handshake

The TCP three-way handshake is a process used to establish a connection between a client and a server:

## 1. SYN (Synchronize):

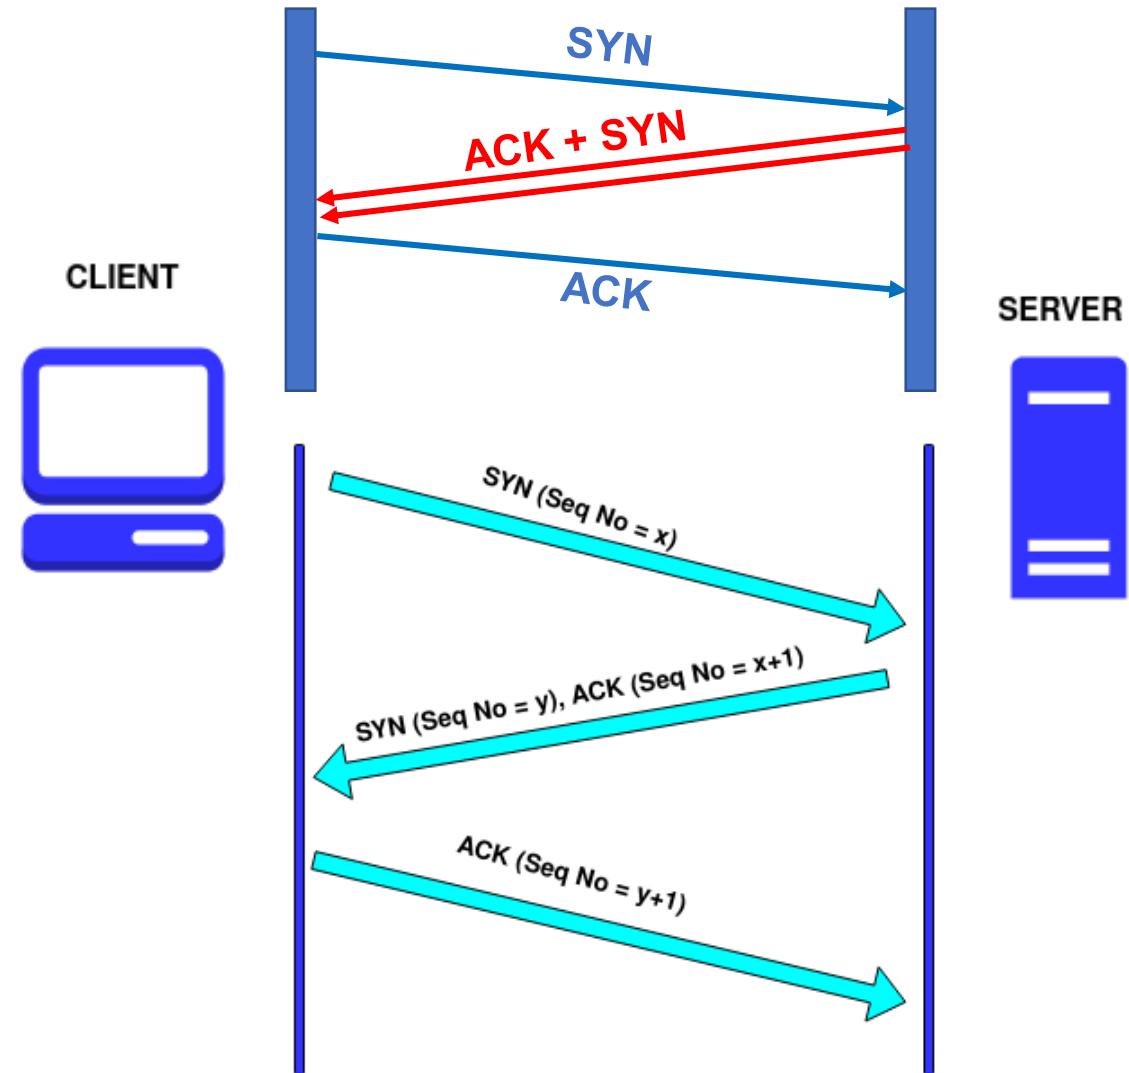
The client sends a SYN packet to the server to initiate a connection.

## 2. SYN-ACK (Synchronize-Acknowledge):

The server responds with a SYN-ACK packet to acknowledge the client's request and indicate it is ready to establish a connection.

## 3. ACK (Acknowledge):

The client sends an ACK packet back to the server, confirming the connection is established



# UDP: User Datagram Protocol

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- *connectionless*:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

## Why is there a UDP?

- no connection establishment (which can add RTT delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control
  - UDP can blast away as fast as desired!
  - can function in the face of congestion

# UDP: User Datagram Protocol



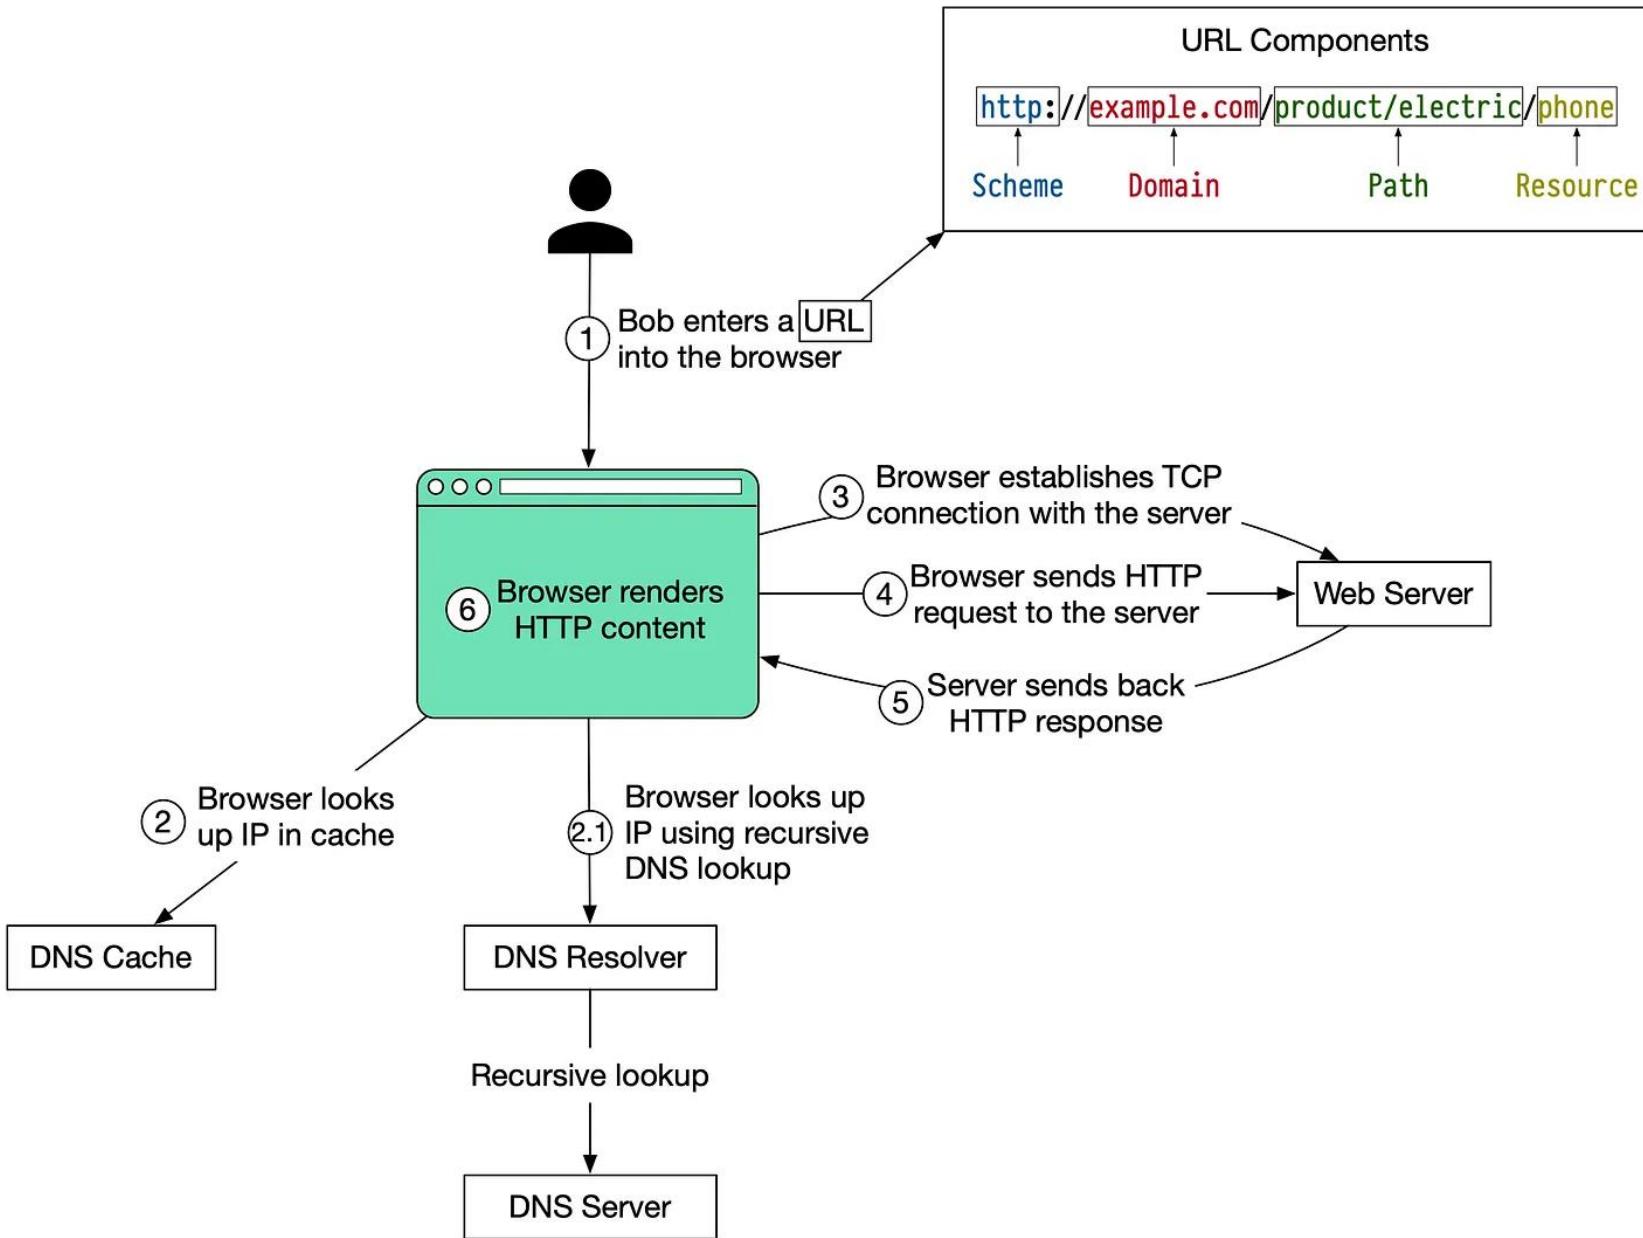
- UDP use:
  - streaming multimedia apps (loss tolerant, rate sensitive)
  - DHCP, DNS, NTP, TFTP, SNMP, SIP (VoIP), ...
- if reliable transfer needed over UDP (e.g., HTTP/3):
  - add needed reliability at application layer
  - add congestion control at application layer

# A day in the life of a web request

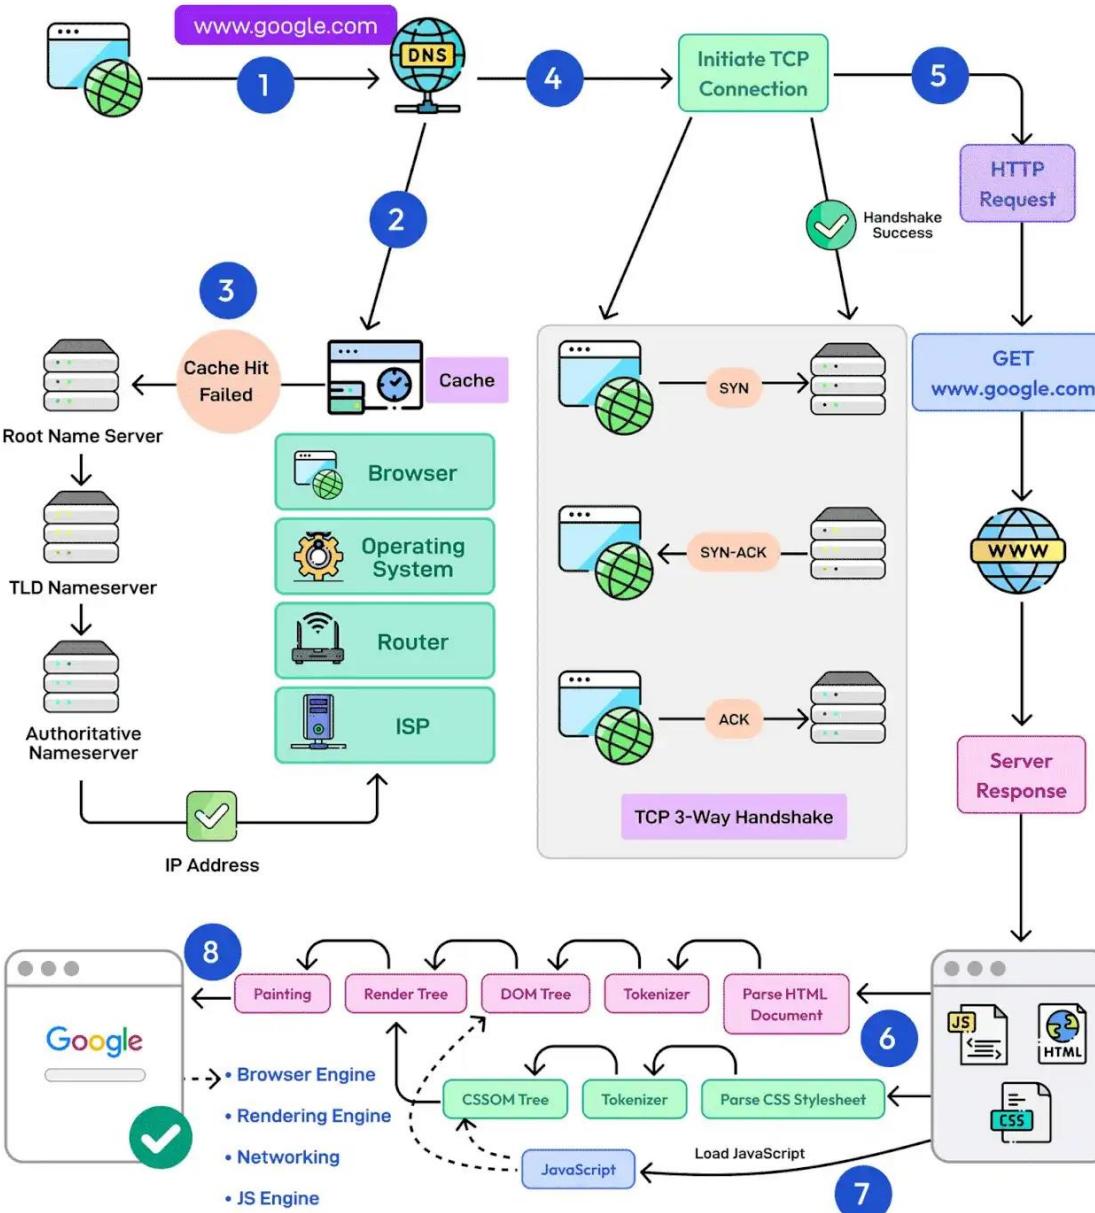
- *goal*: identify, review, understand protocols (at all layers) involved in seemingly simple scenario: requesting www page
- *scenario*: student attaches laptop to campus network, requests/receives www.google.com

THERE IS NOT A SINGLE AND SIMPLE ANSWER TO THIS QUESTION





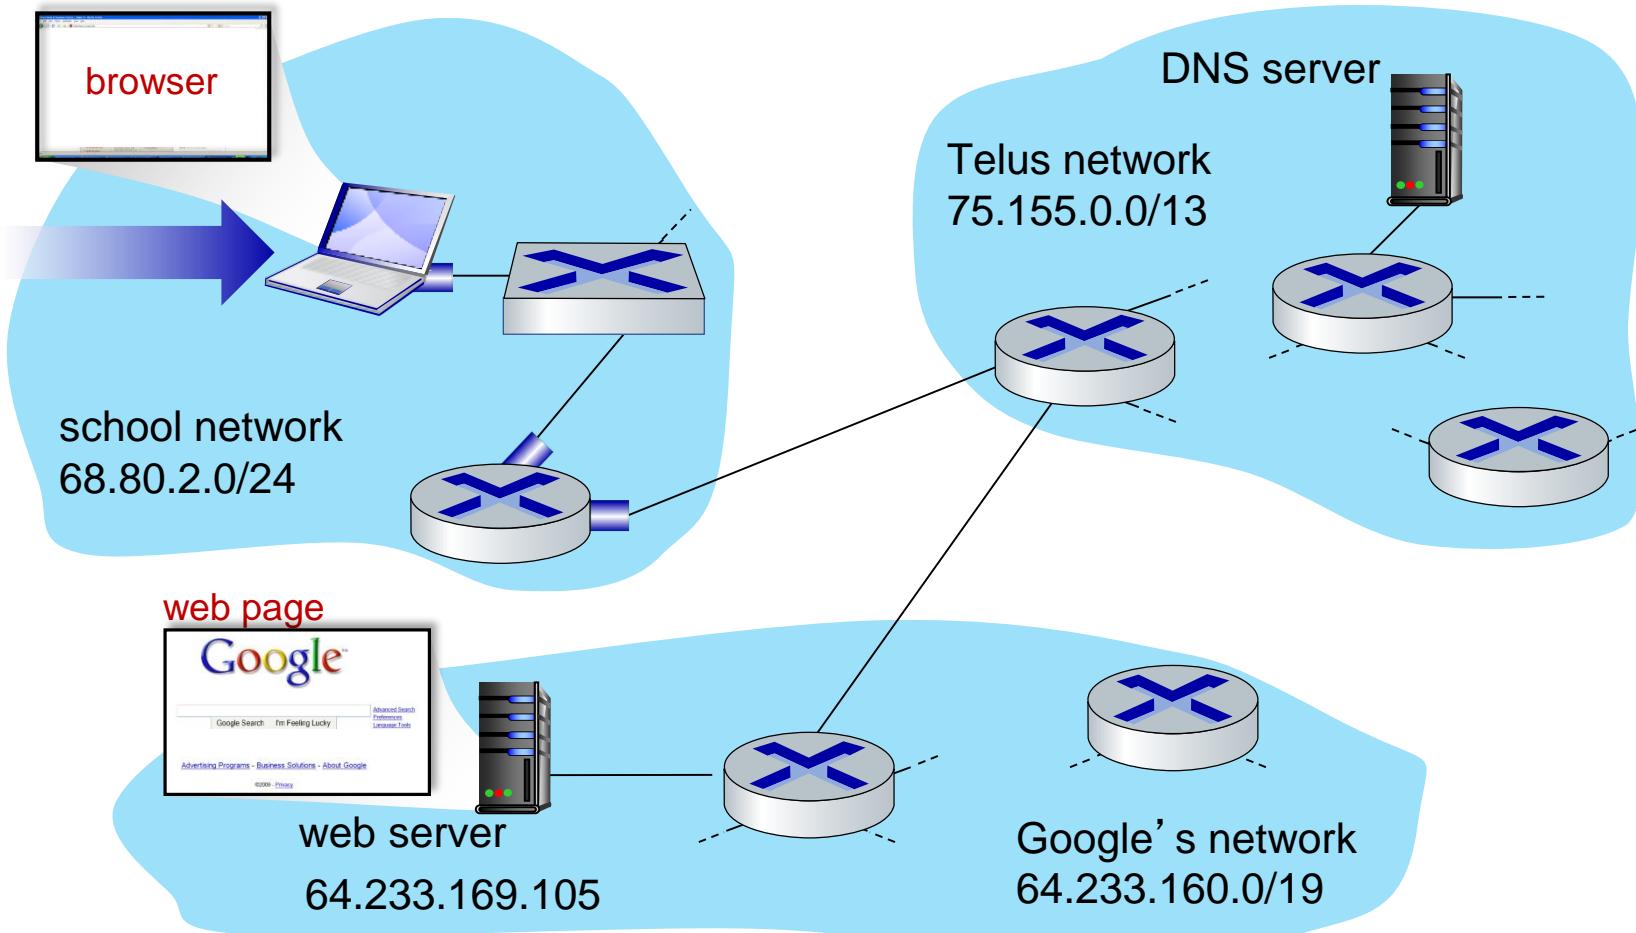
# What Happens When You Type Google.com in Your Browser?



1. First up, you type the website address in the browser's address bar.
2. The browser checks its cache first. If there's a cache miss, it must find the IP address.
3. DNS lookup begins (think of it as looking up a phone number). The request goes through different DNS servers (root, TLD, and authoritative). Finally, the IP address is retrieved.
4. Next, your browser initiates a TCP connection like a handshake. For example, in the case of HTTP 1.1, the client and server perform a TCP 3-way handshake with SYN, SYN-ACK, and ACK messages.
5. Once the handshake is successful, the browser makes an HTTP request to the server and the server responds with HTML, CSS, and JS files.
6. Finally, the browser processes everything. It parses the HTML document and creates DOM and CSSOM trees.
7. The browser executes the JavaScript and renders the page through various steps (tokenizer, parser, render tree, layout, and painting).
8. Finally, the webpage appears on your screen.

<https://blog.bytebybytego.com/p/ep145-infrastructure-as-code-landscape>

# A day in the life: scenario

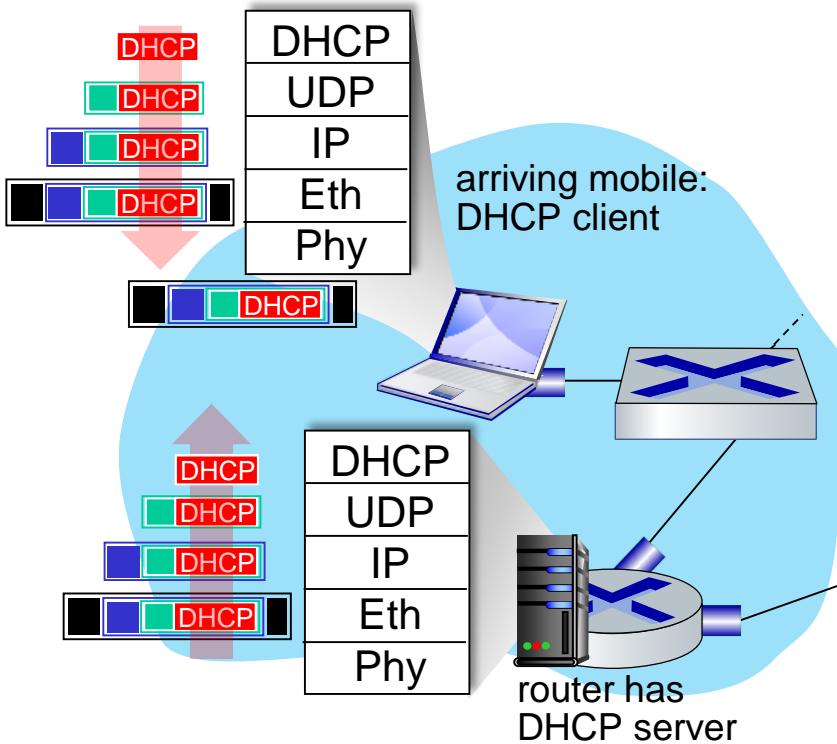


scenario:

- arriving mobile client attaches to network ...
- requests web page:  
[www.google.com](http://www.google.com)

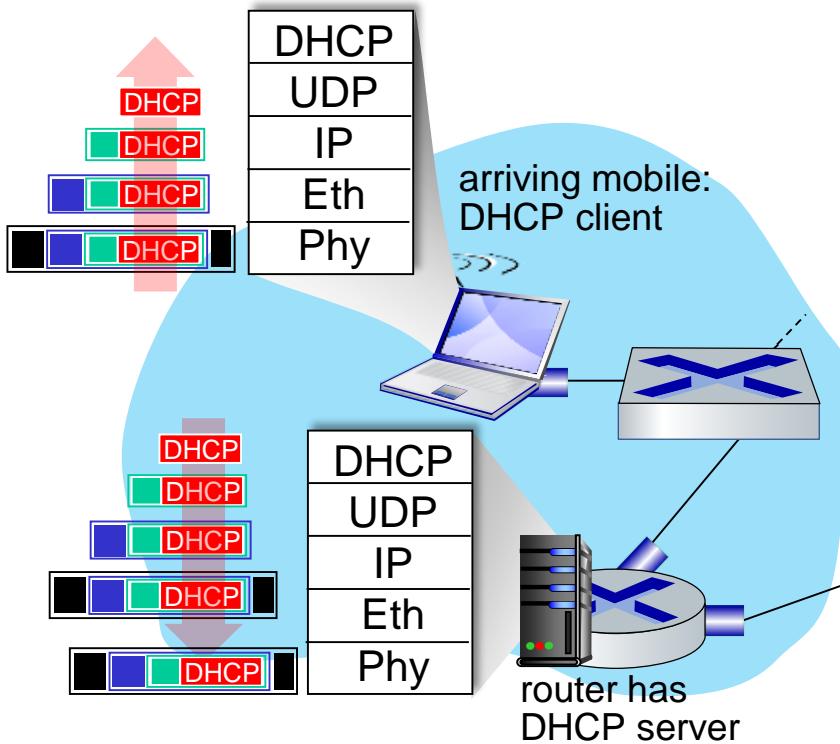
*Sounds simple!* !

# A day in the life: connecting to the Internet



- connecting laptop needs to get its own IP address, addr of first-hop router, addr of DNS server: use **DHCP**
- DHCP request **encapsulated** in **UDP**, encapsulated in **IP**, encapsulated in **802.3 Ethernet**
- Ethernet frame **broadcast** (dest: FFFFFFFFFFFF) on LAN, received at router running **DHCP** server
- Ethernet **de-muxed** to IP de-muxed, UDP de-muxed to DHCP

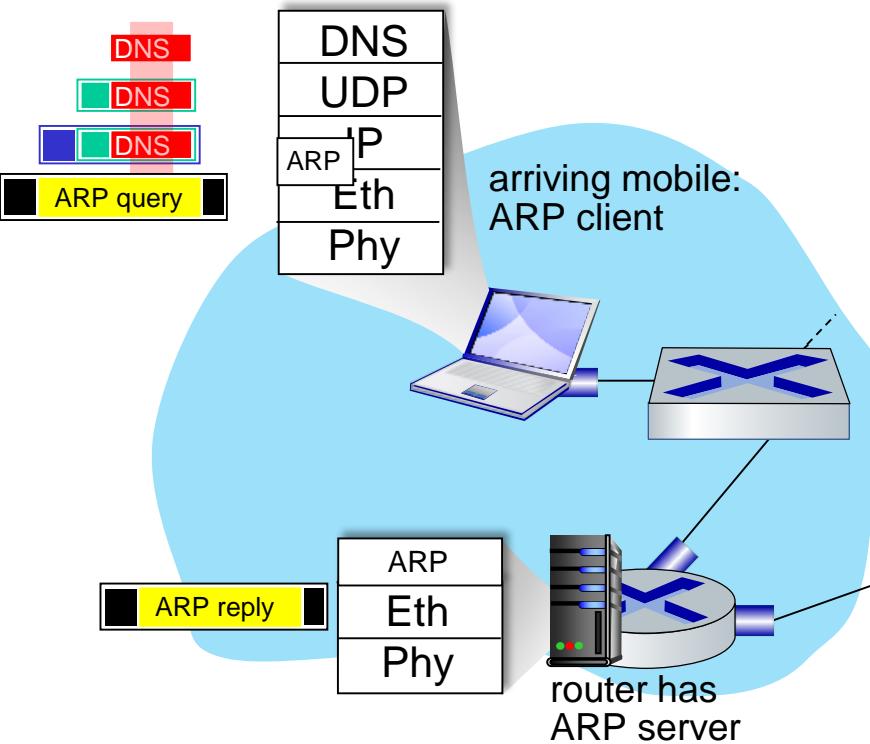
# A day in the life: connecting to the Internet



- DHCP server formulates **DHCP ACK** containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation at DHCP server, frame forwarded (**switch learning**) through LAN, demultiplexing at client
- DHCP client receives DHCP ACK reply

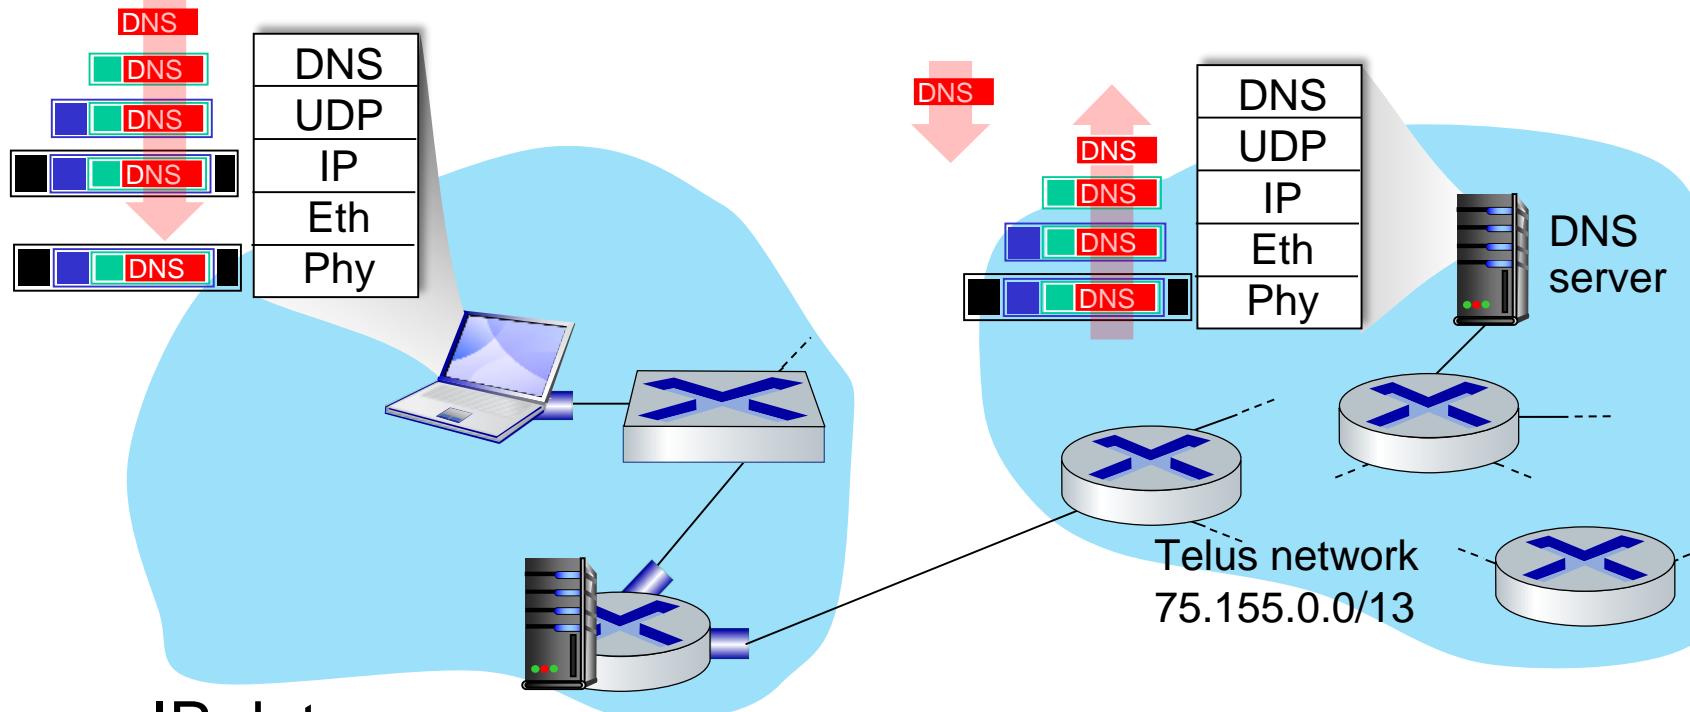
*Client now has IP address, knows name & addr of DNS server, IP address of its first-hop router*

# A day in the life... ARP (before DNS, before HTTP)



- before sending **HTTP** request, need IP address of www.google.com: **DNS**
- DNS query created, encapsulated in UDP, encapsulated in IP, encapsulated in Eth. To send frame to router, need MAC address of router interface: **ARP**
- **ARP query** broadcast, received by router, which replies with **ARP reply** giving MAC address of router interface
- client now knows MAC address of first hop router, so can now send frame containing DNS query

# A day in the life... using DNS

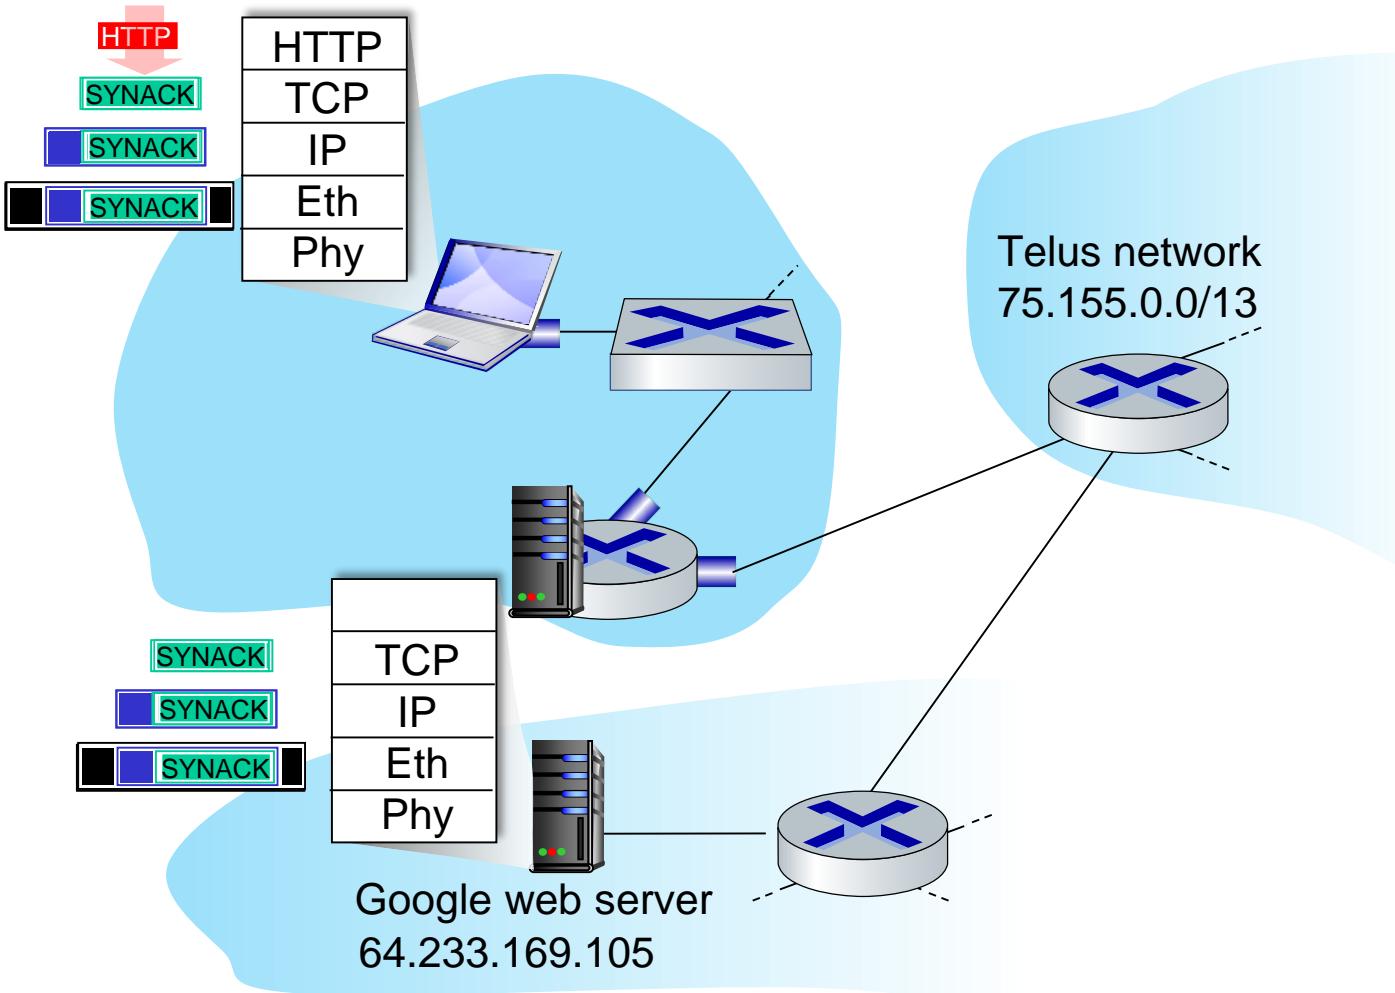


- IP datagram containing DNS query forwarded via LAN switch from client to 1<sup>st</sup> hop router

- IP datagram forwarded from campus network into Telus network, routed (tables created by **RIP, OSPF, IS-IS** and/or **BGP** routing protocols) to DNS server

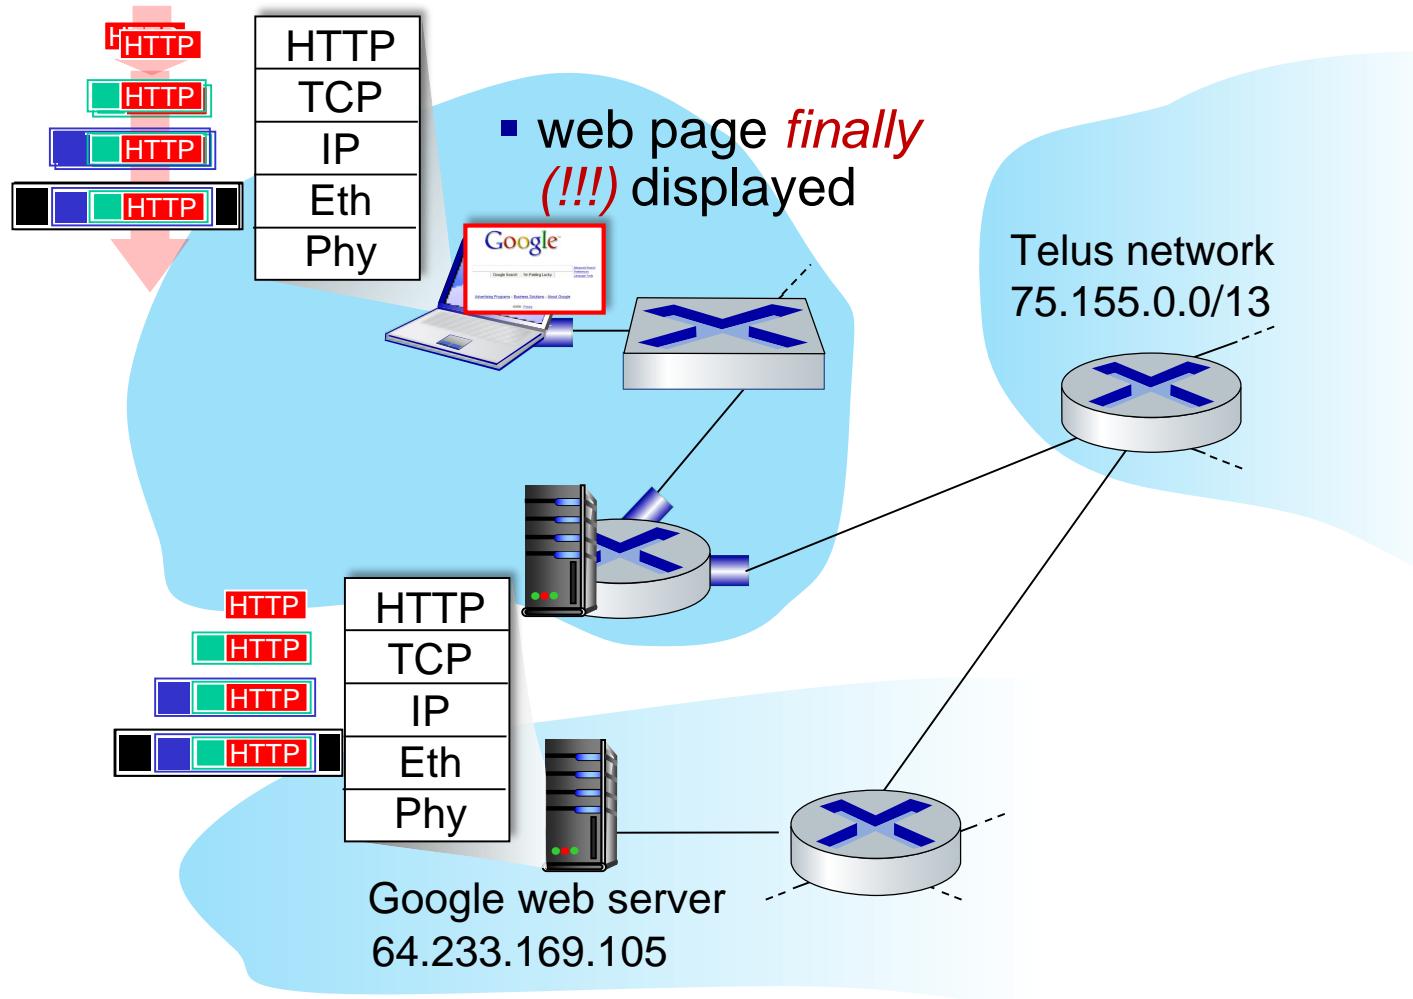
- de-muxed to DNS
- DNS replies to client with IP address of [www.google.com](http://www.google.com)

# A day in the life... TCP connection carrying HTTP



- to send HTTP request, client first opens **TCP socket** to web server
- TCP **SYN segment** (step 1 in TCP 3-way handshake) inter-domain routed to web server
- web server responds with **TCP SYNACK** (step 2 in TCP 3-way handshake)
- **TCP connection established!**

# A day in the life... HTTP request/reply



- **HTTP request** sent into TCP socket
- IP datagram containing HTTP request routed to [www.google.com](http://www.google.com)
- web server responds with **HTTP reply** (containing web page)
- IP datagram containing HTTP reply routed back to client

# A Brief Intro to Linux

## Origins of Linux

**Unix Inspiration:** Developed in the late 1960s at AT&T's Bell Labs.

**GNU Project (1983):** Launched by Richard Stallman to create a free OS.

## Linux Kernel Creation

**Linus Torvalds (1991):** Created the Linux kernel as a personal project.

**First Release (1991):** v0.01 released on Minix; gained a developer community.

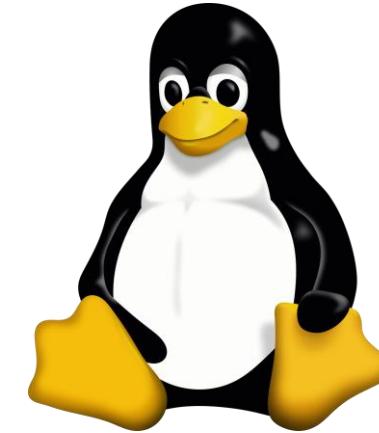
## Linux Growth

**1990s:** Collaboration with GNU tools created a fully functioning OS (GNU/Linux).

**90%+ of public cloud servers** run on Linux (AWS, Google Cloud, Azure).

**Android** powers over **2.5 billion active devices** worldwide,

**Billions of IoT devices** run Linux, particularly due to lightweight distributions



**debian**



**Red Hat**



**ubuntu**

**fedora** 



**CentOS**



**gentoo linux™**



**Amazon  
Linux**



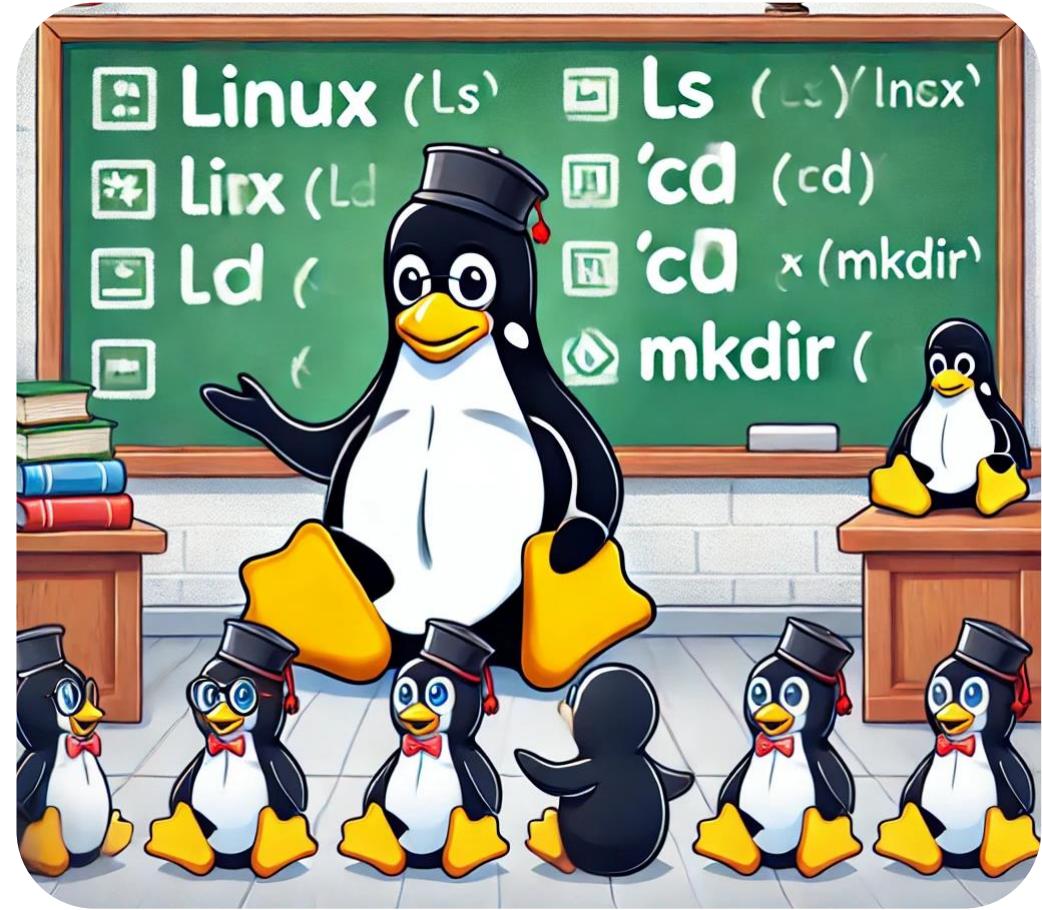
**alpine  
linux**

# A Brief Intro to Linux

- ✓ Master the Basics/Concepts
- ✓ Learn by Building Small Projects
- ✓ Embrace Curiosity
- ✓ Take Advantage of Online Materials
- ✓ Join Linux Communities

Some basic commands

{ cd, ls, pwd, hostname  
cp, mv, rm, mkdir, rmdir, man  
cat, vim, nano, touch, find, grep  
  
uname, df, du, tar, fdisk, whoami  
  
top, ps, kill, jobs, init, uptime, date  
  
sudo, yum/apt, ssh, scp, history  
  
curl, wget, ping, mtr, ip, ifconfig  
netstat, iptables

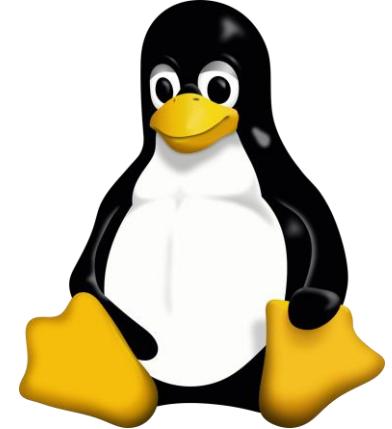


## Some Nice References

<https://linuxcommand.org/>

<https://linuxjourney.com/>

<https://linuxsurvival.com/>



1. Visit [\*\*https://cocalc.com/features/terminal\*\*](https://cocalc.com/features/terminal)
2. Sign up the site.
3. Open “**New > Linux Terminal**” in your browser.

**cd, ls, pwd**  
**cp, mv, rm, mkdir, rmdir, man**  
**cat, vim, nano, touch, find, grep**

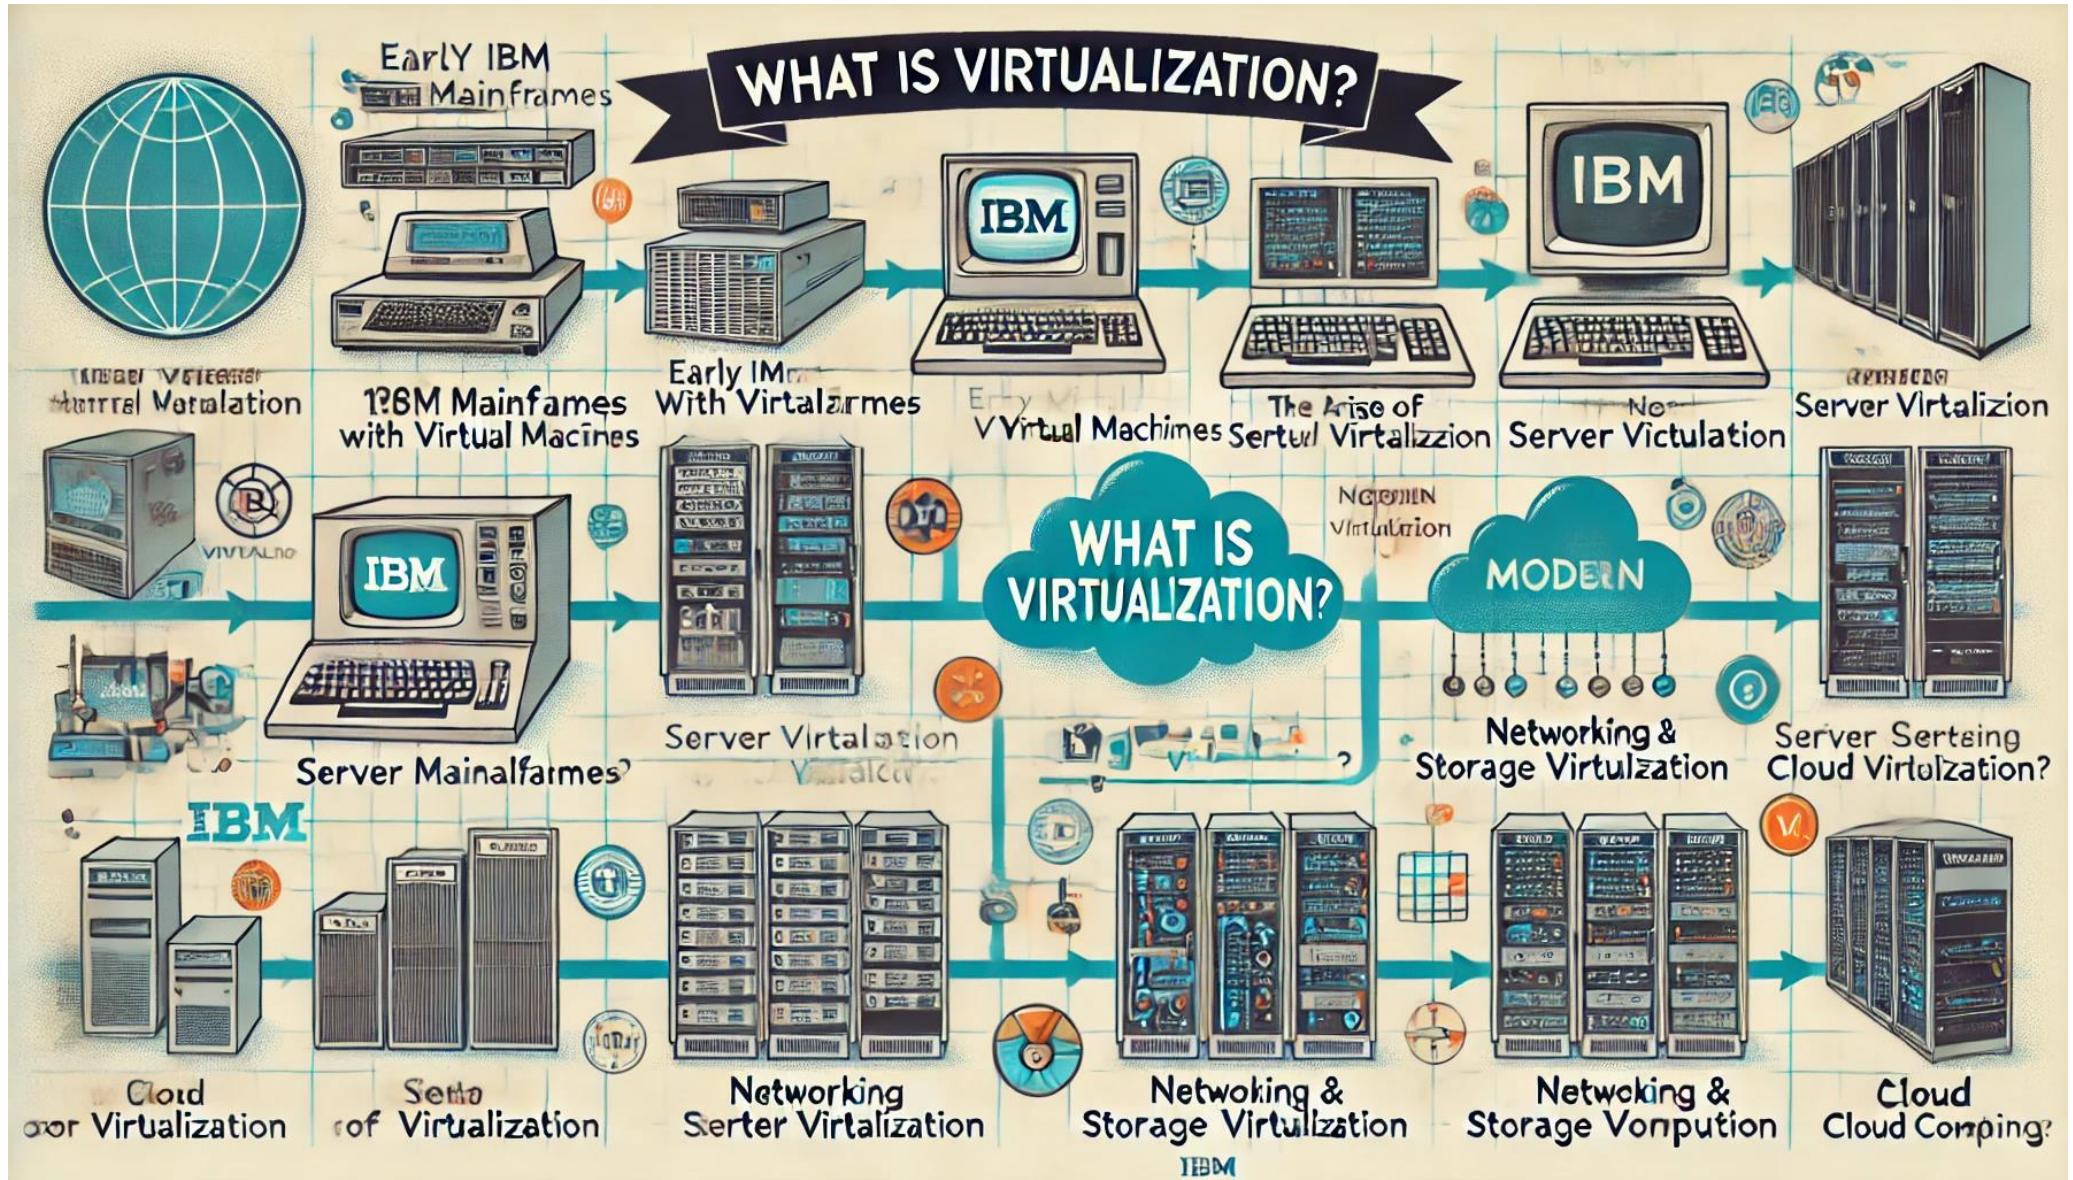
**uname, df, du, tar, fdisk**

**top, ps, kill, jobs, init**

**sudo, yum/apt, ssh, scp**

**curl, wget, ping, mtr, ip**  
**netstat, iptables**

# What is Virtualization?

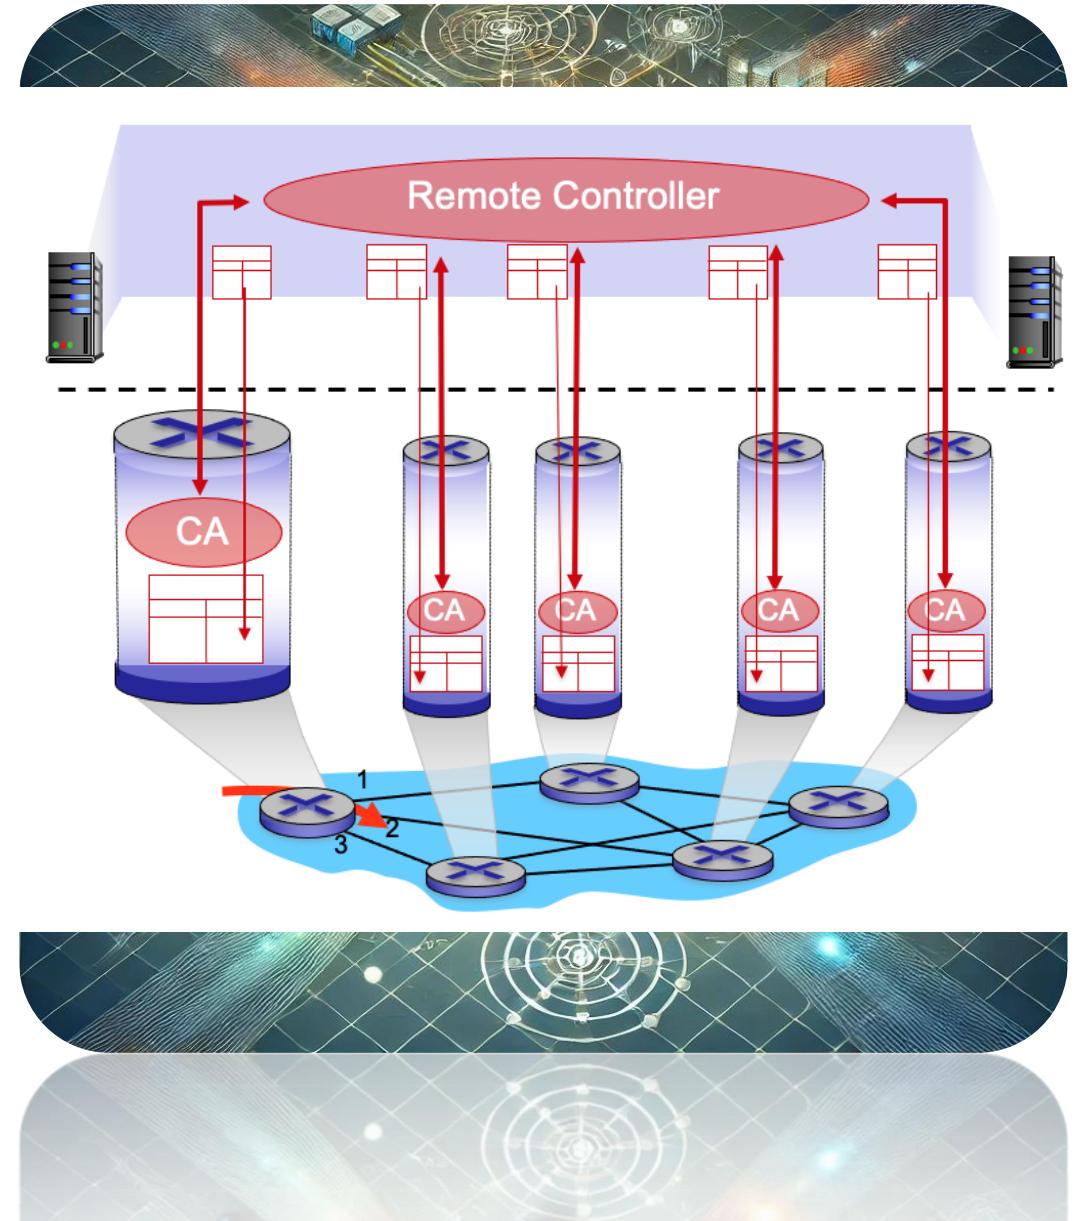


# Network Virtualization

Abstracts physical network resources, such as Switches and Routers into Logical or Virtual Networks.

**VLANs** (Virtual LANs) or **VPNs** (Virtual Private Networks), **NFV** (Network Function Virtualization), **Overlay Networks** allow multiple networks to run over a single infrastructure.

Software-Defined Networking (**SDN**) is one the these technologies that separates control and data planes, allowing dynamic network management

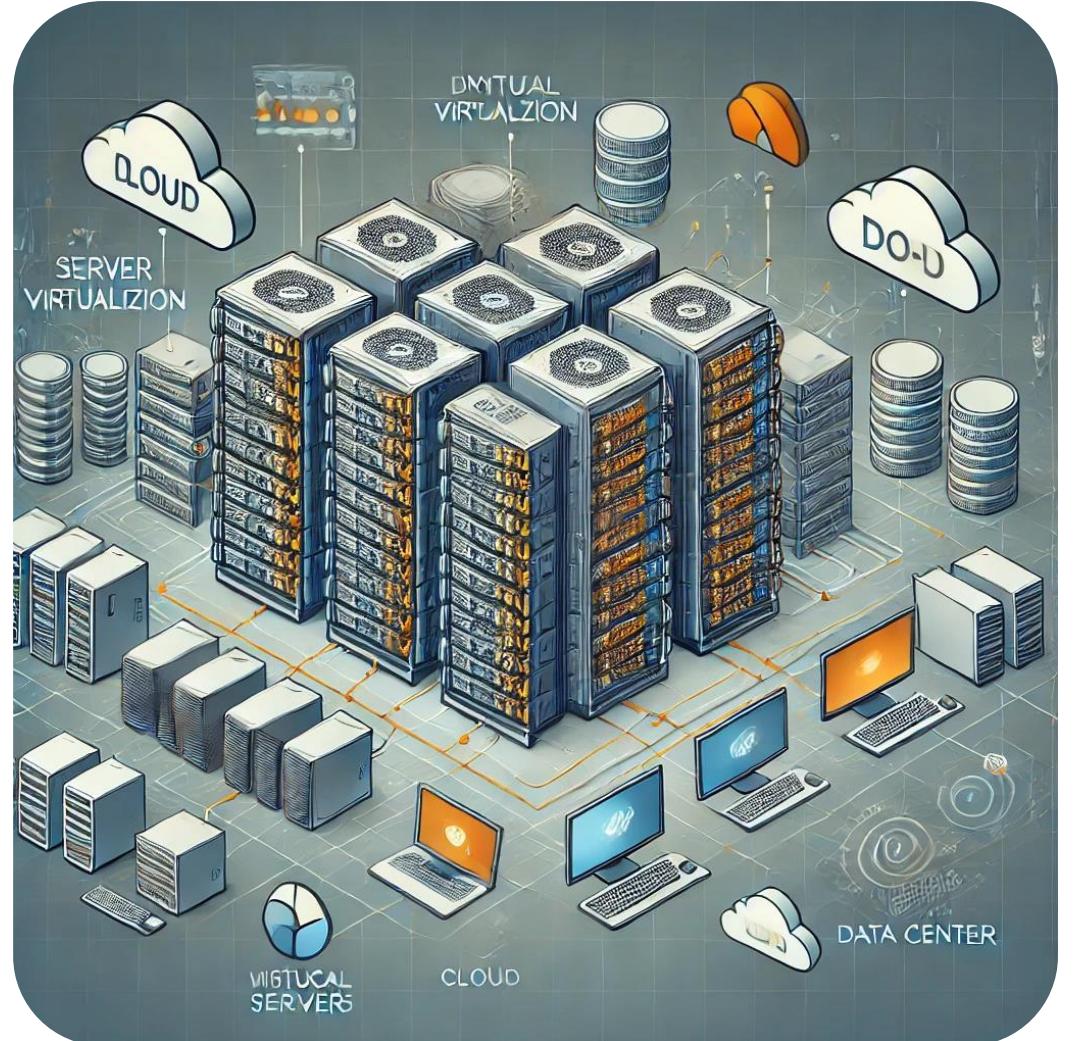


# Server Virtualization

Multiple virtual machines (VMs) run on a single physical server

VMware, Hyper-V, and KVM allow servers to host isolated environments (VMs), each running a different Operating System or Application (can share the same Kernel).

Many Cloud providers such as AWS/GCP/Azure use virtualization to provision virtual servers to users.

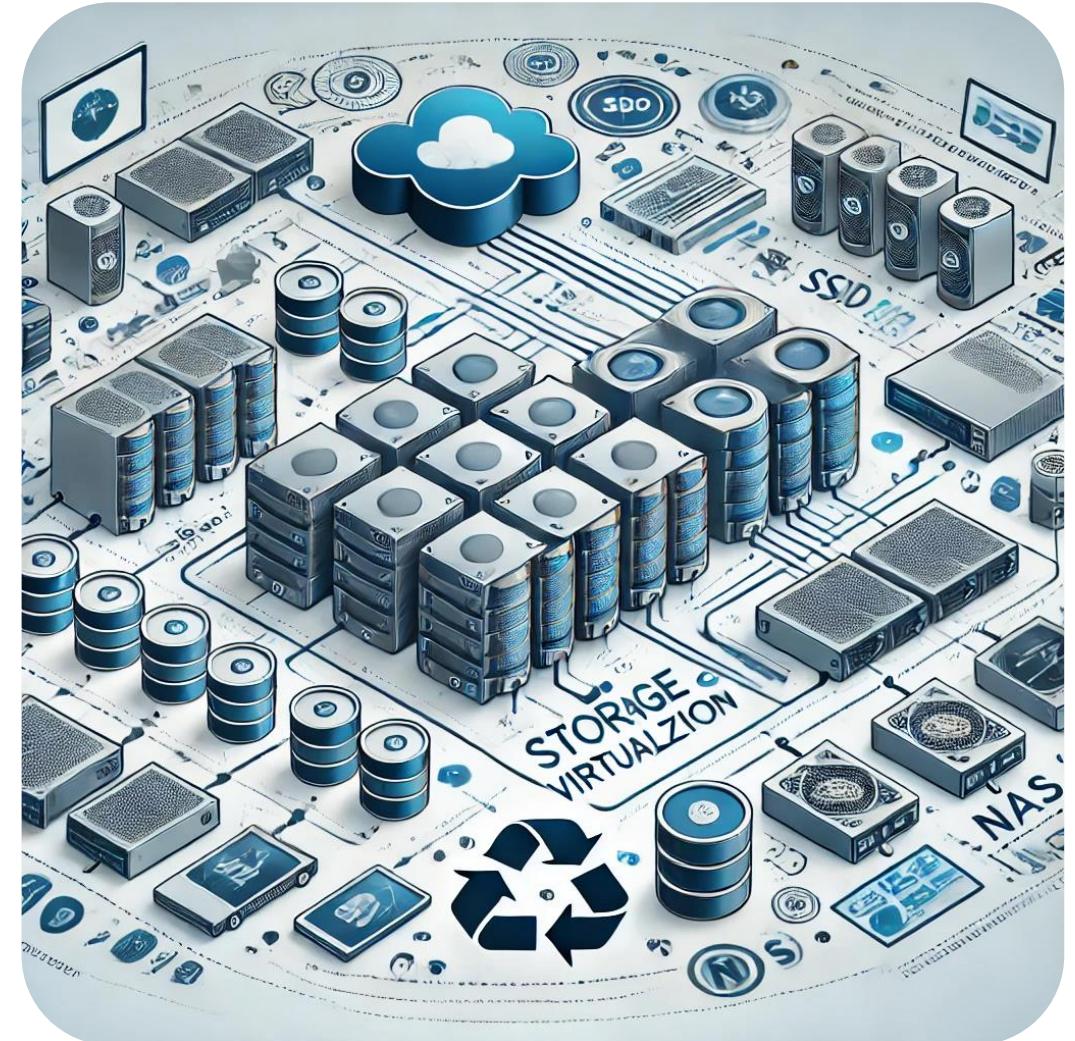


# Storage Virtualization

Combines multiple physical storage resources into a single virtual storage pool.

SAN (Storage Area Networks) or cloud storage solutions like Amazon S3. Ceph, vSAN, EBS, and Cinder also some of the known virtual storage technologies.

Enterprises use virtual storage to manage and allocate space efficiently across various applications.



# Cloud Computing (Cloud Technology)

## What is Cloud Technology?

**Cloud Technology** enables access to computing resources like storage, databases, servers, and software over the internet.

## Types of Cloud Models

**Public Cloud:** Shared resources provided by third-party vendors.

**Private Cloud:** Dedicated resources for one organization.

**Hybrid Cloud:** Combination of public and private (on-premise) clouds.

**Multi-Cloud:** Use of multiple cloud services from different providers.

## Benefits

**Scalability:** Adjust resources based on demand.

**Cost Efficiency:** Pay-as-you-go pricing models.

**Flexibility:** Access from anywhere with internet.

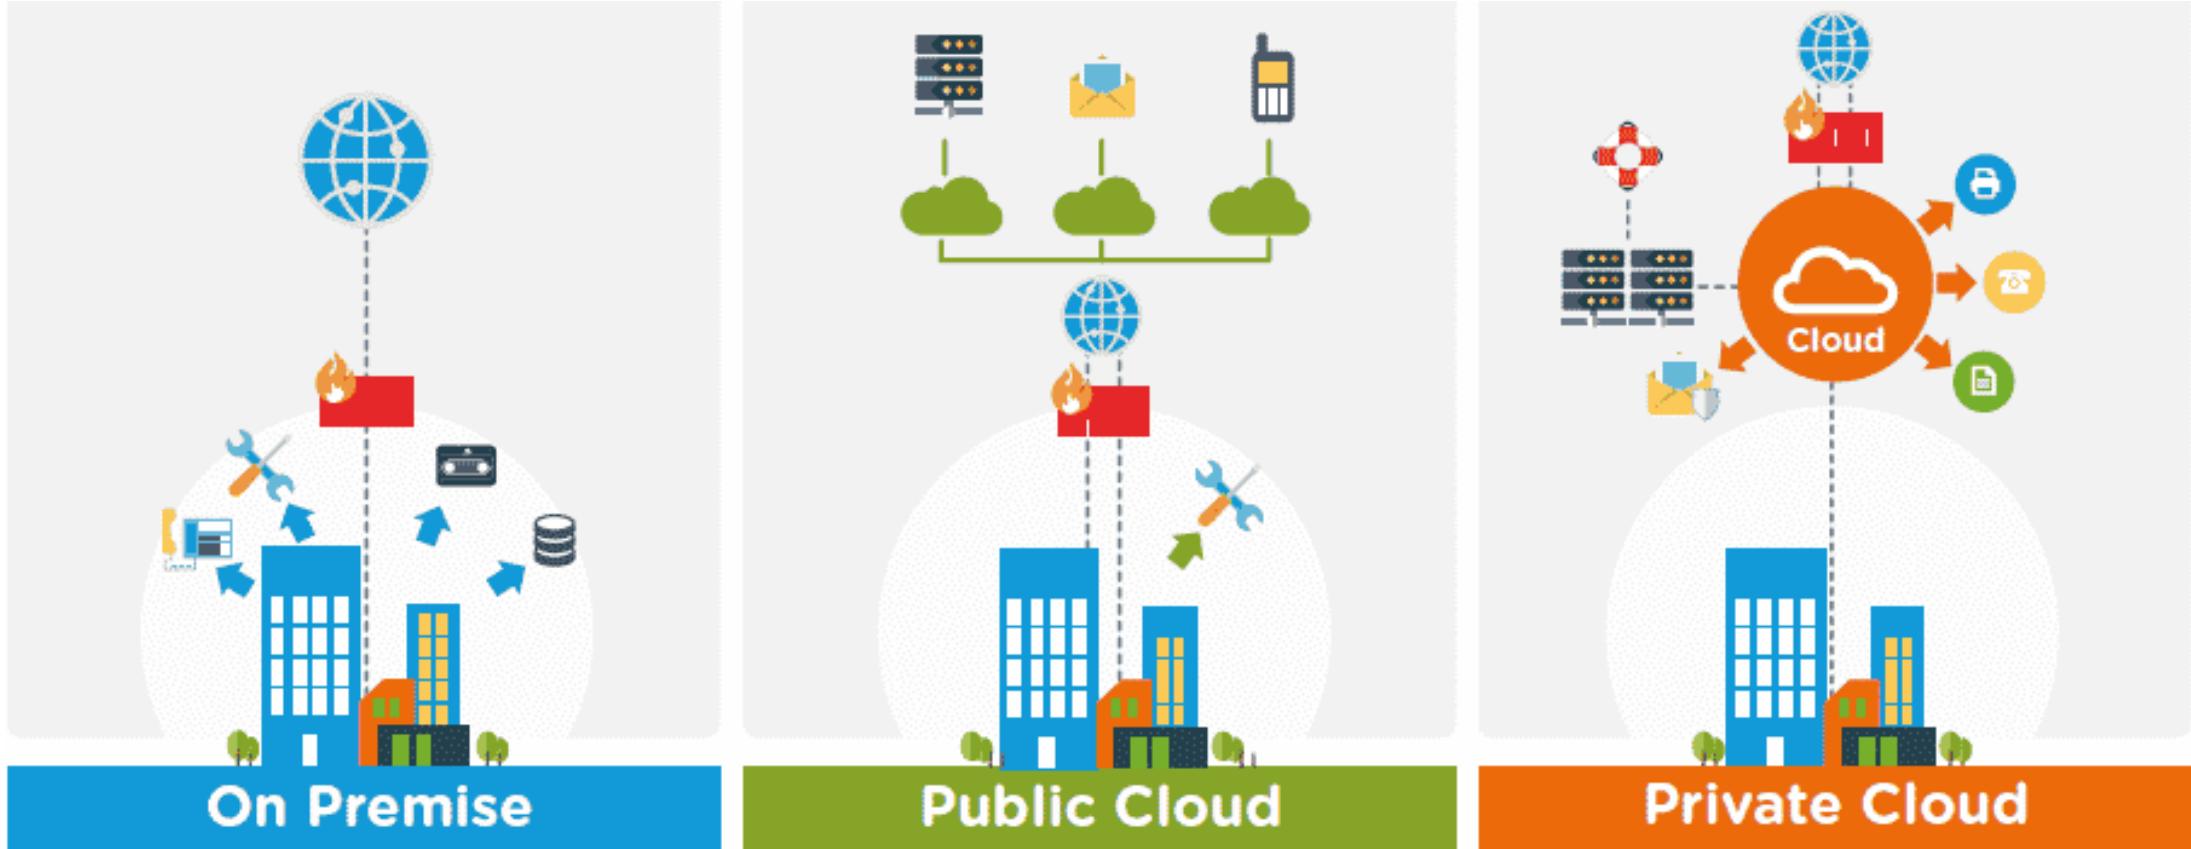
**Collaboration:** Shared environments for teams and organizations.



# Cloud Computing (Public Cloud)



# Cloud Computing (Cloud Technology)



TM

# Cloud Computing (Cloud Technology)

Cloud service models provide different levels of management, from full applications to bare infrastructure

## 1. SaaS (Software-as-a-Service):

Minimal user management (just access and settings).



## 2. PaaS (Platform-as-a-Service):

User manages applications and data, but not infrastructure.



## 3. IaaS (Infrastructure-as-a-Service):

User manages applications, data, and operating systems, but the provider manages hardware.



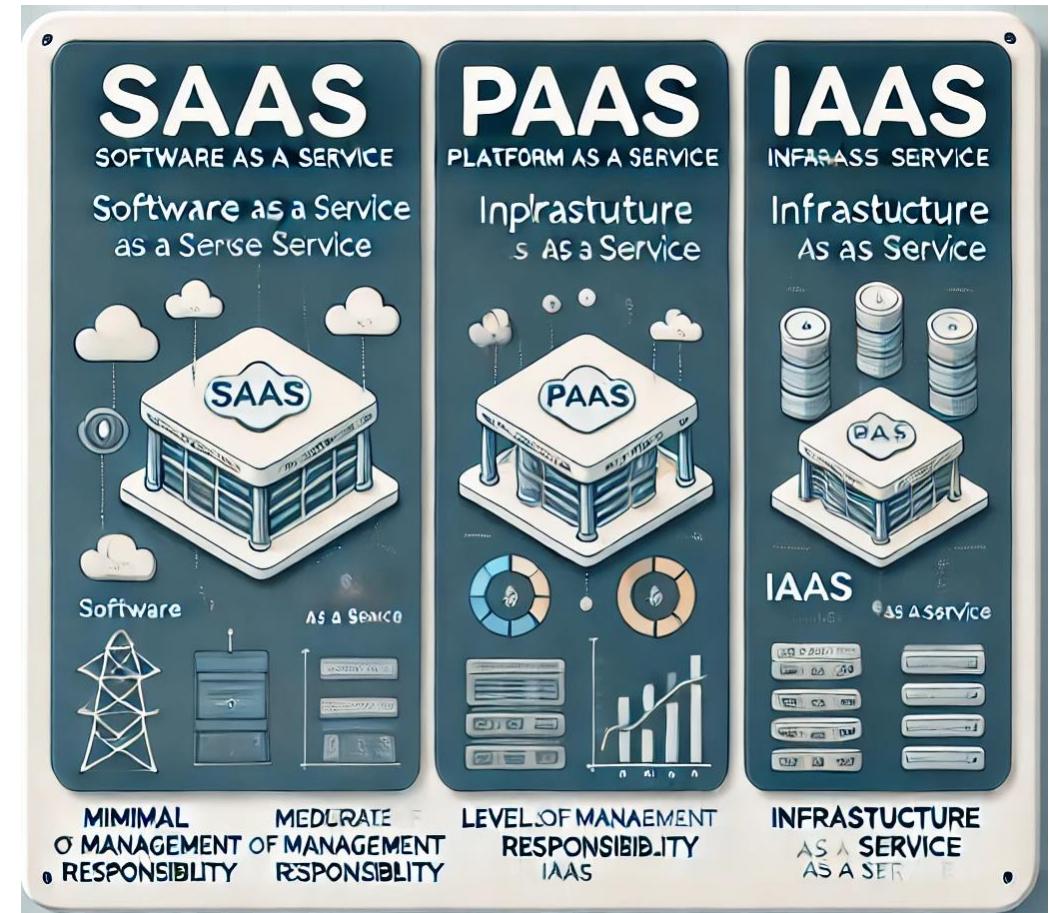
Amazon EC2



Google Cloud Platform

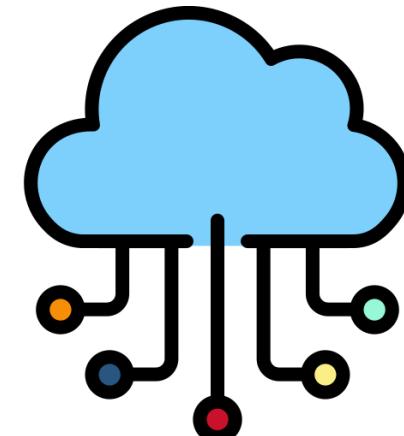
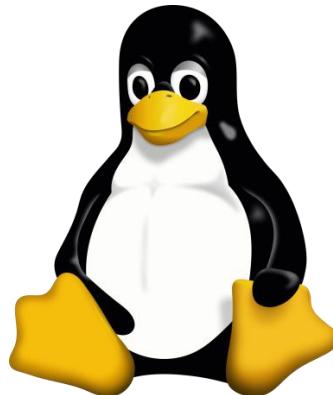
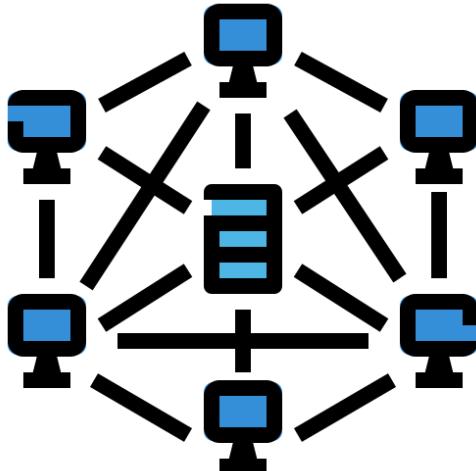


Azure

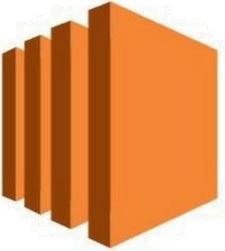


# Lecture #1 Summary

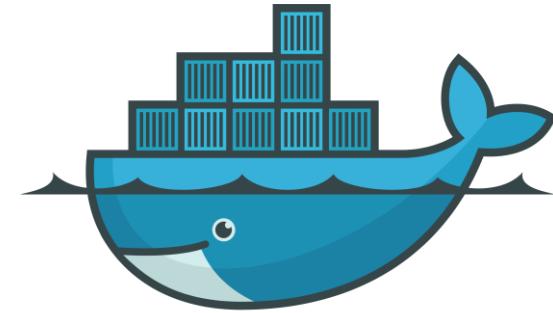
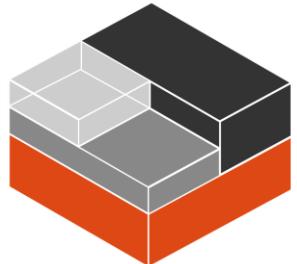
1. Quick review of **Computer Networking** (Routing/Switching/IP/TCP/UDP)
2. Quick recap of Basic **Linux** Commands
3. Quick recap of the **Virtualization** and **Cloud Computing** Concepts



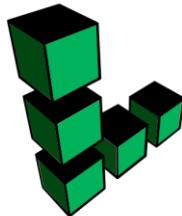
# LAB/Project Requirements



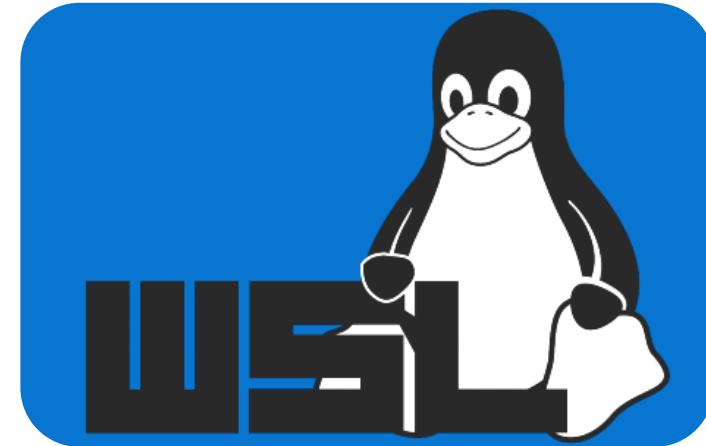
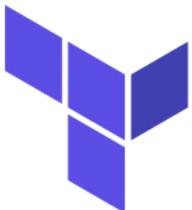
Amazon EC2



docker



linode



# End of Lecture #1



THANK  
YOU